

Tutorial


filtering, morphology and canny edge detection

Ziang Cheng

Outline

- Image filtering
 - Review
 - Boundary handling (padding)
 - Image morphology
- Edge detection algorithm
 - F1 evaluation metric


Spatial image filtering



$$e_{processed} = \mathcal{G}(\Omega_e * F)$$

1. Compute the stacked response of F on Ω_e with some element-wise operator $*$
2. Gather the response to a single value through some reduction function $\mathcal{G}(\cdot)$
3. Slide the window to next pixel


Cross-correlation/convolution



$$e_{processed} = \mathcal{G}(\Omega_e * F) = v * e + r * a + s * b + t * c + u * d + w * f + x * g + y * h + z * i$$

- $*$ is element-wise multiplication operator
- $\mathcal{G}(\cdot)$ is the summation function
- $*$ is not to be confused with convolution operator $*$
- Processed image is linear w.r.t. input

Cross-correlation/convolution



*

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Laplacian




$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

Sharpen




$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Mean blur




$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

Gaussian 3x3




$$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

Gaussian 5x5



Median filter




- $*$ is the element-wise multiplication
- $\mathcal{G}(\cdot)$ is the median function
- F is $\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

Filters with content-aware kernels

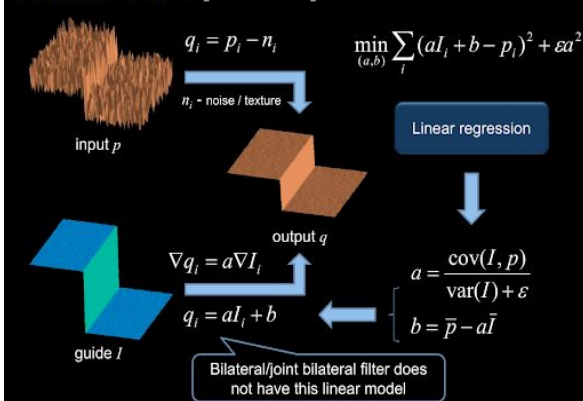
Bilateral Filter [Aurich 95, Smith 97, Tomasi 98]

No Averaging across Edges




The kernel shape depends on the image content.

Guided filter [He 2010]



Doubly-guided filters [Lu 2017]




Gaussian filter (implementation)

- Separability: Gauss kernels can be factorized $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$
 - Since convolution operator $*$ is associative

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \text{Image} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \text{Image} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * (\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \text{Image})$$

- Complexity of filtering with $n \times n$ image and $m \times m$ kernel is $O(n \times n \times m \times m)$
- Complexity of filtering with $n \times n$ image and **separable** $m \times m$ kernel is $O(2 \times n \times n \times m)$
- $m/2$ times faster implementation for Gaussian filtering!
- 3-sigma rule:
 - Pixels more than 3σ away can be ignored
 - No need for kernel size greater than $(6\sigma+1)$



Padding (why padding?)

0	2	2	0	1
2	1	0	1	1
2	1	1	0	2
0	0	2	2	1
1	2	2	0	2



x	x	x
x	x	x
x	x	x

Without padding, the edges of the image are only partially processed, and the result of convolution is smaller than the original image size

0	0	0
0	0	1
1	1	0

0	0	0	0	0	0	0
0	0	2	2	0	1	0
0	2	1	0	1	1	0
0	2	1	1	0	2	0
0	0	0	2	2	1	0
0	1	2	2	0	2	0
0	0	0	0	0	0	0

0	0	0
0	0	1
1	1	0

Padding

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	1	2	3	4	5	0	0	0
0	0	6	7	8	9	10	0	0	0
0	0	11	12	13	14	15	0	0	0
0	0	16	17	18	19	20	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Constant (zero)

13	12	11	12	13	14	15	14	13	
8	7	6	7	8	9	10	9	8	
3	2	1	2	3	4	5	4	3	
8	7	6	7	8	9	10	9	8	
13	12	11	12	13	14	15	14	13	
18	17	16	17	18	19	20	19	18	
13	12	11	12	13	14	15	14	13	
8	7	6	7	8	9	10	9	8	

Mirror/Symmetric

1	1	1	2	3	4	5	5	5	
1	1	1	2	3	4	5	5	5	
1	1	1	2	3	4	5	5	5	
6	6	6	7	8	9	10	10	10	
11	11	11	12	13	14	15	15	15	
16	16	16	17	18	19	20	20	20	
16	16	16	17	18	19	20	20	20	
16	16	16	17	18	19	20	20	20	

Replicate



Original Image




Filtered Image with Black Border




Filtered Image with Border Replication

Morphological filter: Dilation



- Input is a binary image (True=1, False=0), padded with constant 0
- $*$ is the logical and
- $\mathcal{G}(\cdot)$ is the logical or reduction
- F is $\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

Morphological filter: Erosion




$$e_{processed} = \mathcal{G}(\Omega_e * F)$$

- Input is a binary image (True=1, False=0), padded with constant 1
- $*$ is the logical and
- $\mathcal{G}(\cdot)$ is the logical and reduction
- F is $\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

Q3 Edge detection algorithm (overview)

Key steps:

1. Pre-processing: smoothing e.g. Gauss filter (q.3b)
2. Compute intensity of gradients: e.g. Sobel filter, etc.
 - a. With appropriate boundary handling (q.3a)
3. NMS along gradient direction (q.3c)
4. Returns boundary score



How to evaluate the performance?


- We have a test set with ground truth binary labels for each pixel
 - True: is edge pixel
 - False: not edge pixel
- We have a model that outputs a score in $[0,1]$
 - How to make this model output binary decision? ($\text{bin_model} = \text{score} + \text{threshold}$)
 - How to select threshold?
- Edges are sparse: very few pixel are actually edge pixels
 - Measuring accuracy may not be the best idea...

Precision, recall and accuracy

- Precision: $TP/(TP+FP)$
- Recall: $TP/(TP+FN)$
- Accuracy: $(TP+TN)/(TP+FP+TN+FN)$

Suppose only 1% of pixels are true edge pixels in test sets, what are the precision, recall and accuracy of a model that randomly predicts edge score with $th=0.99$?

- Precision: 1%
- Recall : 1%
- Accuracy: 98%




How are precision and recall related?


Generally, by increasing threshold

- Model gets more pessimistic (less likely to say yes)
- Precision likely increase
- Recall likely decrease


The precision-recall curve



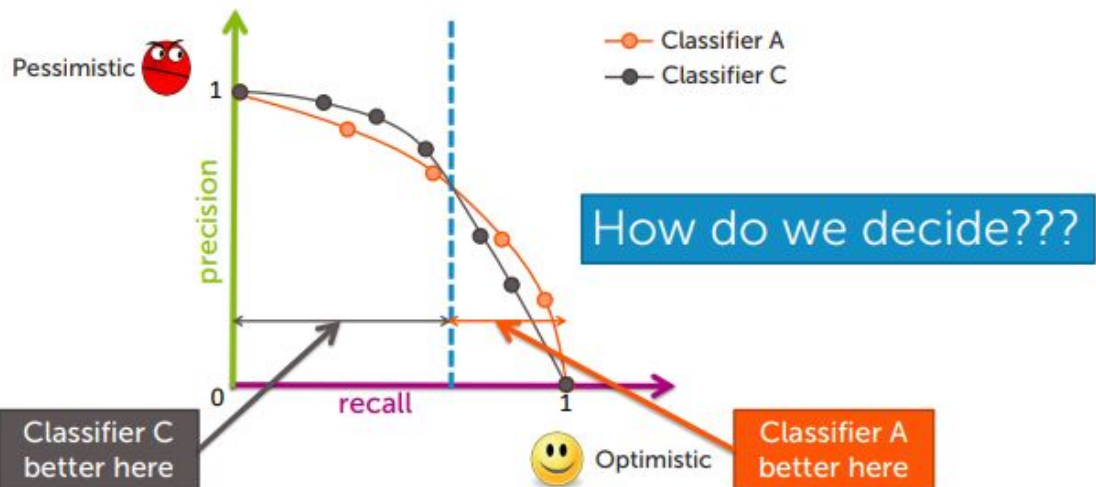
What does the perfect algorithm look like?



Which classifier is better? A or B?



Which classifier is better? A or C?



How to select the best predictor/threshold for a task?

Whichever maximizes F_β score:

$$F_\beta = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

$\beta > 1$ attaches greater importance to recall than to precision.

F_1 Score:

$$F_1 = \left(\frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} \right) = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$