

# ENGN6528 Computer Vision

## Semester 2, 2020



**Thalaiyasingam Ajanthan**

Majority of the slides from Miaomiao Liu

# Week-1-part2

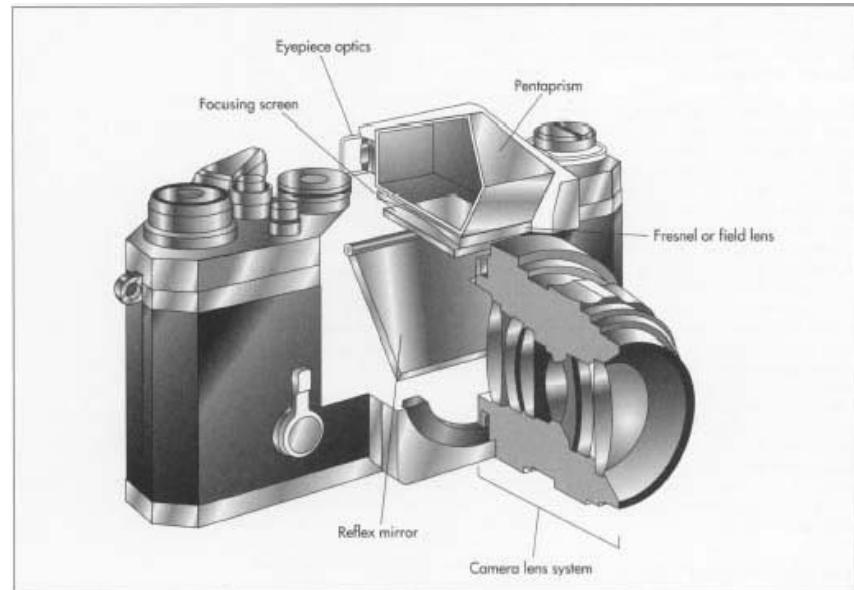
Image Formation, Representation

# Outline

---

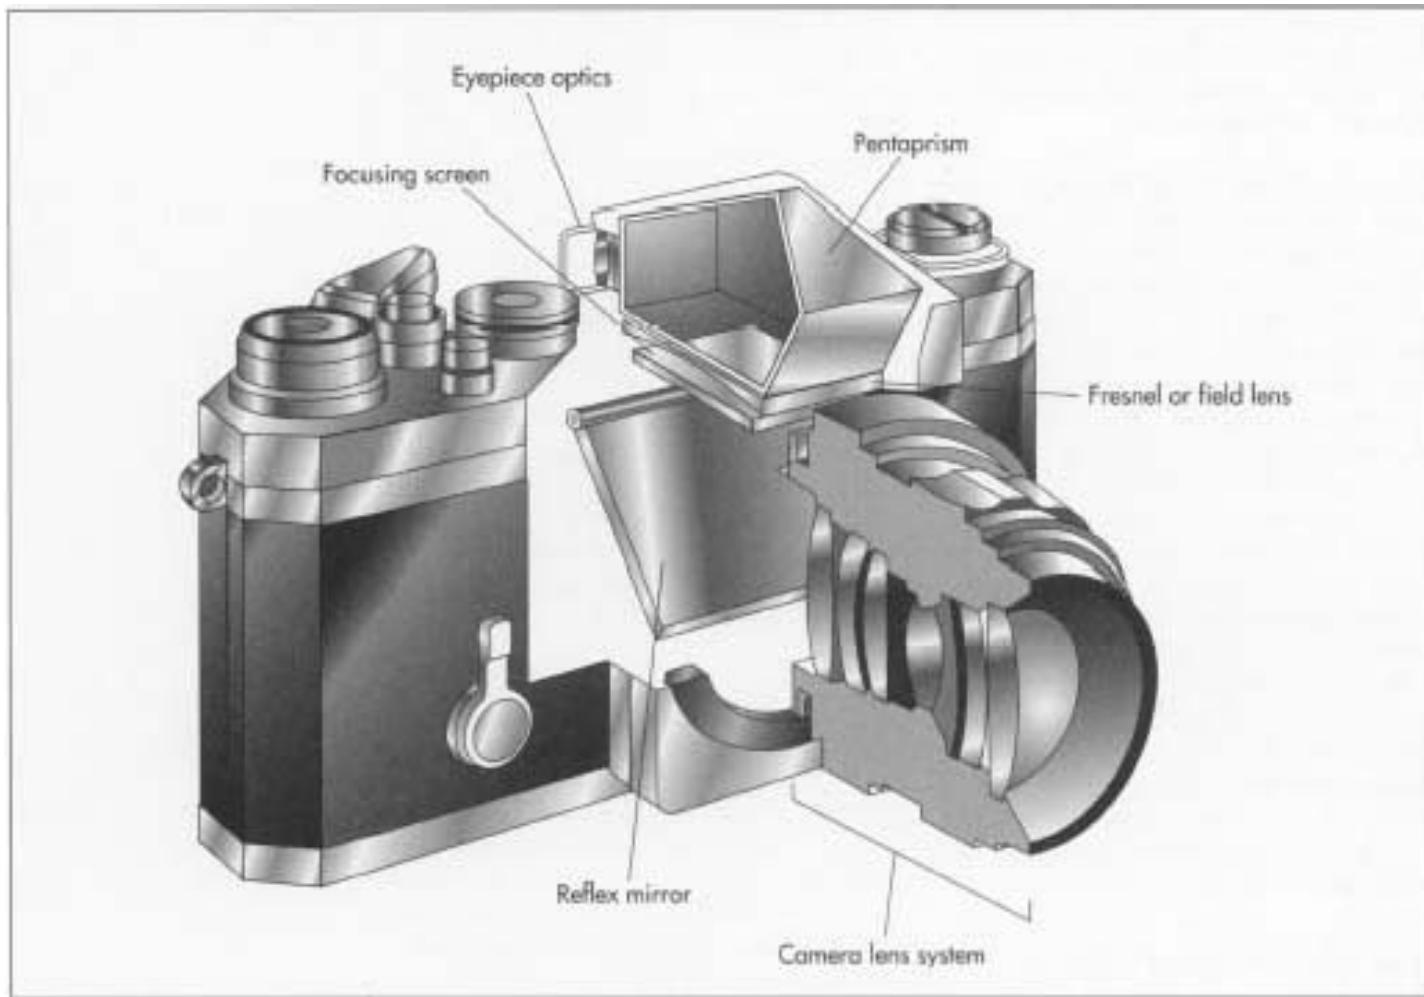
- What is a camera ?
- Geometric image formation
- Photometric image formation
- Image representation
  - Colour space

# What is a camera ?



**Slides from: F. Durand, S. Seitz, S. Lazebnik, S. Palmer**

# What is a camera ?



**Slides from: F. Durand, S. Seitz, S. Lazebnik, S. Palmer**

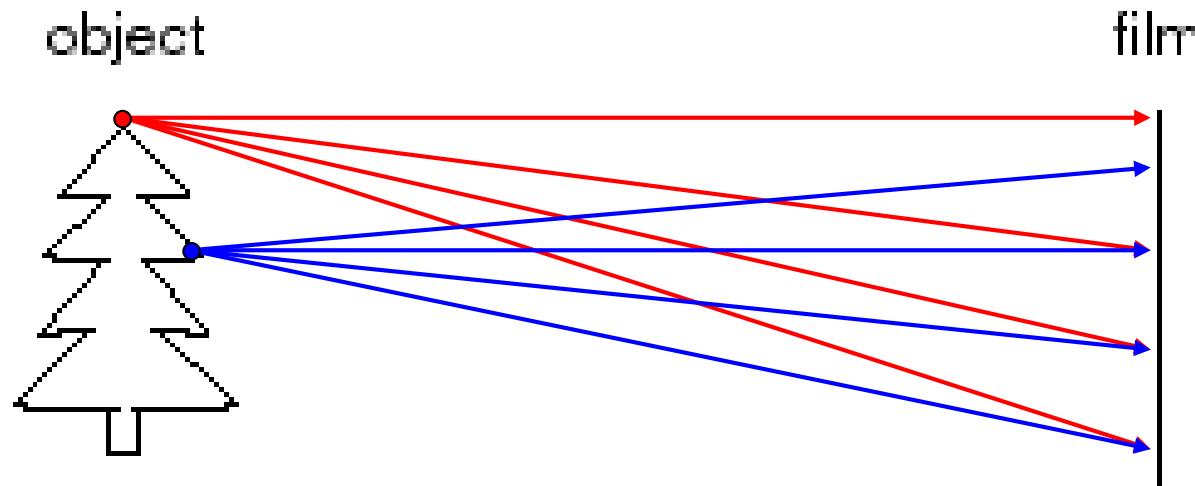
# Let us design a camera

---



# Let's design a camera

---



Idea 1: put a piece of film in front of an object

Will we get an image?

# Turn any room into a camera

(Accidental pinhole, A. Torralba, MIT, published in IEEE CVPR 2012)

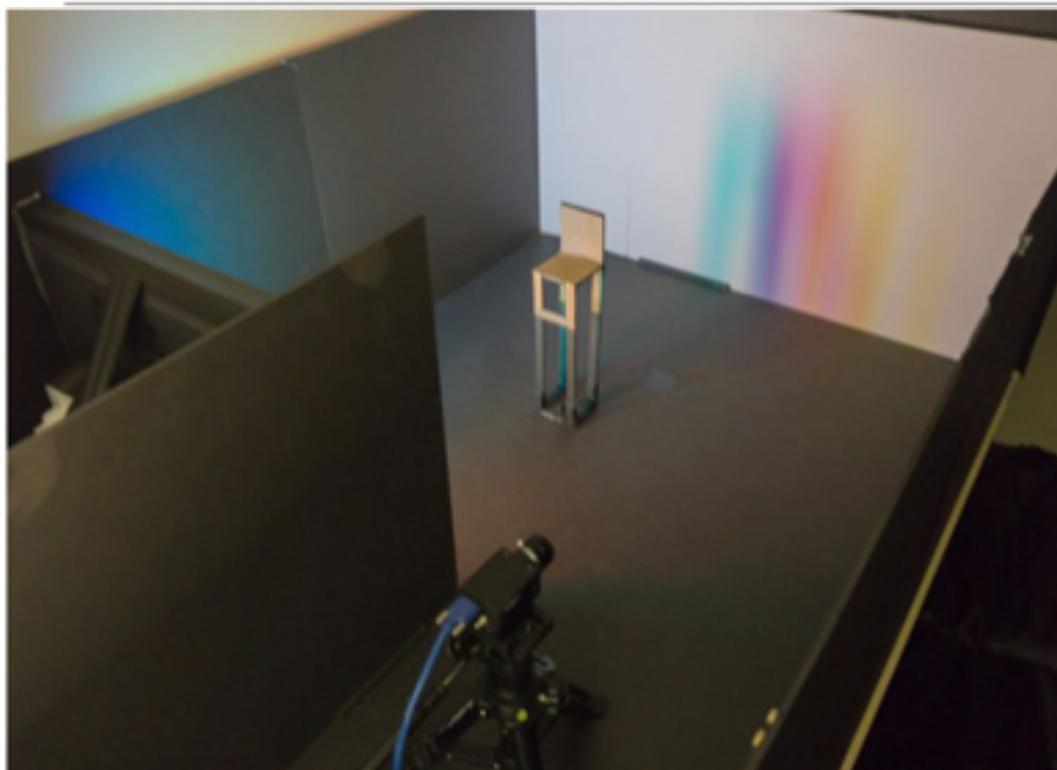


# A new paper on *Nature*

NEWS · 23 JANUARY 2019

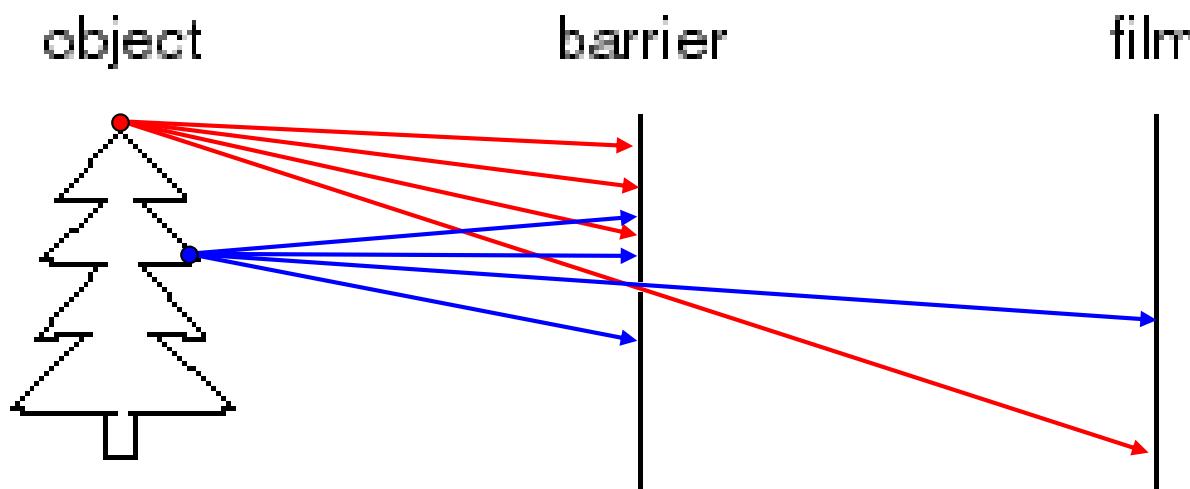
## How an ordinary camera can see around corners

Digital cameras have been used to reconstruct rough images of hidden objects just by analysing light that bounces off a wall.



# Pinhole camera

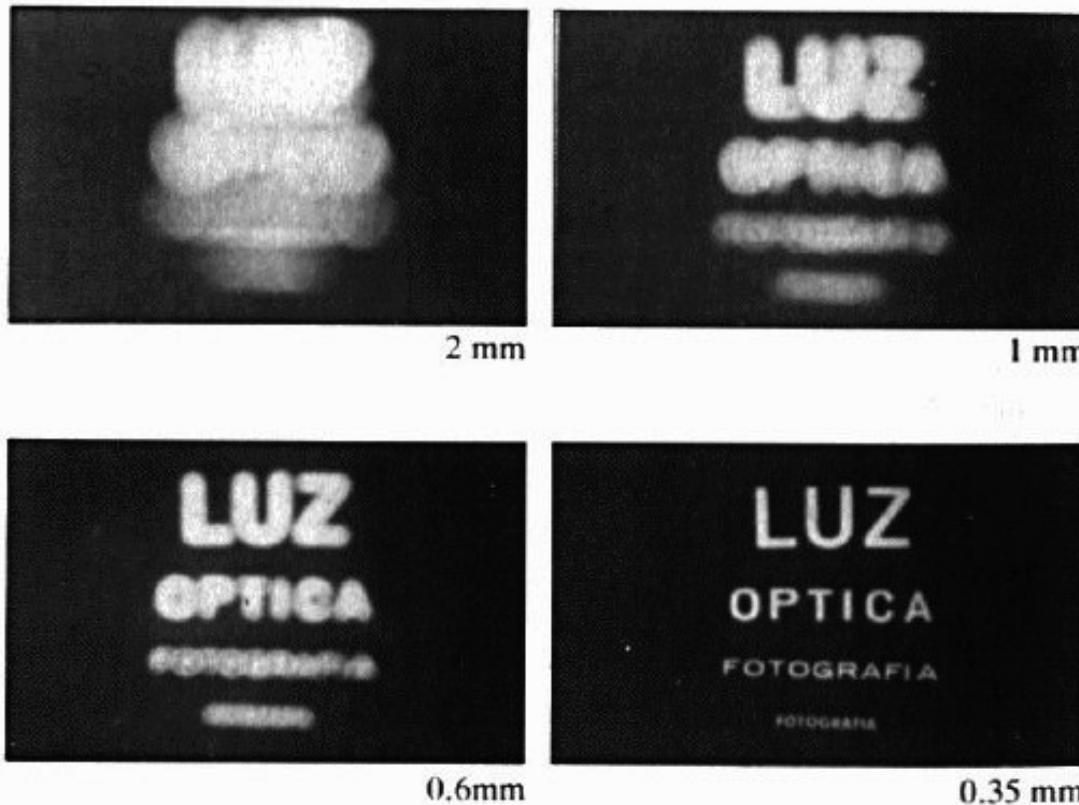
---



Add a barrier to block off most of the rays

- This reduces blurring
- The opening is known as the ***aperture***

# Shrinking the aperture

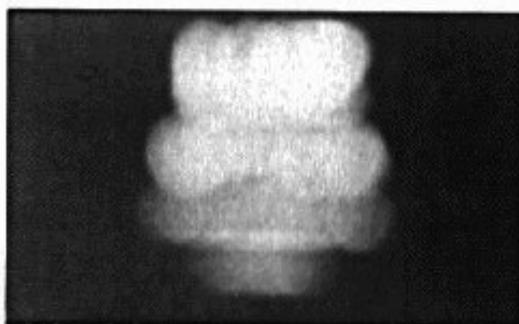


Why not make the aperture as small as possible?

- Less light gets through
- Diffraction effects...

# Shrinking the aperture

---



2 mm



1 mm



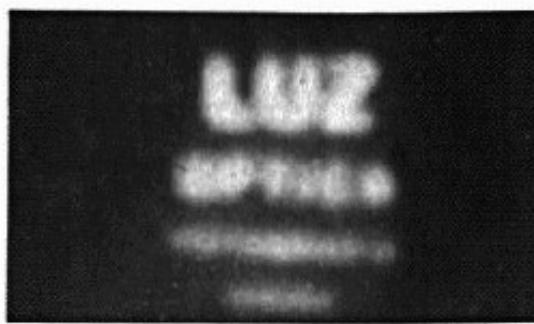
0.6mm



0.35 mm



0.15 mm



0.07 mm

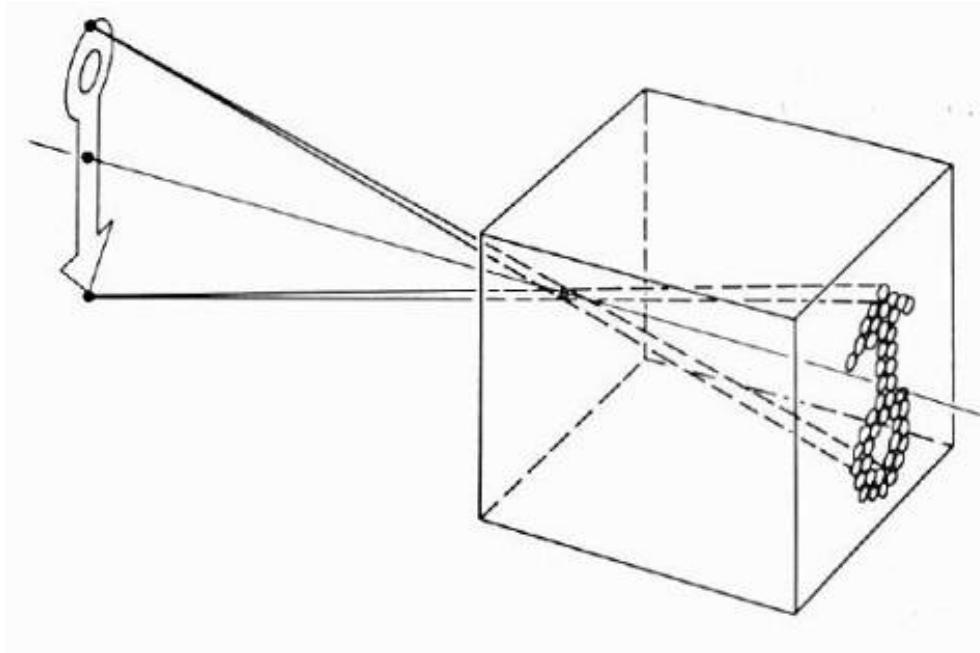
# Home-made pinhole camera

---



# Pinhole camera model

---

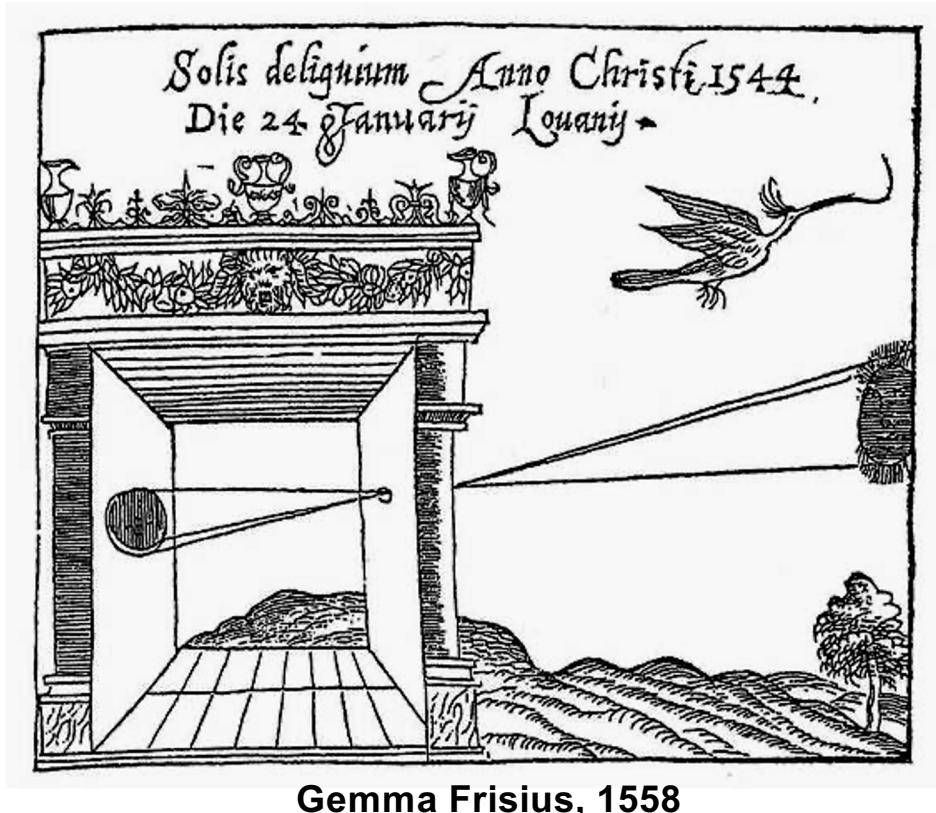


Pinhole model:

- Captures ***pencil of rays*** – all rays through a single point
- This point is called ***Center of Projection (focal point)***
- The image is formed on the **Image Plane**

# Camera Obscura

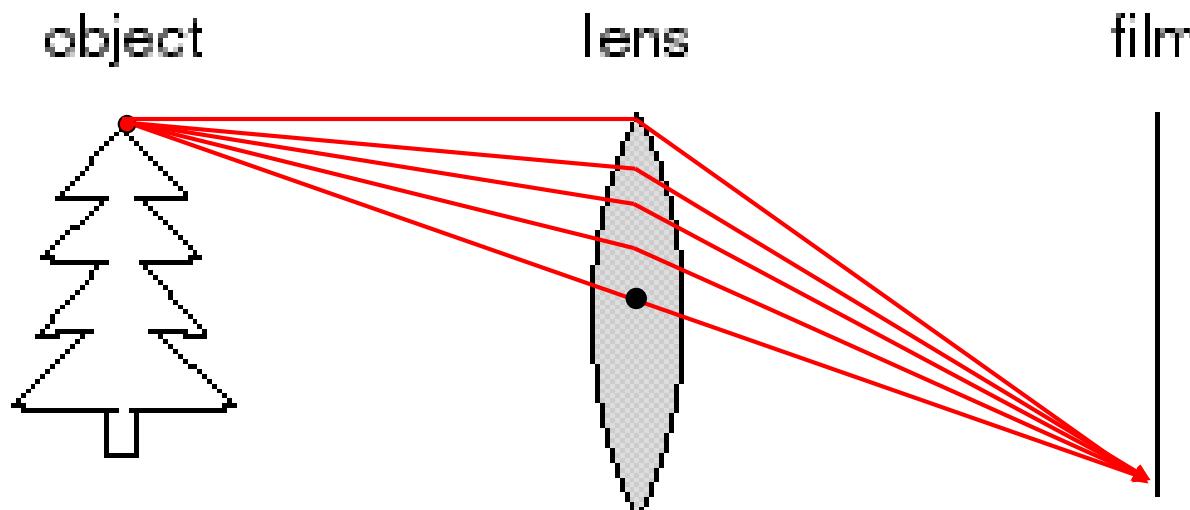
---



- Basic principle known to Mozi (470-390 BC, China), Aristotle (384-322 BC, Greek)
- Drawing aid for artists: described by Leonardo da Vinci (1452-1519)

# Adding a lens... to capture more light

---

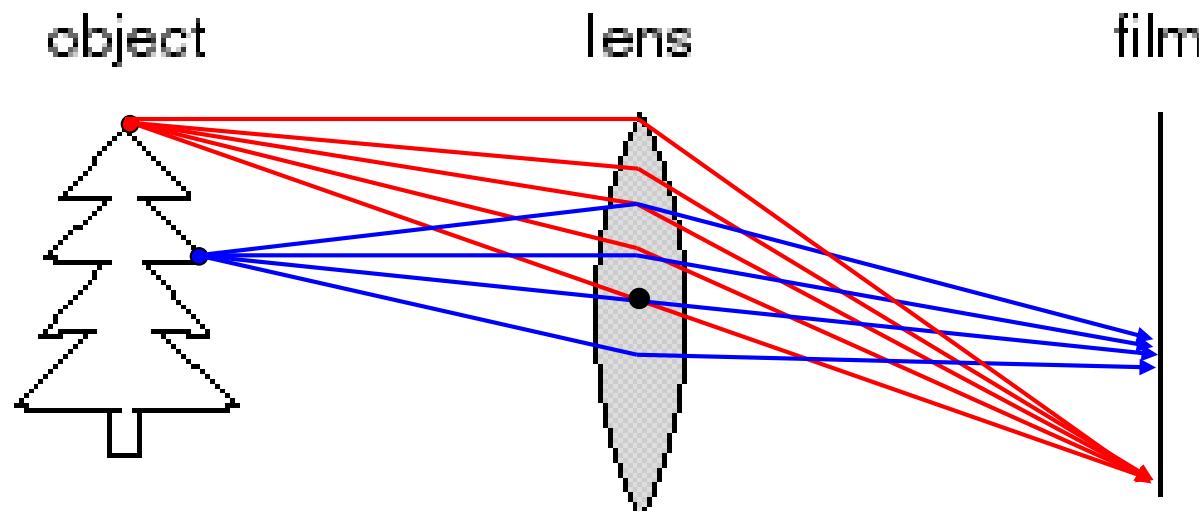


A lens focuses light onto the film

- Rays passing through the center are not deviated

# Adding a lens

---

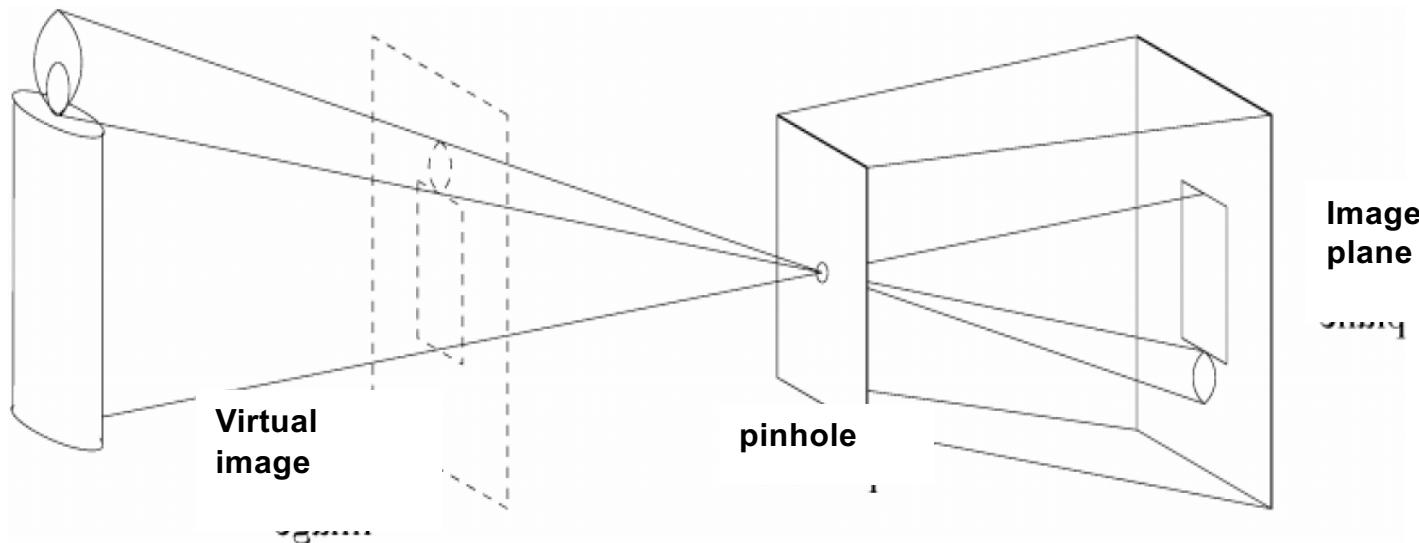


---

# Geometric Image Formation: Pinhole Camera Model

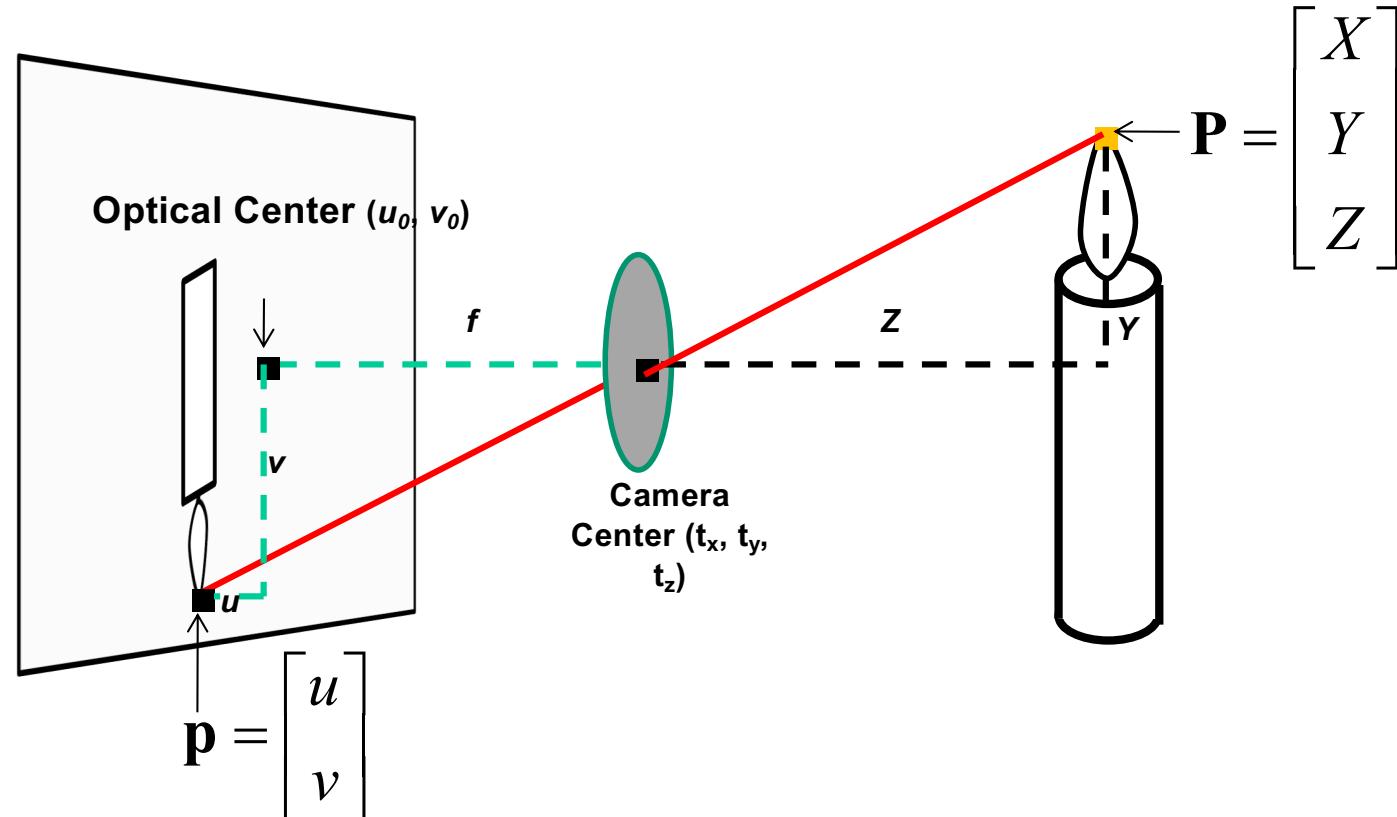
# Pinhole camera

- Pinhole camera is a simple model to approximate imaging process: **perspective projection**.



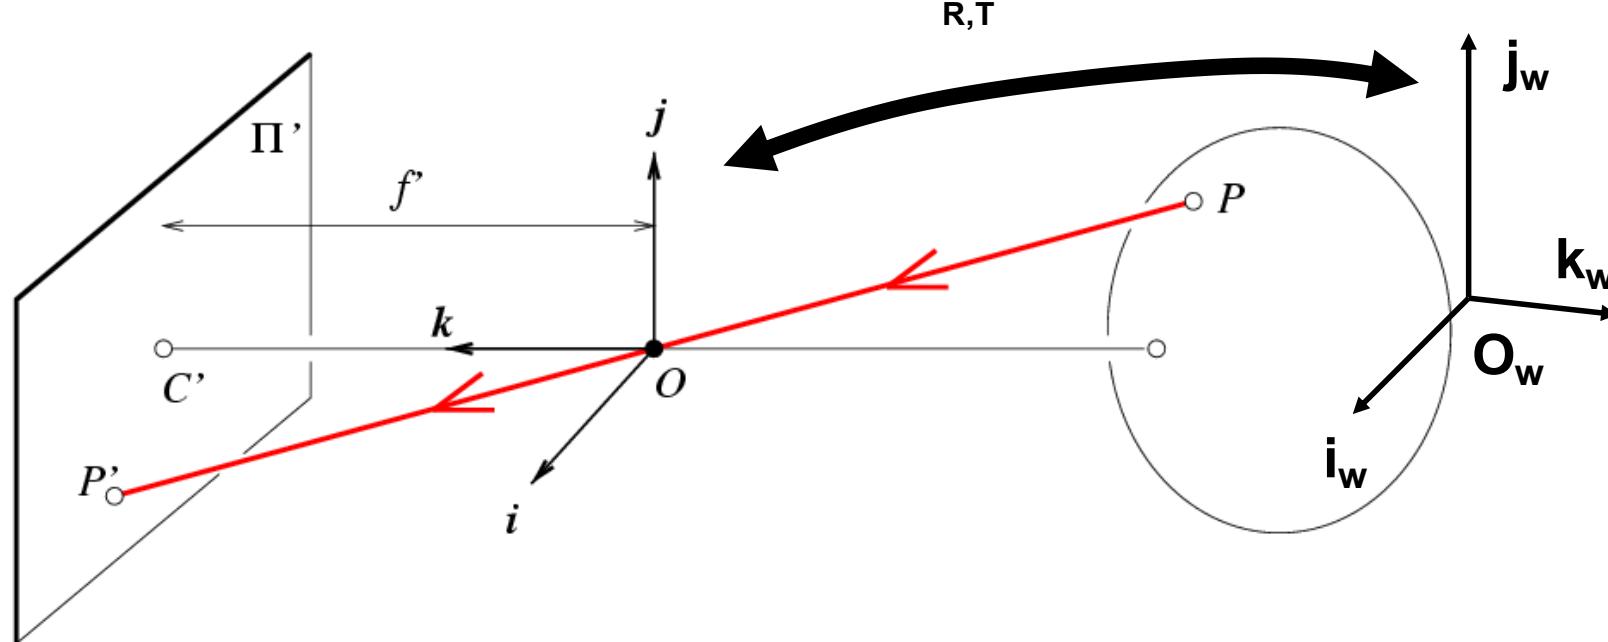
- If we treat pinhole as a point, only one ray from any given point can enter the camera.

# Projection: world coordinates → image coordinates



# General Camera Projection Matrix

---



$$\mathbf{x} = \mathbf{K} [\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$

$\mathbf{x}$ : Image Coordinates:  $(u, v, 1)$

$\mathbf{K}$ : Intrinsic Matrix (3x3)

$\mathbf{R}$ : Rotation (3x3)

$\mathbf{t}$ : Translation (3x1)

$\mathbf{X}$ : World Coordinates:  $(X, Y, Z, 1)$

# Homogeneous coordinates Representation

---

## Conversion

- Converting *from* Cartesian *to* homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous **image**  
coordinates

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

homogeneous **scene**  
coordinates

- Converting *from* homogeneous coordinates *to* cartesian

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

$$\begin{bmatrix} x \\ y \\ z \\ w \end{bmatrix} \Rightarrow (x/w, y/w, z/w)$$

# Homogeneous Coordinates

---

- Invariant to scaling

$$k \begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} kx \\ ky \\ kw \end{bmatrix} \Rightarrow \begin{bmatrix} \frac{kx}{kw} \\ \frac{ky}{kw} \end{bmatrix} = \begin{bmatrix} \frac{x}{w} \\ \frac{y}{w} \end{bmatrix}$$

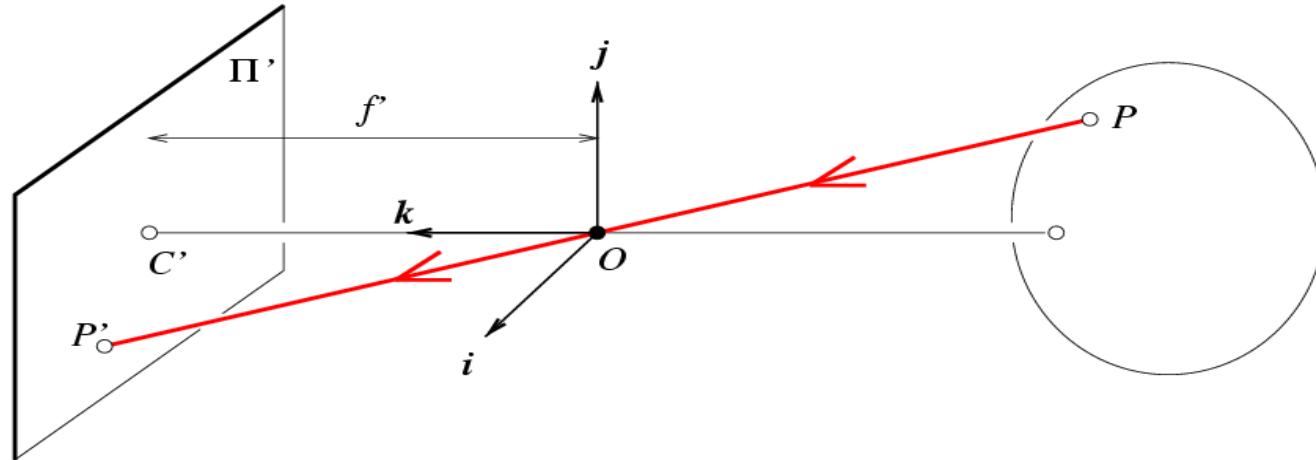
**Homogeneous  
Coordinates**

**Cartesian  
Coordinates**

- Point in **cartesian** is a ray in **homogeneous**

# Projection matrix

---



## Intrinsic Assumptions

- Unit aspect ratio
- Optical center at  $(0,0)$
- No skew

## Extrinsic Assumptions

- No rotation
- Camera at  $(0,0,0)$

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Remove assumption of known optical center

---

## Intrinsic Assumptions

- Unit aspect ratio
- No skew

## Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & u_0 \\ 0 & f & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Remove assumption: square pixels

---

## Intrinsic Assumptions

- No skew

## Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X}$$

$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Remove assumption: non-skewed pixels

---

Intrinsic Assumptions

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X}$$

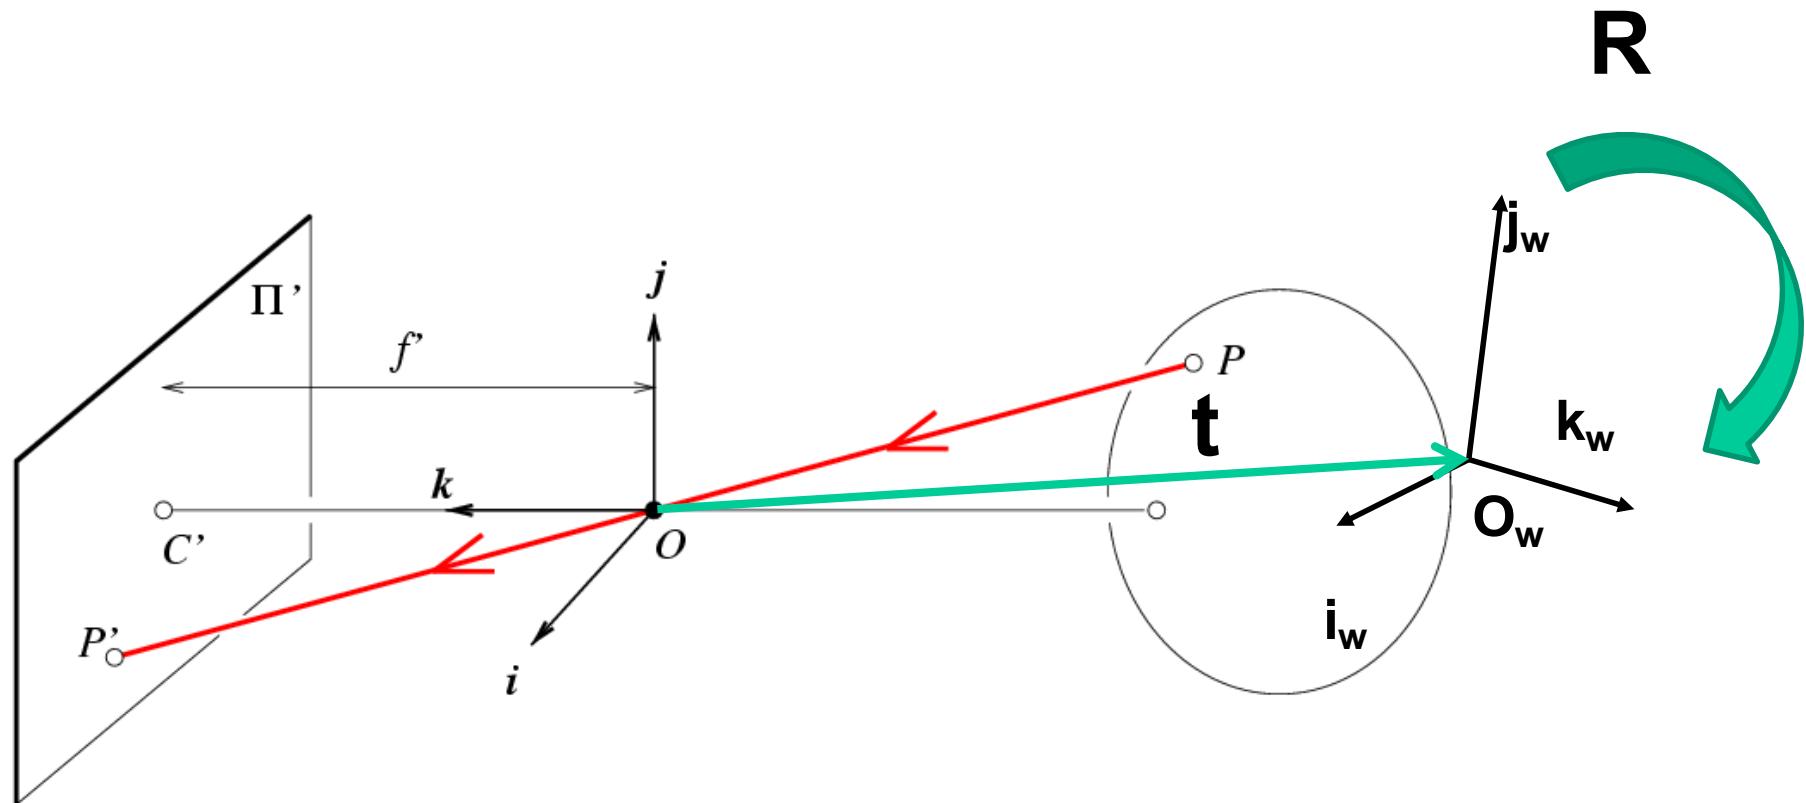
Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# Rotated and Translated Camera

---



# Allow Camera Translation

---

Intrinsic Assumptions

Extrinsic Assumptions

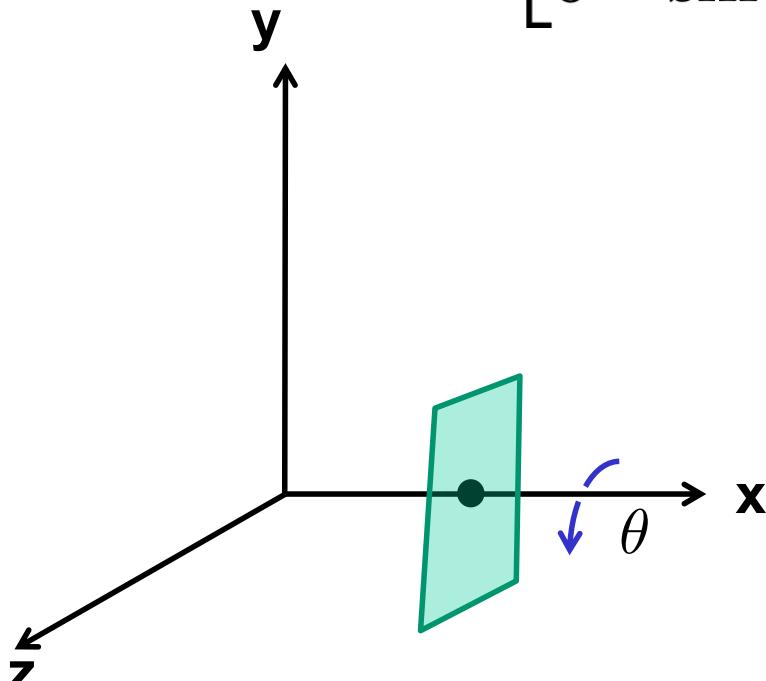
- No rotation

$$\mathbf{x} = \mathbf{K}[\mathbf{I} \quad \mathbf{t}] \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

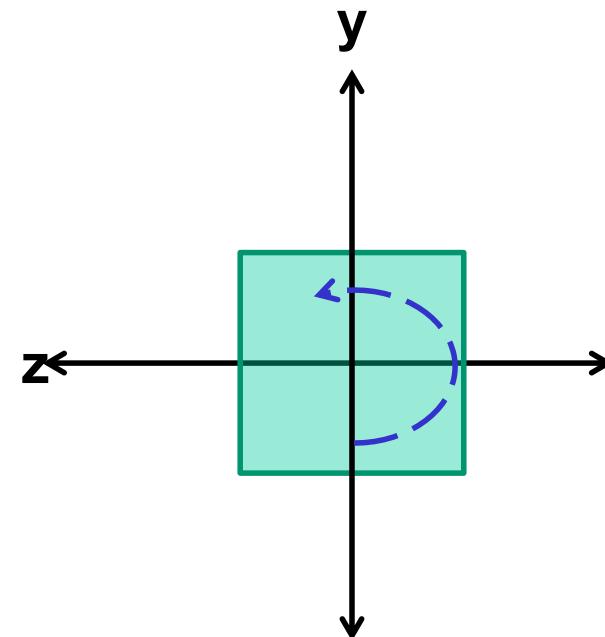
# 3D Rotation of Points

---

$$R_x(\theta) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos(\theta) & -\sin(\theta) \\ 0 & \sin(\theta) & \cos(\theta) \end{bmatrix}$$



x coordinate is unchanged after rotation around x-axis

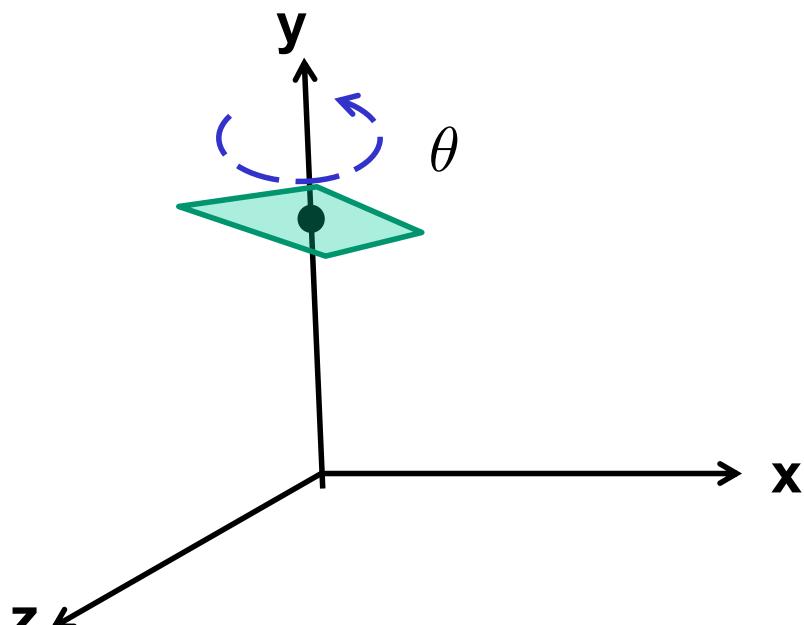


View looking down -x

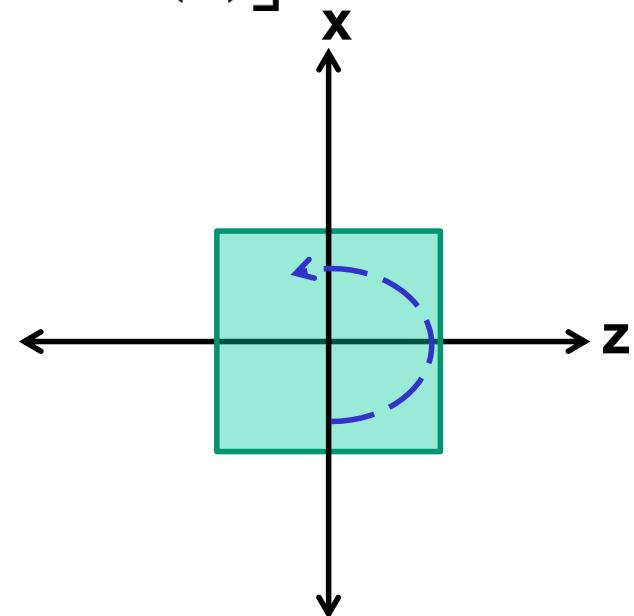
# 3D Rotation of Points

---

$$R_y(\theta) = \begin{bmatrix} \cos(\theta) & 0 & \sin(\theta) \\ 0 & 1 & 0 \\ -\sin(\theta) & 0 & \cos(\theta) \end{bmatrix}$$



y coordinate is unchanged after rotation around y-axis

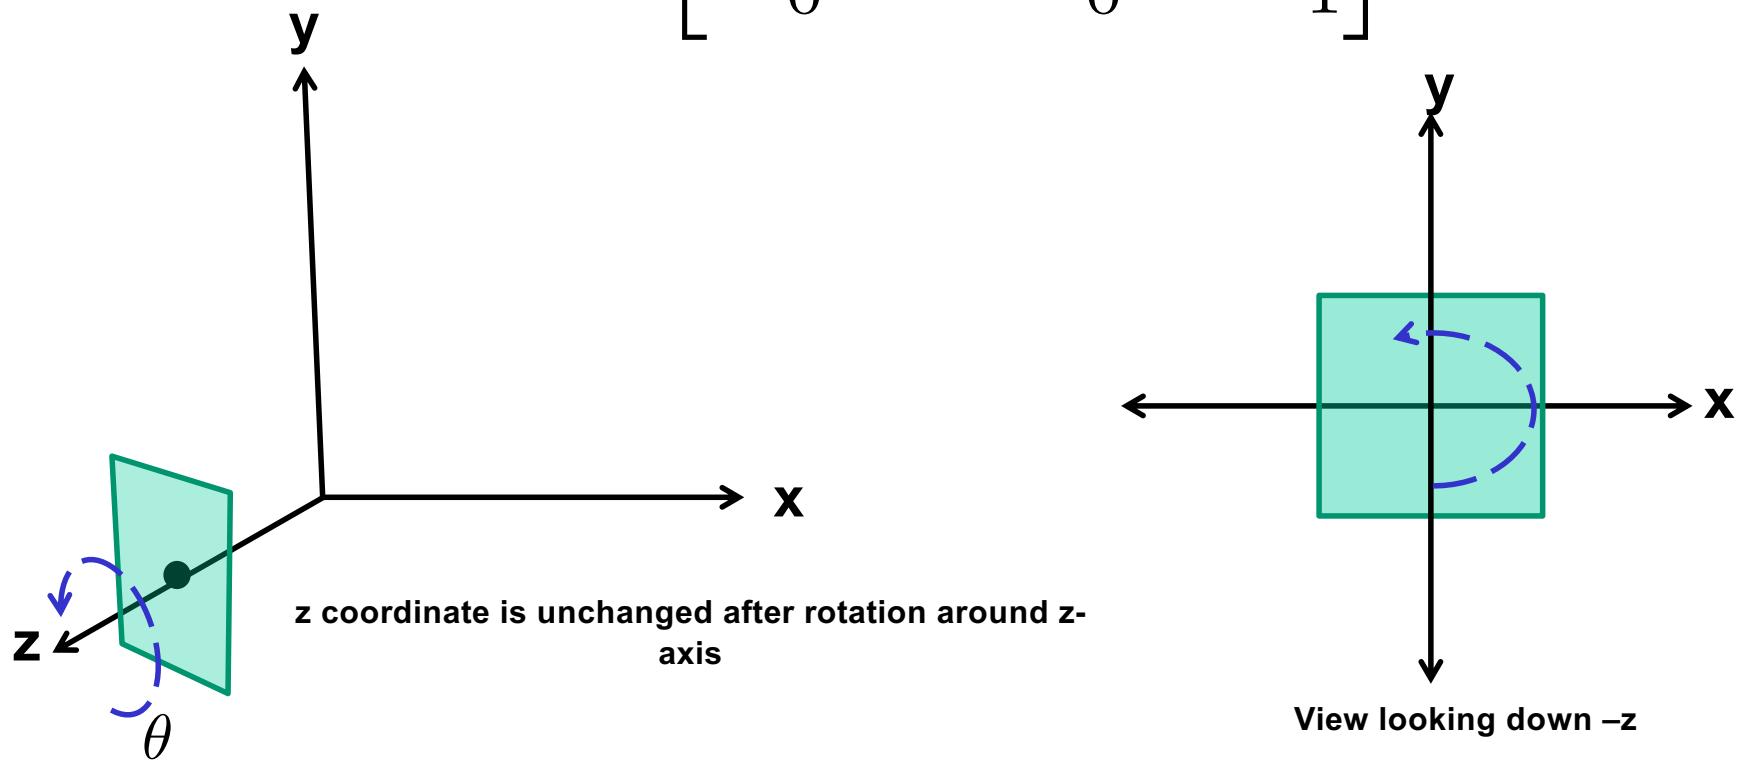


View looking down -y

# 3D Rotation of Points

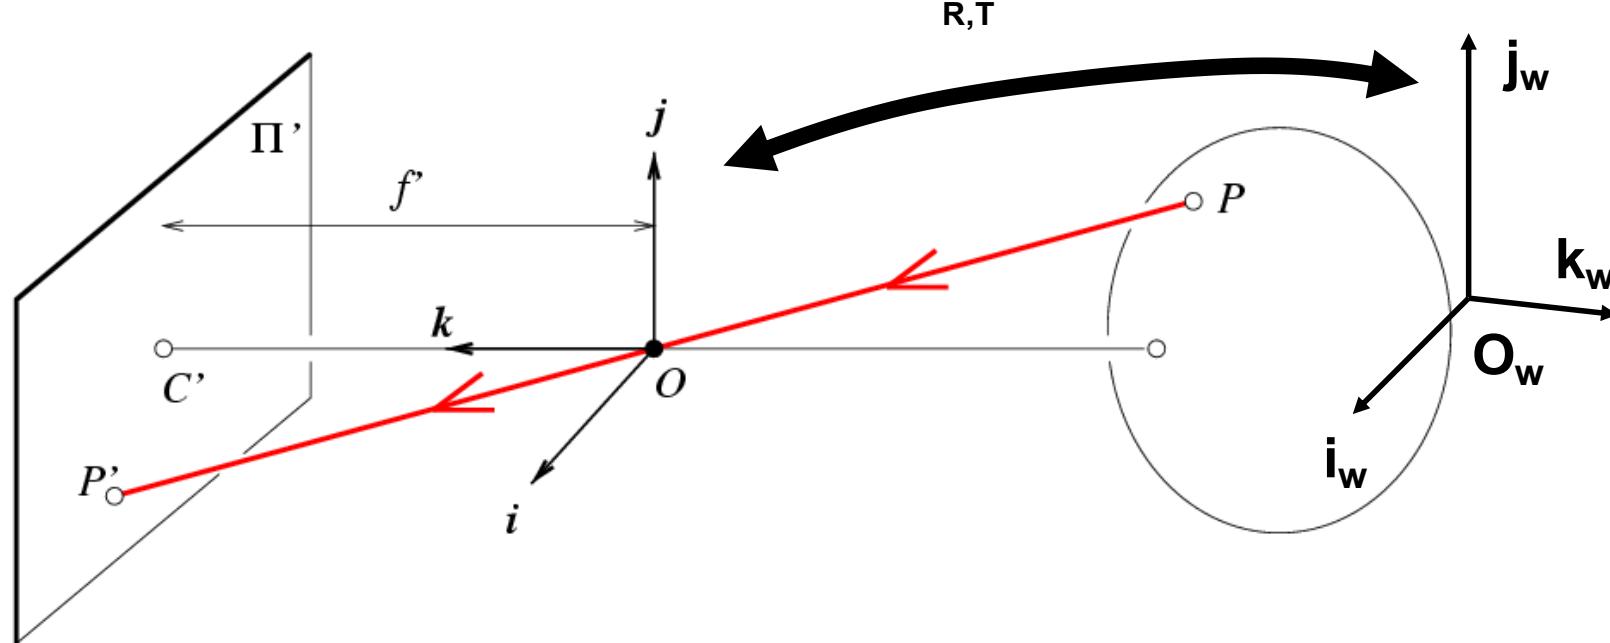
---

$$R_z(\theta) = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & 0 \\ \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



# General Camera Projection Matrix

---



$$\mathbf{x} = \mathbf{K} [\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$

$\mathbf{x}$ : Image Coordinates:  $(u, v, 1)$

$\mathbf{K}$ : Intrinsic Matrix (3x3)

$\mathbf{R}$ : Rotation (3x3)

$\mathbf{t}$ : Translation (3x1)

$\mathbf{X}$ : World Coordinates:  $(X, Y, Z, 1)$

# General Camera Projection Matrix

---

$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$



$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# General Camera Projection Matrix

Slide Credit: Saverese

$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$

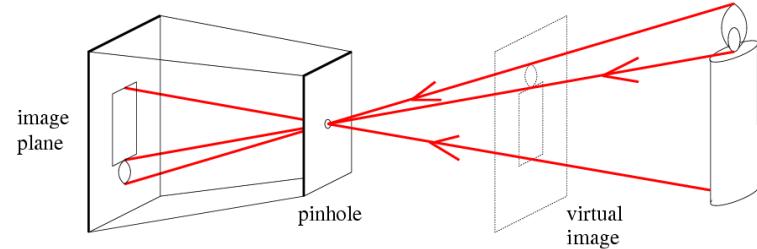


$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

# A quick summary

---

- Pinhole camera model and camera projection matrix
- Homogeneous coordinates



$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Perspective effects



3D computer vision aims to solve an inverse problem of computer graphics

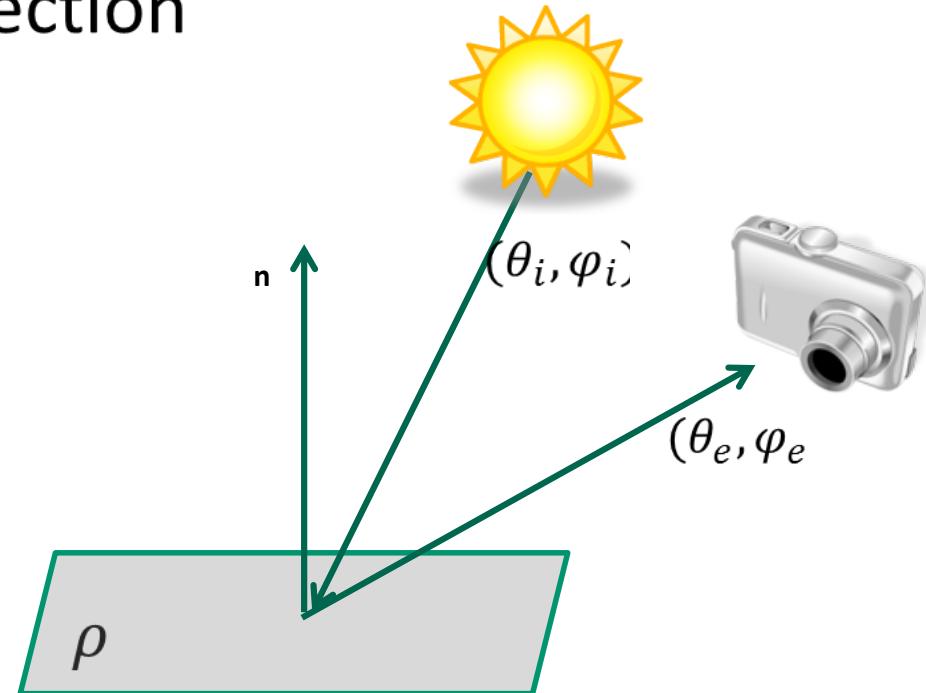
---

# Photometric (radiometric) image formation

# What determines scene radiance?

---

- The amount of light that falls on the surface
- The fraction of light that is reflected (albedo)
- Geometry of light reflection
  - Shape of surface
  - Viewpoint

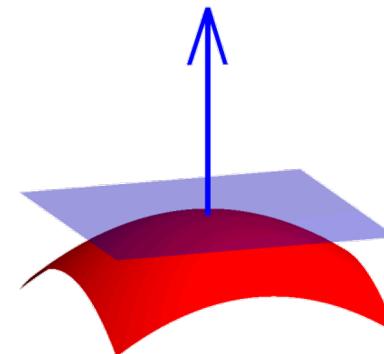


# Surface Normal

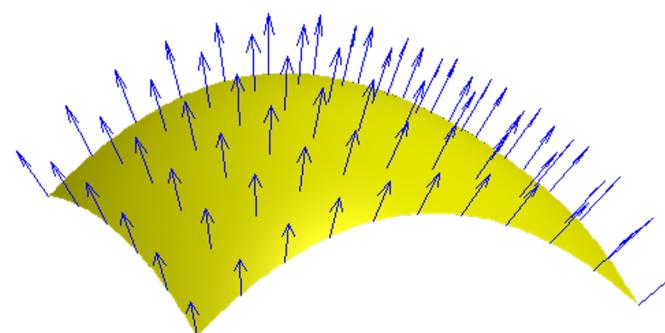
---

Convenient notation for surface orientation

A smooth surface has a tangent plane at every point



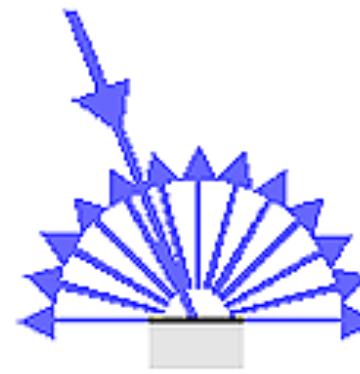
We can model the surface using the normal at every point



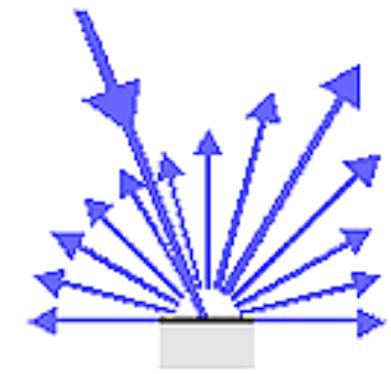
# Lambertian surface

---

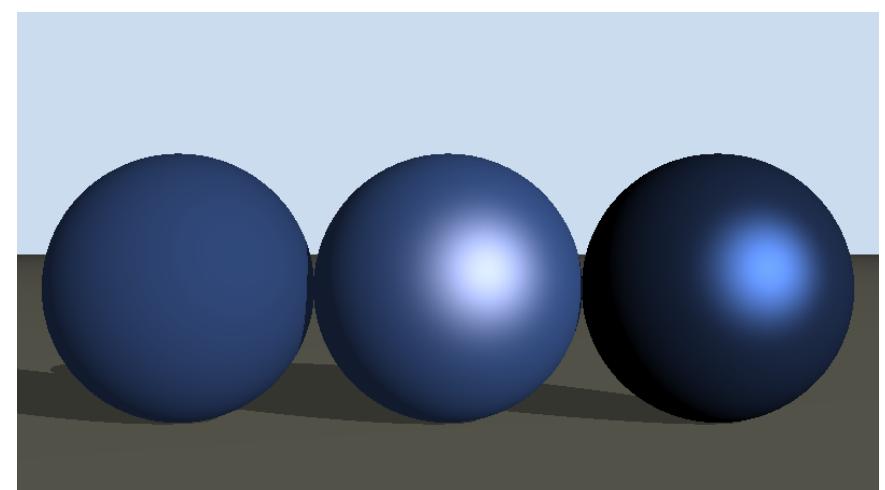
- Ideal diffuse reflectors.
- The apparent brightness of such a surface to an observer is the same regardless of the observer's angle of view.



*Ideal diffuse reflection  
(Lambertian surface)*

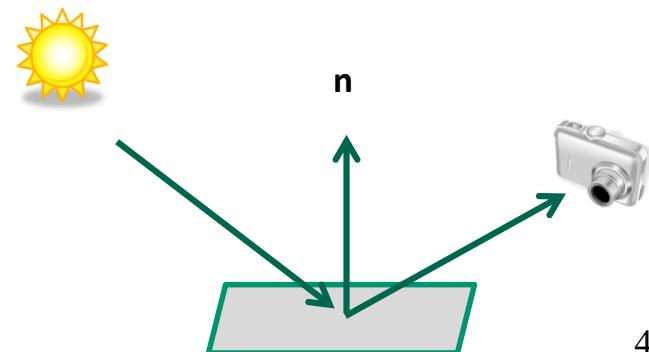


*Diffuse reflection with  
directional component*

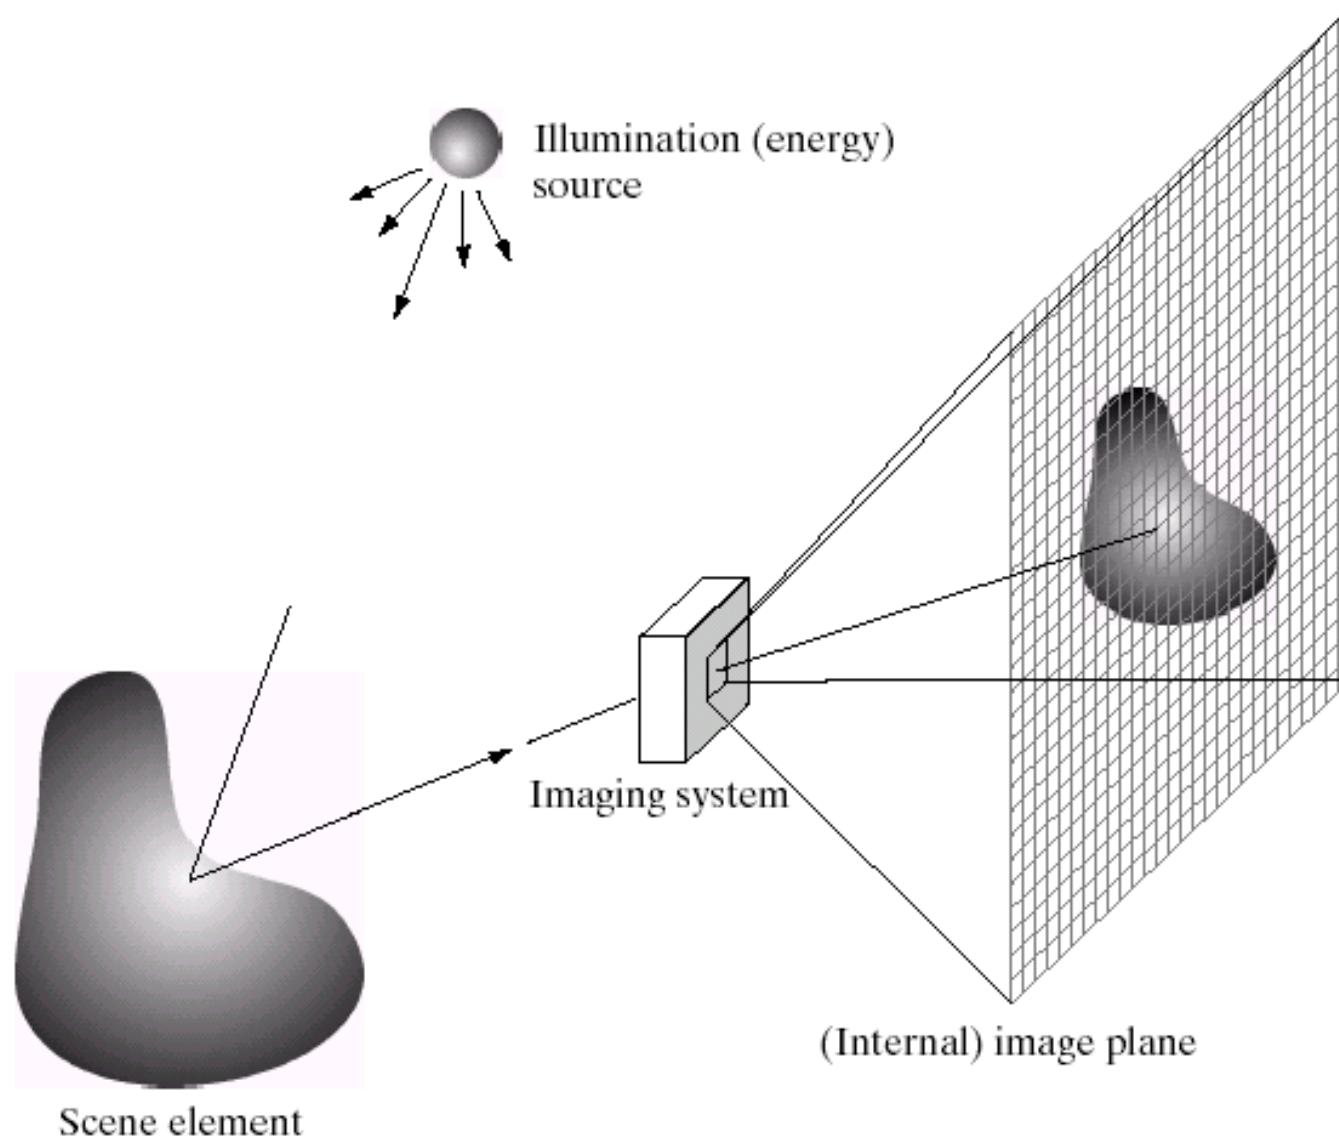


# Lambertian Surface

- Appears equally bright from all viewing directions
- Reflects all light without absorbing
- Matte surface, no “shiny” spots
- Brightness of the surface as seen from camera is linearly correlated to the amount of light falling on the surface

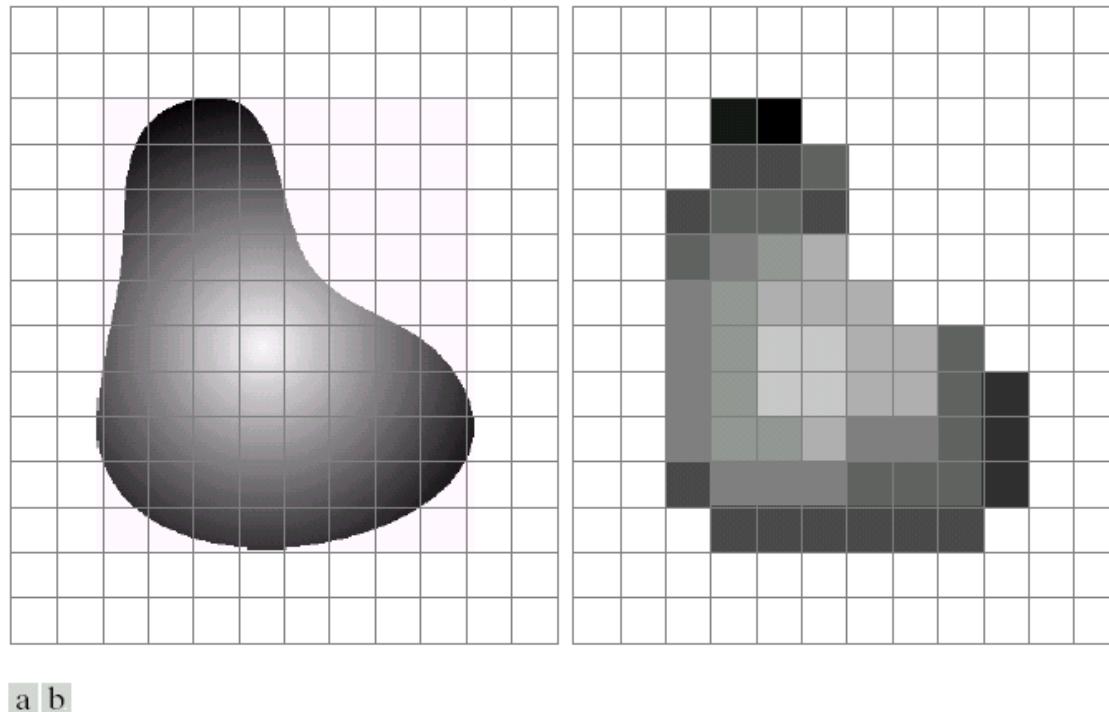


# Photometric Image Formation



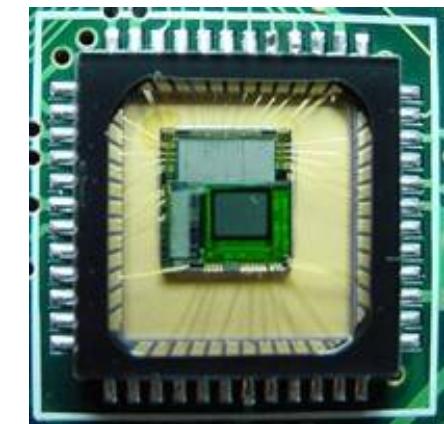
# Sensor Array

---



a b

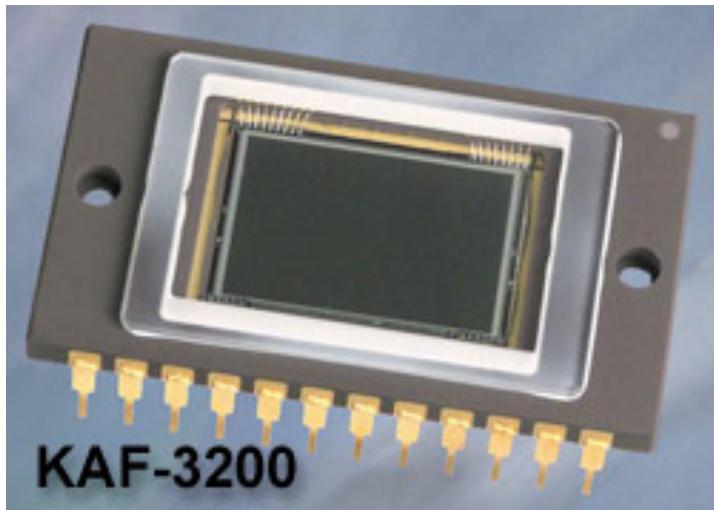
**FIGURE 2.17** (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.



**CCD/CMOS sensor**

# Image Sensors : Array Sensor

## Charge-Coupled Device (CCD)



CCD KAF-3200E from Kodak.  
(2184 x 1472 pixels,  
Pixel size 6.8 microns<sup>2</sup>)

- ◆ Used for convert a continuous image into a digital image
- ◆ Contains an array of light sensors
- ◆ Converts photon into electric charges accumulated in each sensor unit

# Human Color Perception

---

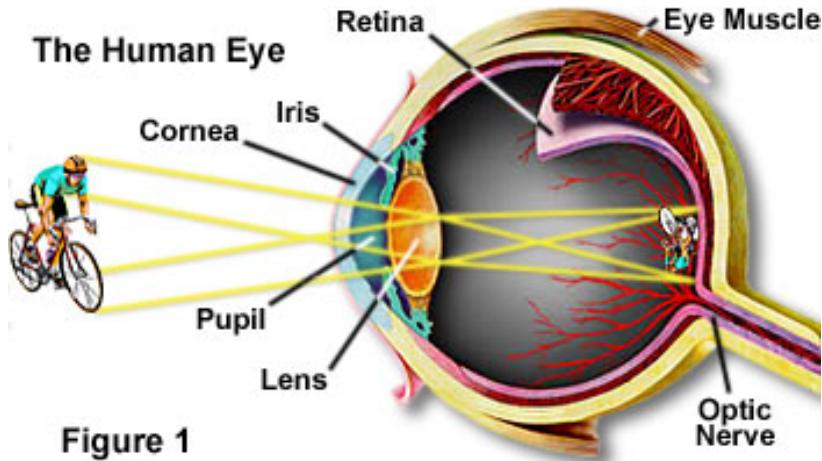
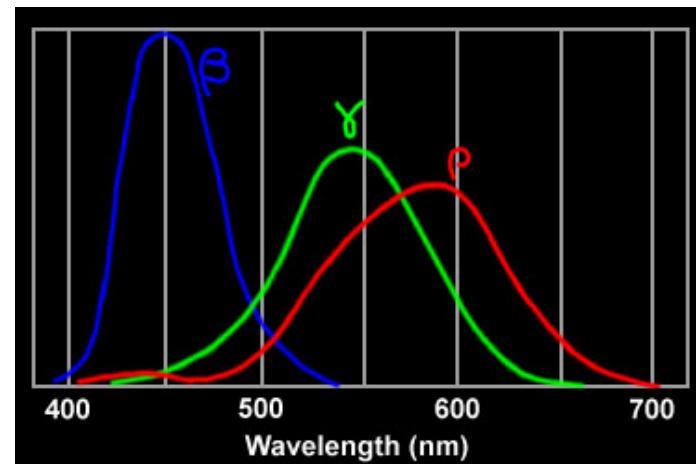
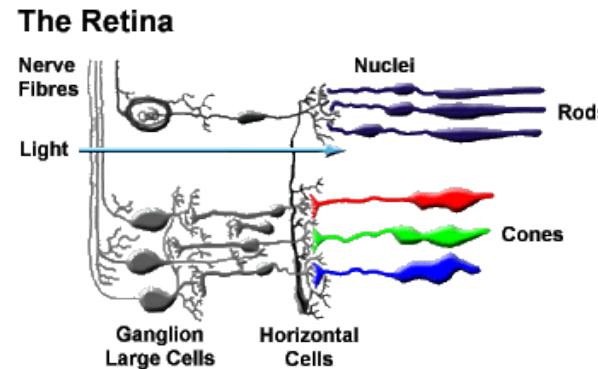


Figure 1

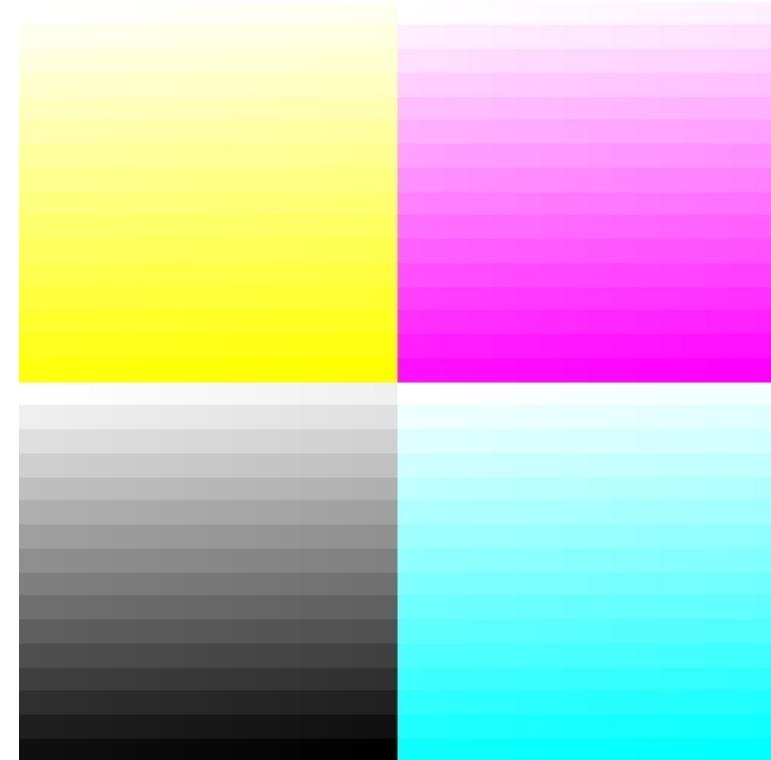


A website about human color perception:

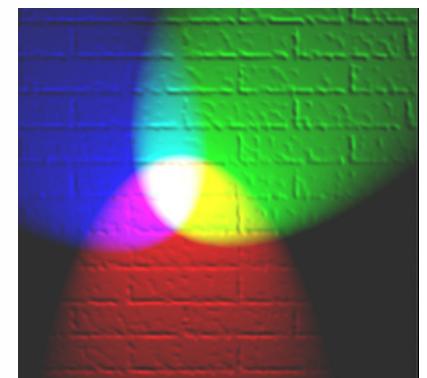
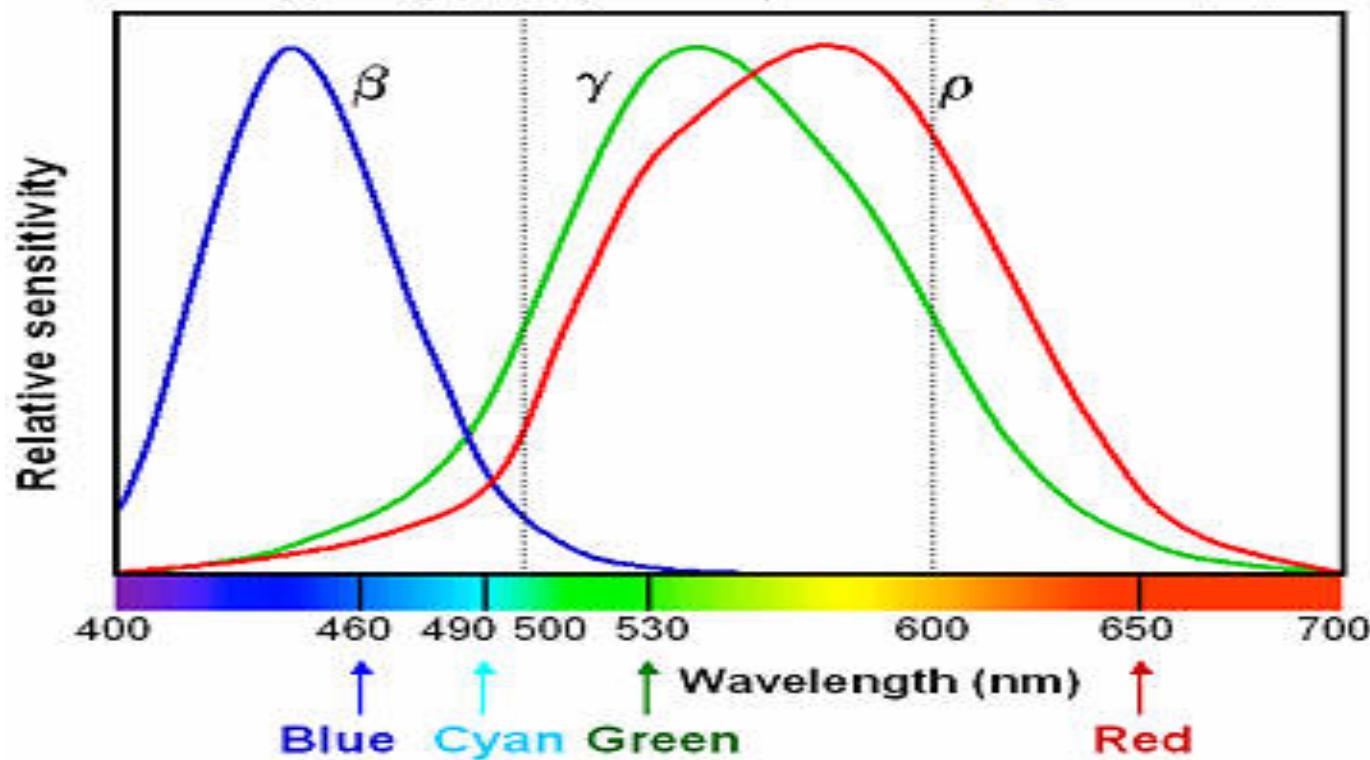
- <http://www.photo.net/photo/edscott/vis00010.htm>

# How many Colors can be Seen?

- Human eyes can distinguish about
  - 128 different hues
  - 130 different saturation levels
- We can distinguish about 380,000 colours!!



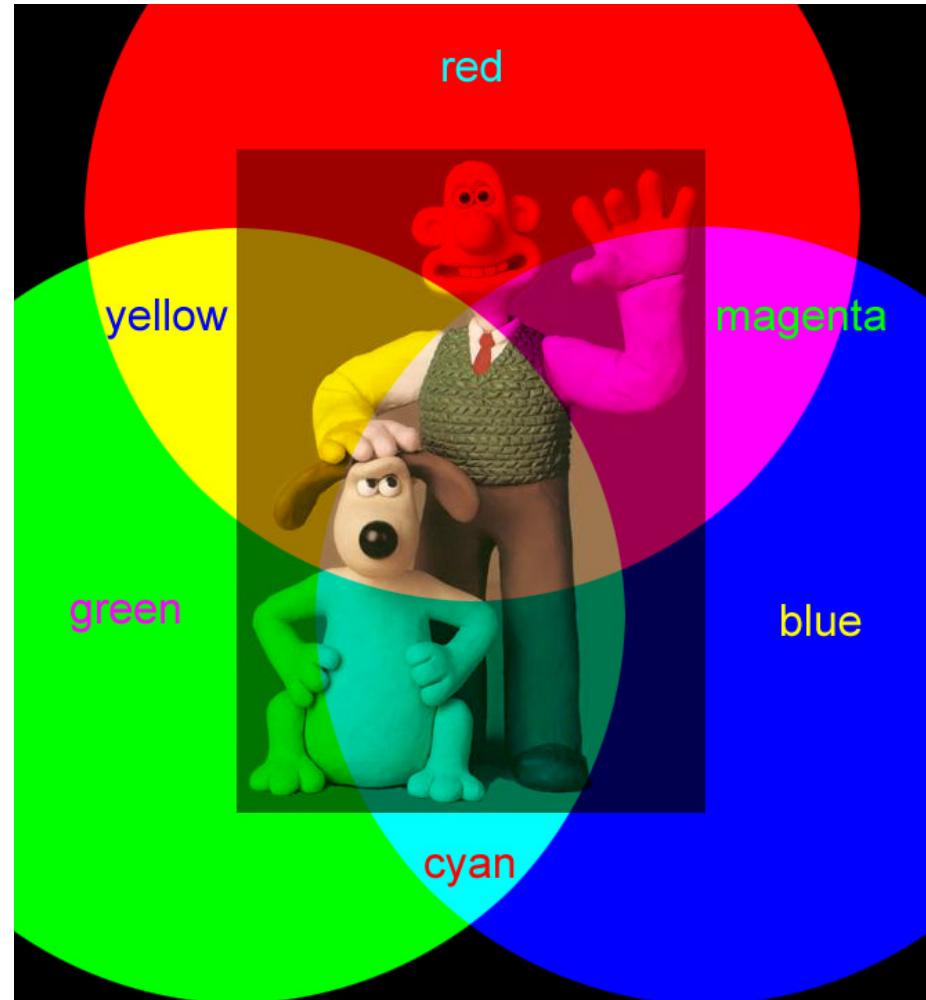
- 
- Wavelength of the light



# Colour Images

---

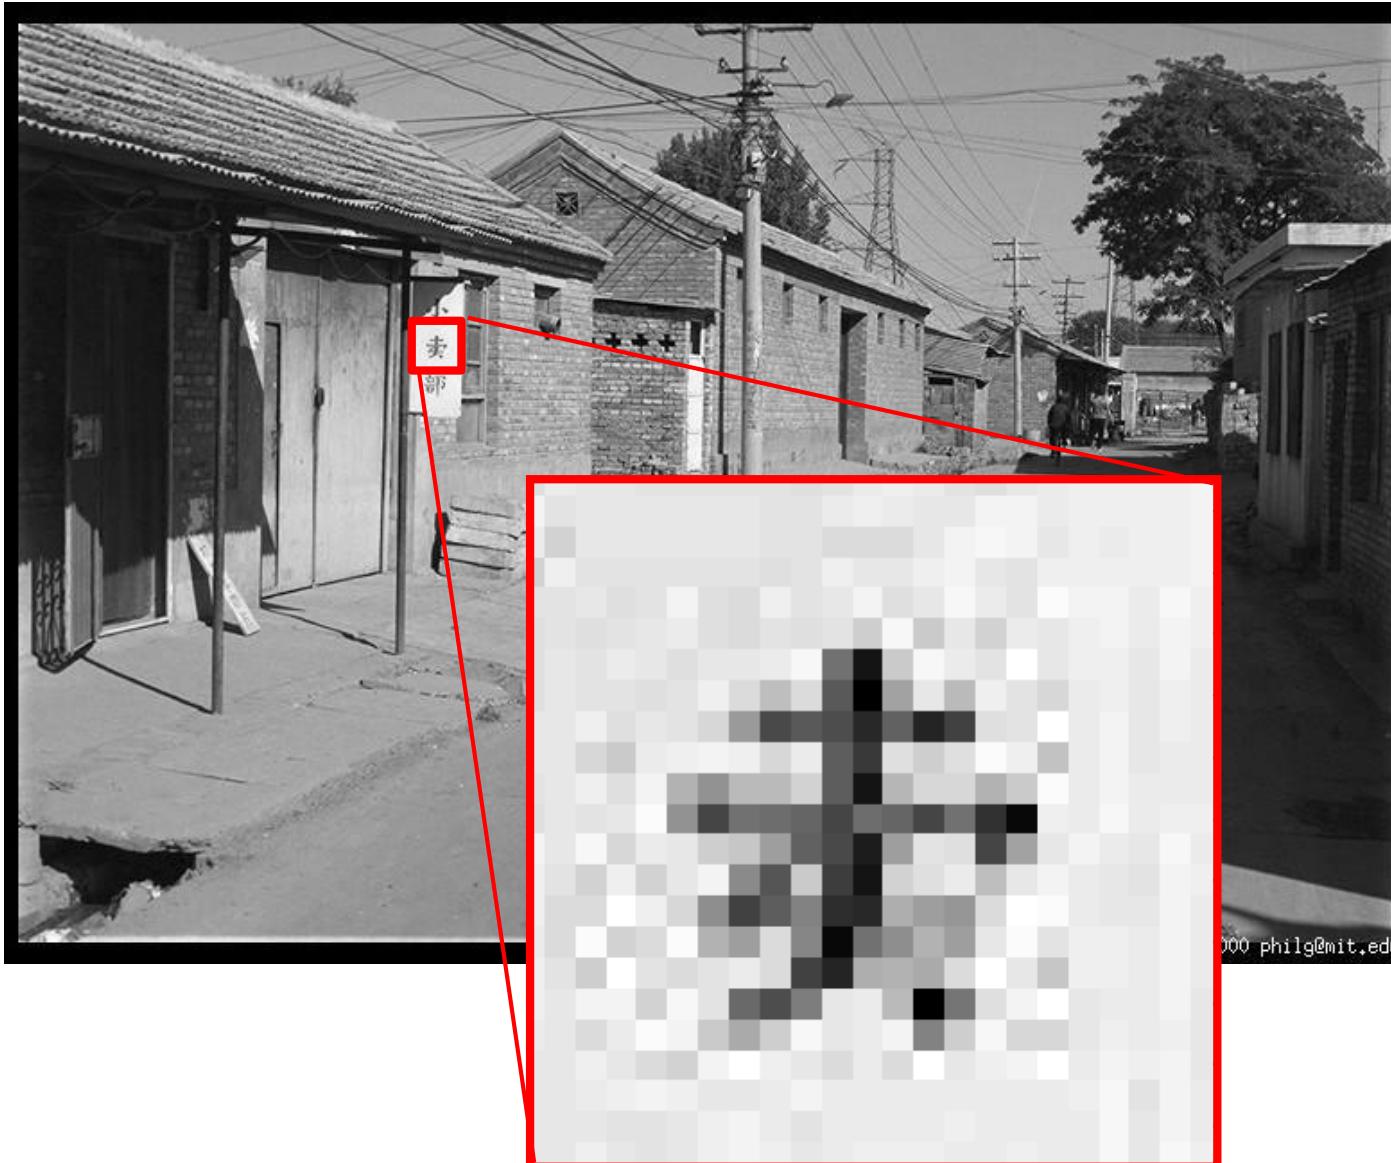
- are constructed from three intensity maps.
- Each intensity map is projected through a colour filter (e.g., red, green, or blue, or cyan, magenta, or yellow) to create a monochrome image.
- The intensity maps are overlaid to create a color image.
- Each pixel in a color image is a three dimensional vector.



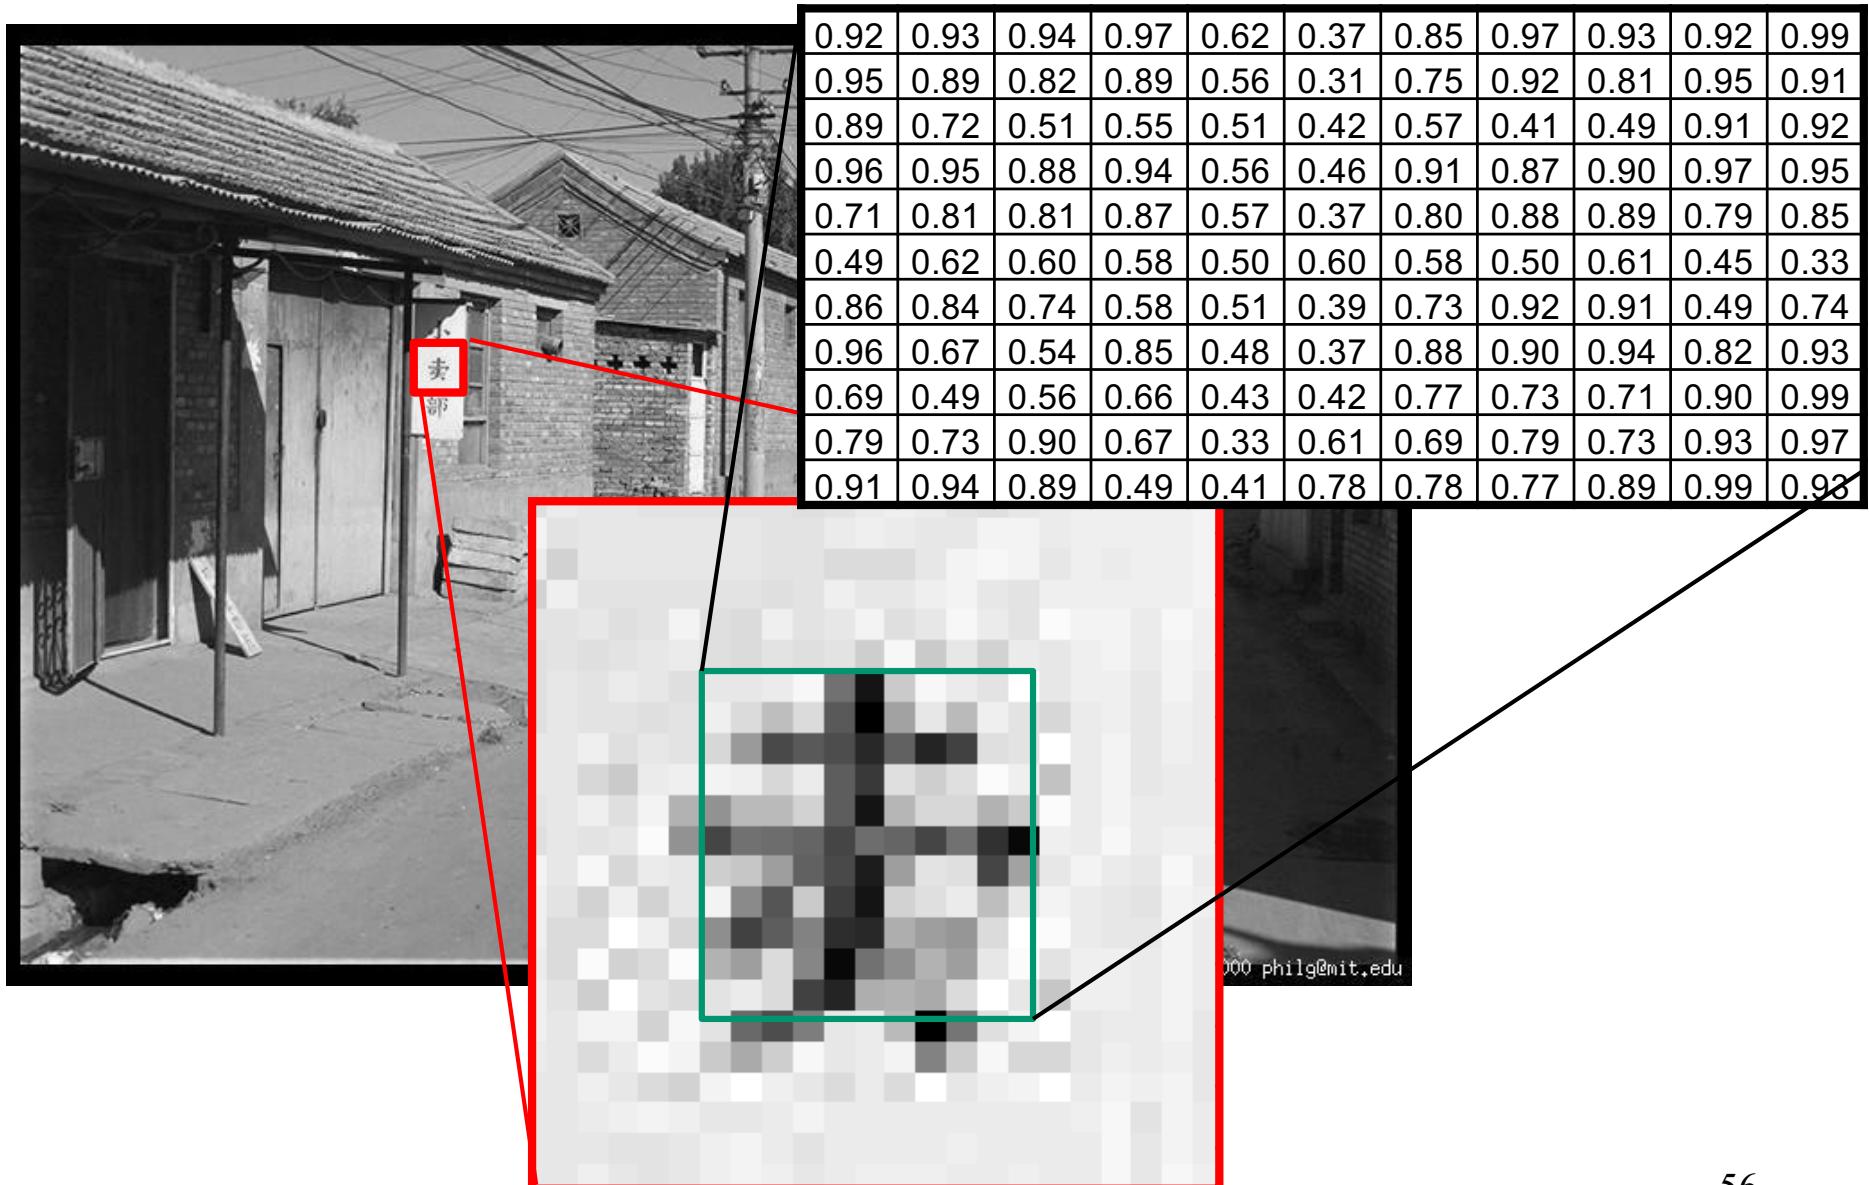
---

# Digital image representation

# The digital image (pixel matrix)



# The digital image (pixel matrix)



# Fundamentals of Digital Images

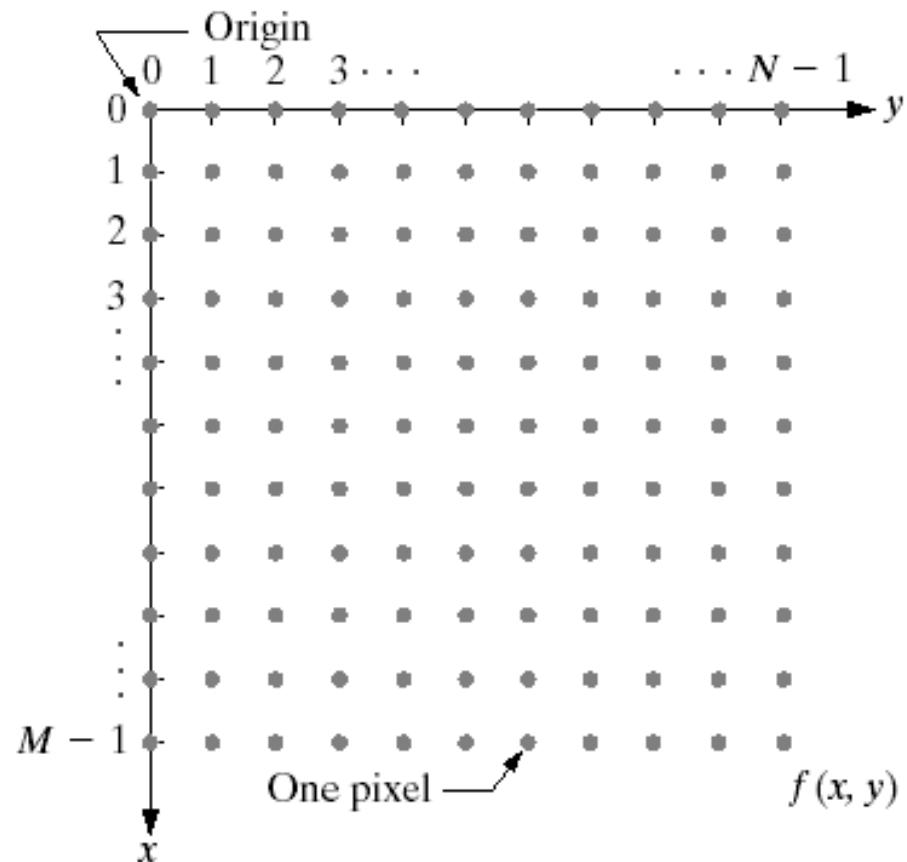
---



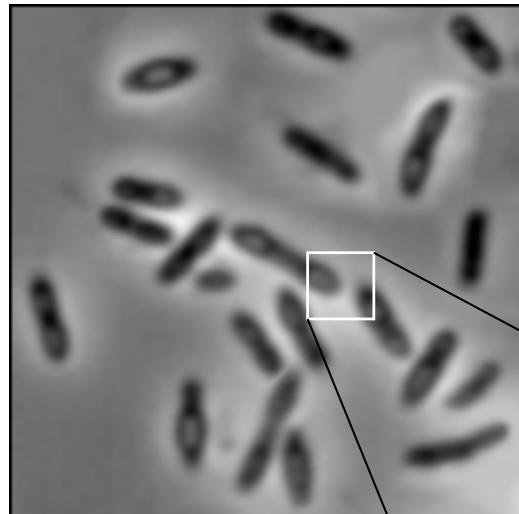
- ◆ An image = a function of spatial coordinates.
- ◆ Spatial coordinate:  $(x,y)$  for 2D case such as photograph,  
 $(x,y,z)$  for 3D case such as CT scan images  
 $(x,y,t)$  for video.
- ◆ The function  $f$  may represent intensity (for greyscale images)  
or color (for color images) or other associated values.

# Conventional Coordinate for Image Representa

---



# Grey-scale Image

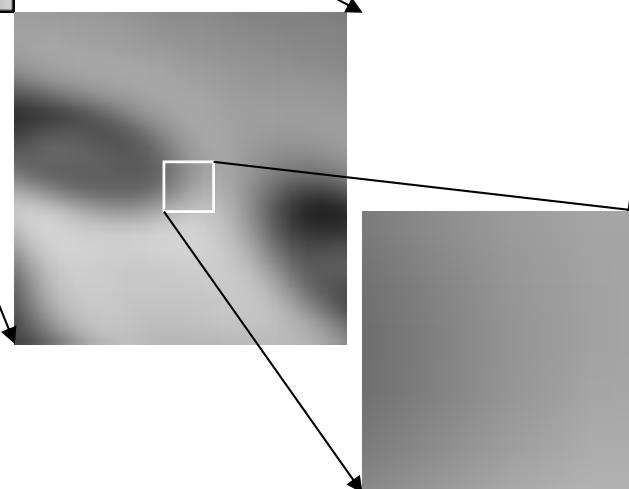


Intensity image or grey-scale image

Each pixel corresponds to light intensity

normally represented in gray scale (gray

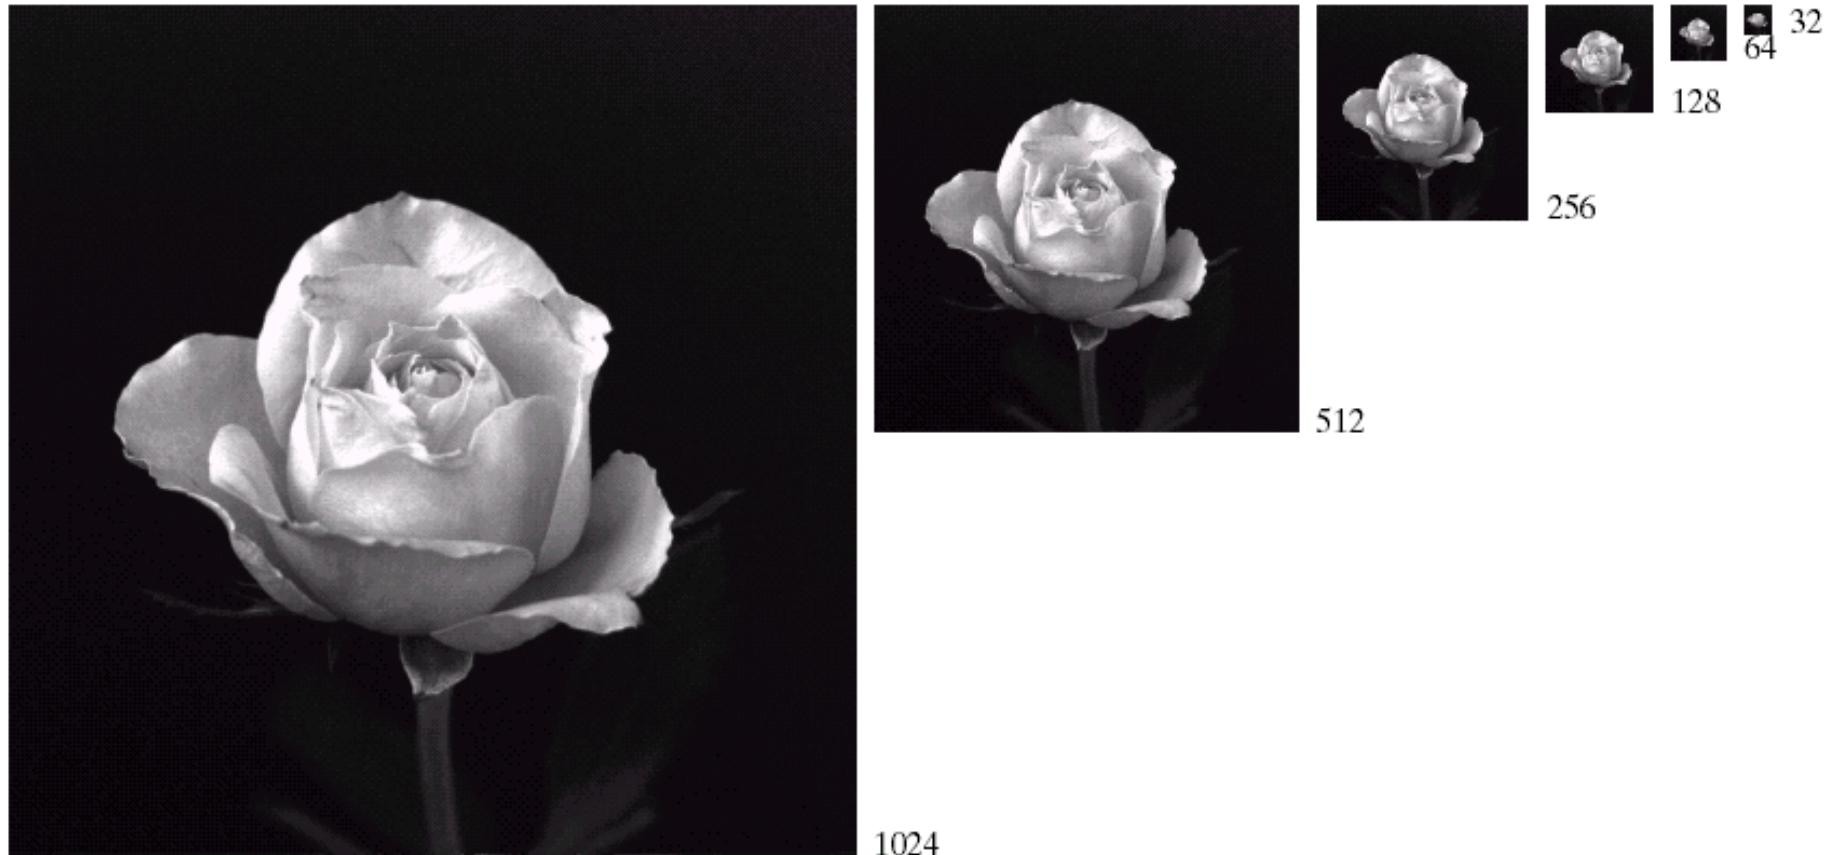
level).



Gray scale values

$$\begin{bmatrix} 10 & 10 & 16 & 28 \\ 9 & 6 & 26 & 37 \\ 15 & 25 & 13 & 22 \\ 32 & 15 & 87 & 39 \end{bmatrix}$$

# Effect of Spatial Resolution

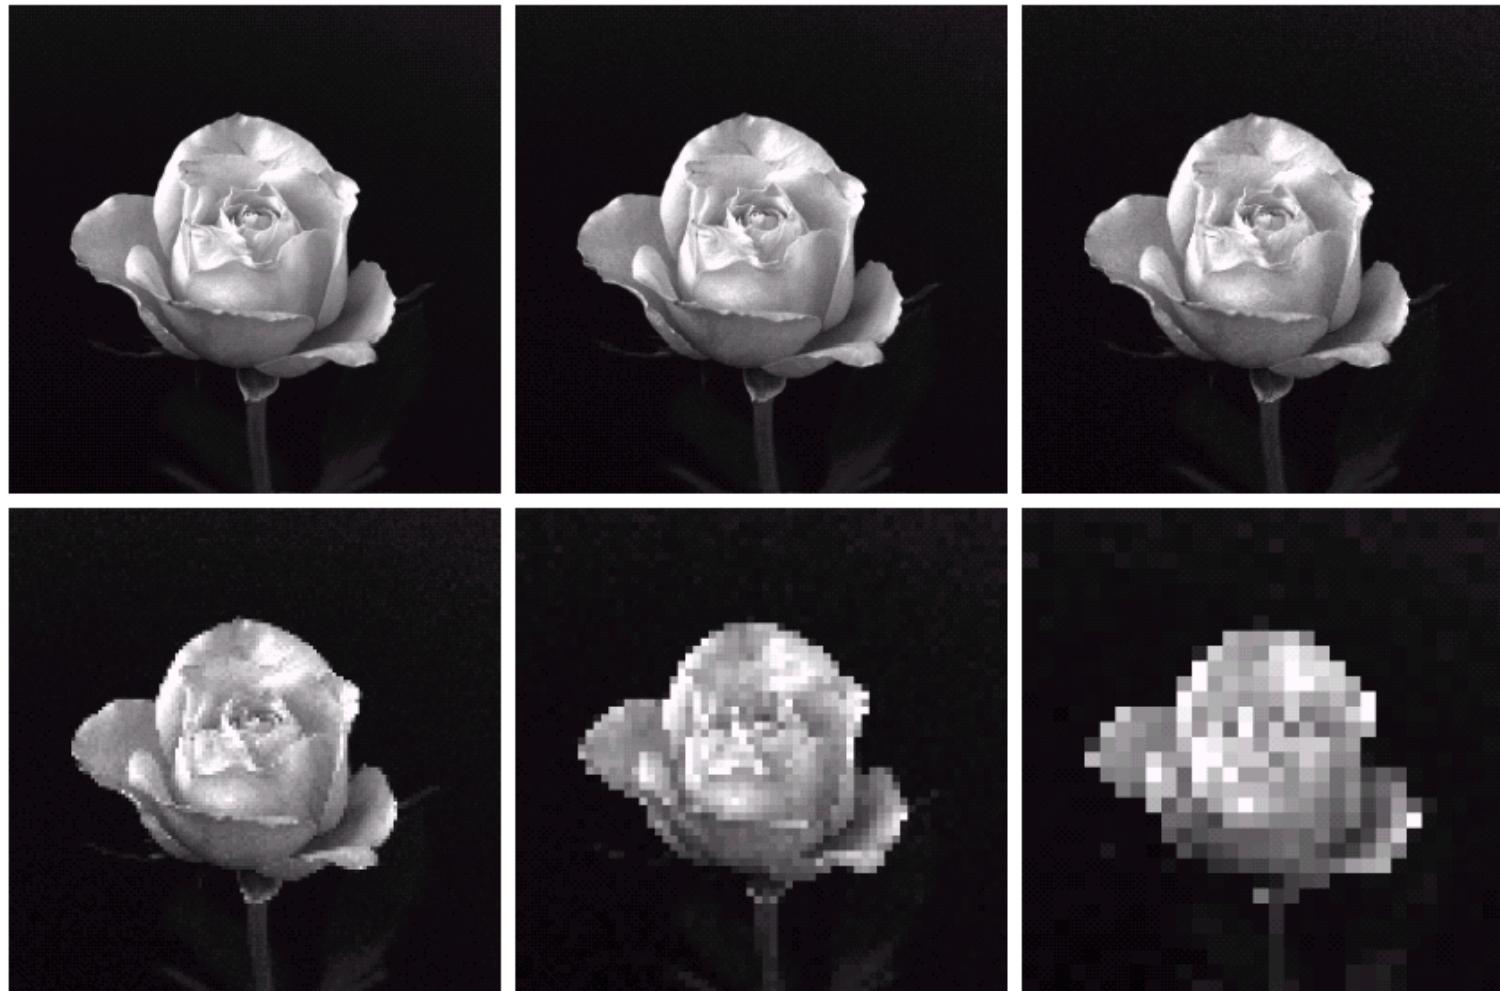


**FIGURE 2.19** A  $1024 \times 1024$ , 8-bit image subsampled down to size  $32 \times 32$  pixels. The number of allowable gray levels was kept at 256.

(Images from Rafael C. Gonzalez and Richard E. Wood, Digital Image Processing, 2<sup>nd</sup> Edition.)

# Effect of Spatial Resolution

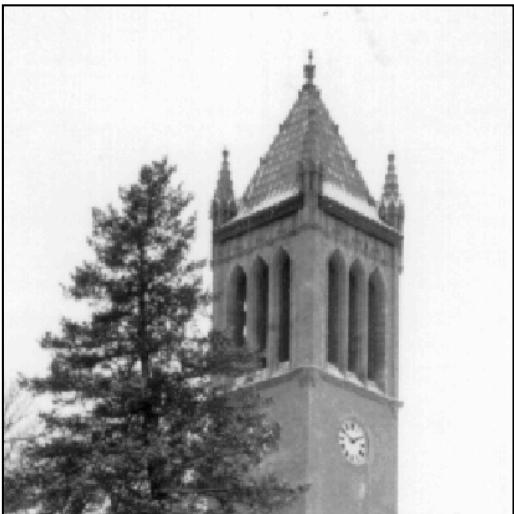
---



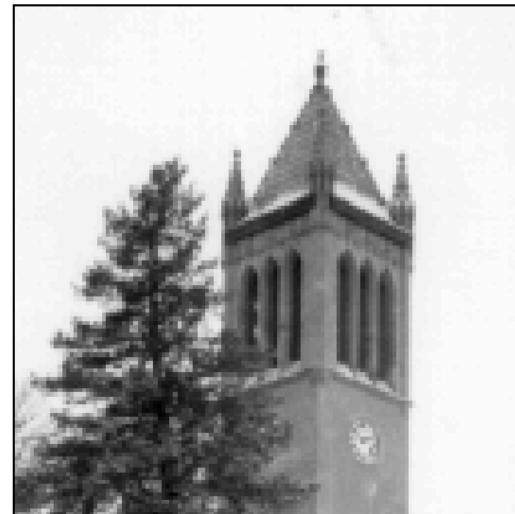
a	b	c
d	e	f

**FIGURE 2.20** (a)  $1024 \times 1024$ , 8-bit image. (b)  $512 \times 512$  image resampled into  $1024 \times 1024$  pixels by row and column duplication. (c) through (f)  $256 \times 256$ ,  $128 \times 128$ ,  $64 \times 64$ , and  $32 \times 32$  images resampled into  $1024 \times 1024$  pixels.

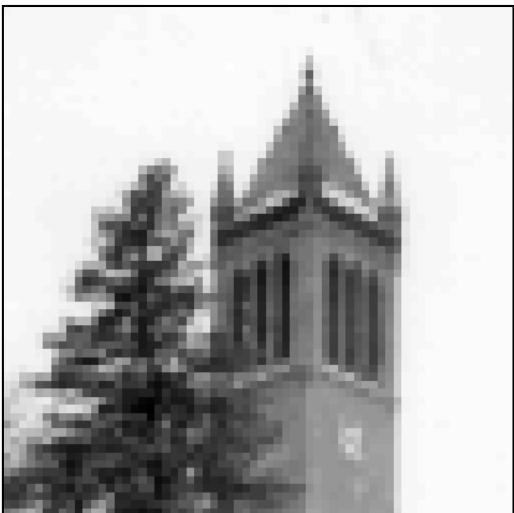
# Effect of Spatial Resolution



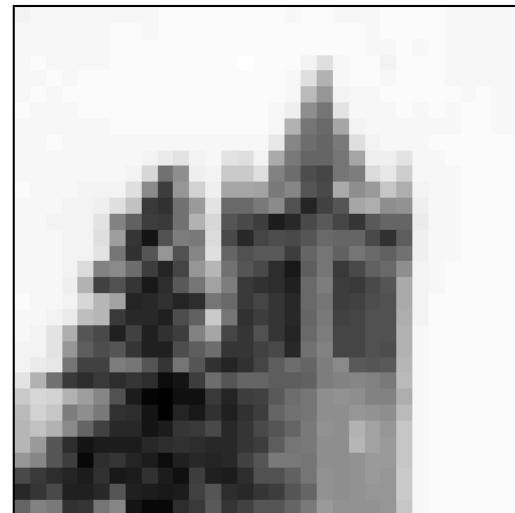
256x256 pixels



128x128 pixels



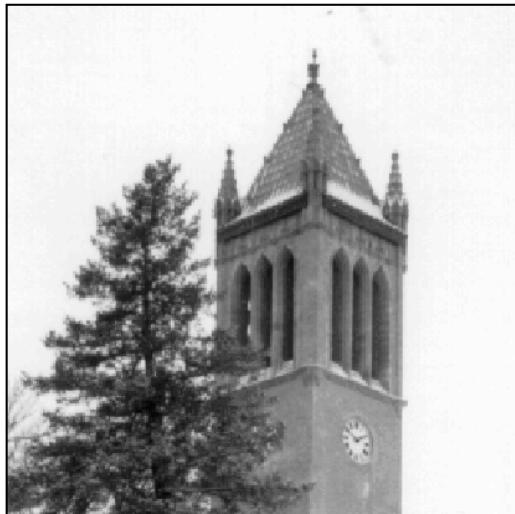
64x64 pixels



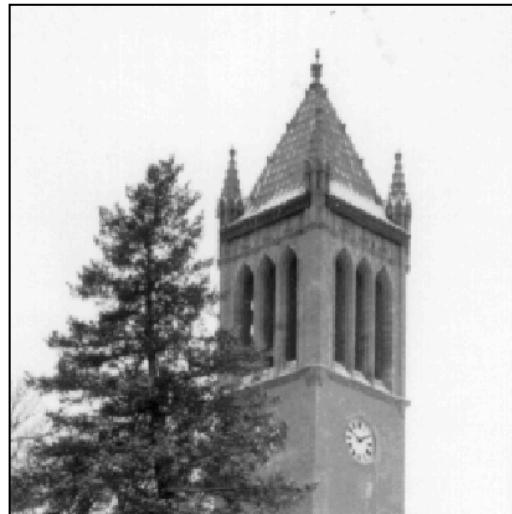
32x32 pixels

# Effect of Quantization Levels

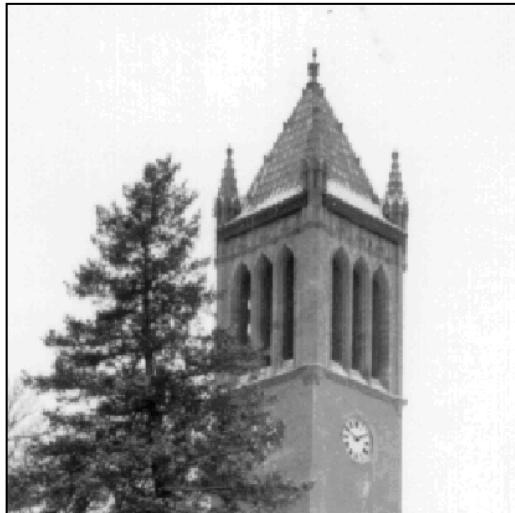
---



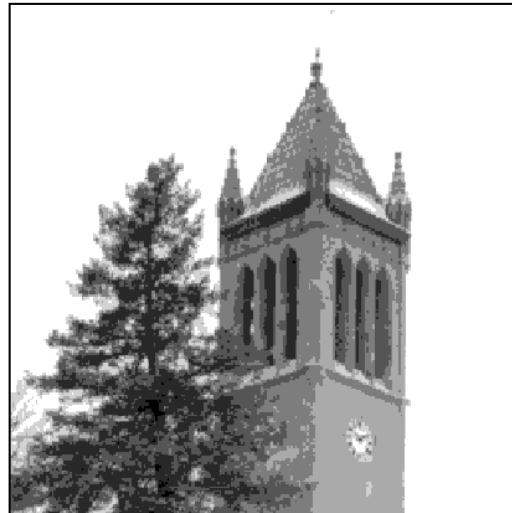
256 levels



128 levels



64 levels

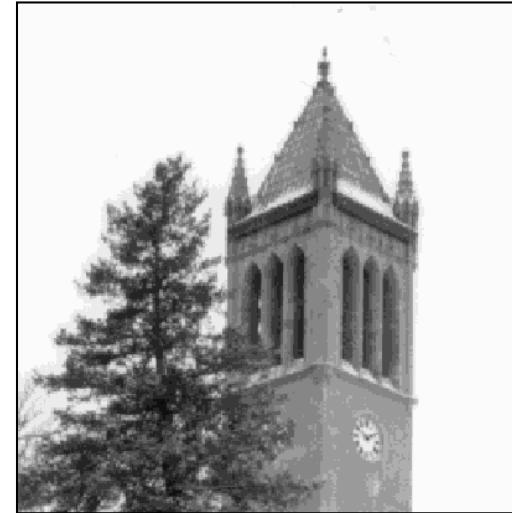


32 levels

# Effect of Quantization Levels (cont.)

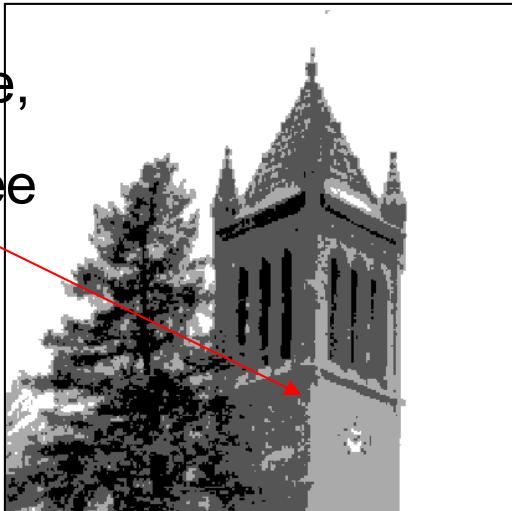


16 levels

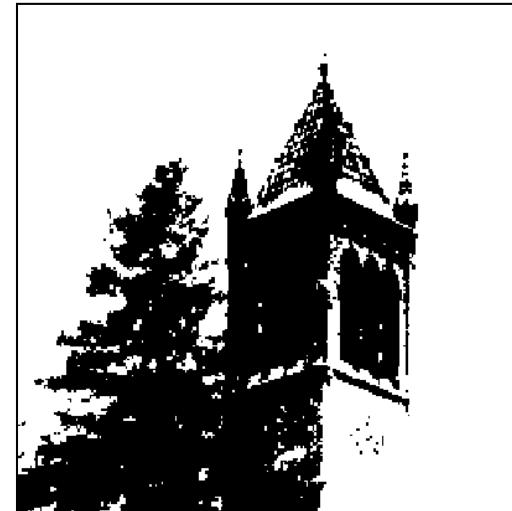


8 levels

In this image,  
it is easy to see  
false contour.



4 levels

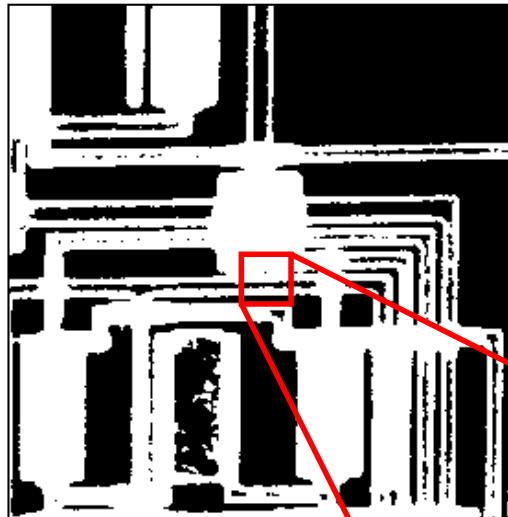


2 levels

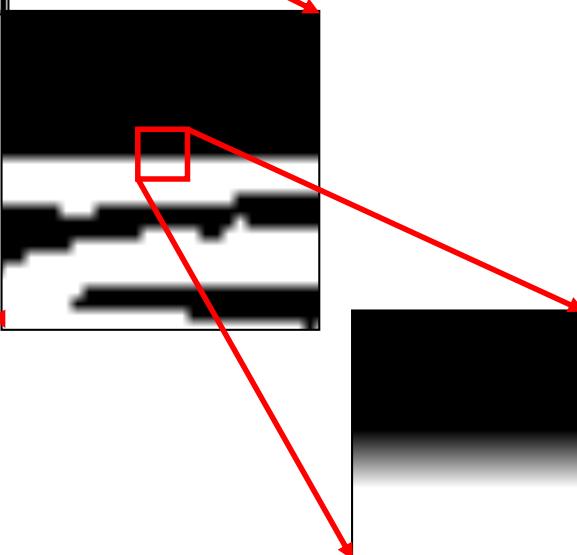
# Two-level image: binary image

---

Binary image or black and white image



- Each pixel contains one bit :
  - 1 represent white
  - 0 represents black



Binary data

$$\begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix}$$

---

# Colour Images

# A colourful image

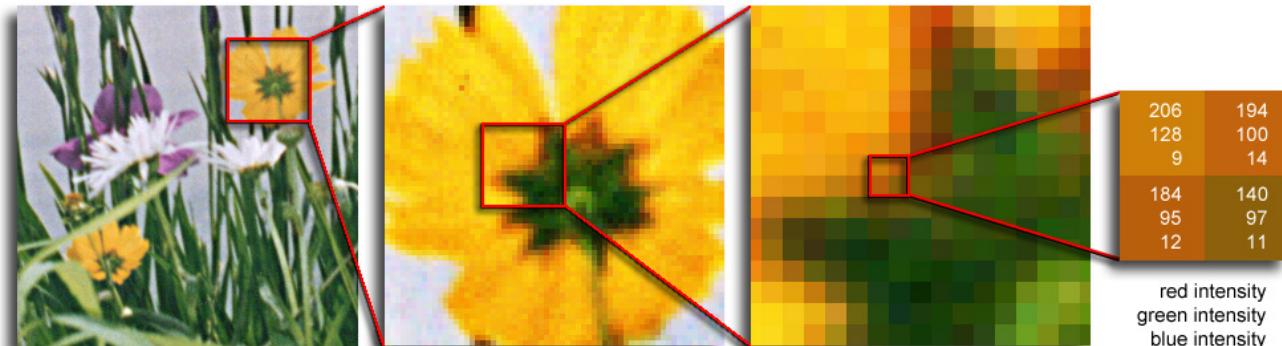
---



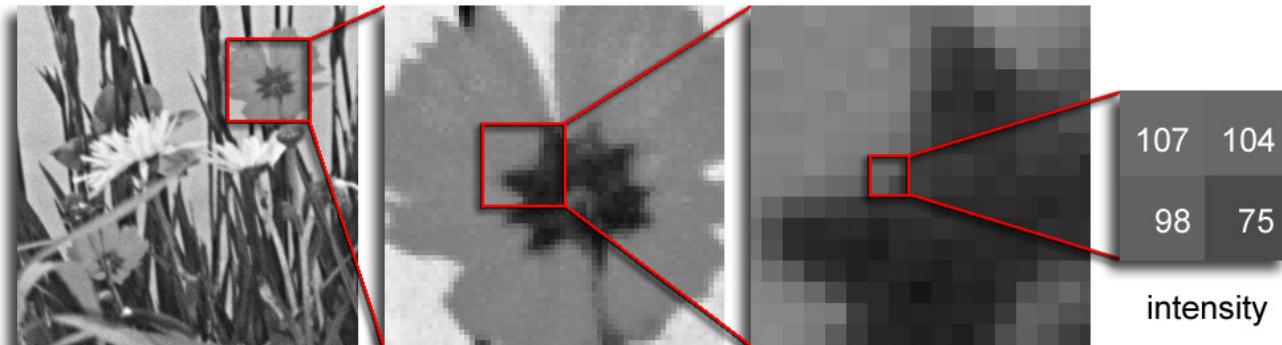
# Colour Image

- Colour images have 3 values per pixel;
- Greyscale images have 1 value per pixel.

a grid of squares, each of which contains a single colour



each square is called a pixel  
(for *picture element*)



# Colour Image

---



**Color (RGB) image:**

- each pixel contains a vector representing red, green and blue components.

RGB components

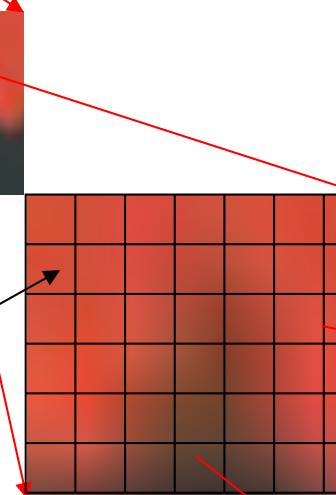
$$\begin{bmatrix} 10 & 10 & 16 & 28 \\ 65 & 70 & 56 & 43 \\ 99 & 70 & 56 & 78 \\ 32 & 60 & 90 & 96 \\ 15 & 21 & 67 & 67 \\ 32 & 85 & 85 & 43 \\ 54 & 85 & 92 & 92 \\ 32 & 32 & 65 & 87 \\ & & & 99 \end{bmatrix}$$

# Colour Image



- Each component in the image called pixel associates with the pixel value (a single number in the case of intensity images or a vector in the case of color images)

- Colour image = a multidimensional array of numbers (such as intensity image) or vectors (such as color image)

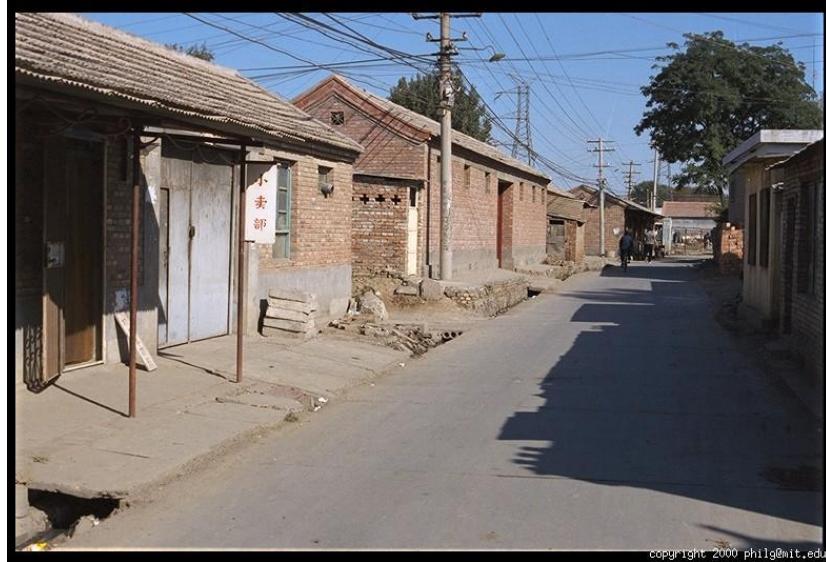


10	10	16	28		
9	65	70	56	43	
15	32	99	70	56	78
32	21	60	90	96	67
54	85	85	43	92	
32	65	87	99		

# Colour Image

---

R



G



B



copyright 2000 philipmit.edu

# Colour Images in Matlab

Images represented as a matrix

Suppose we have a NxM RGB image called “im”

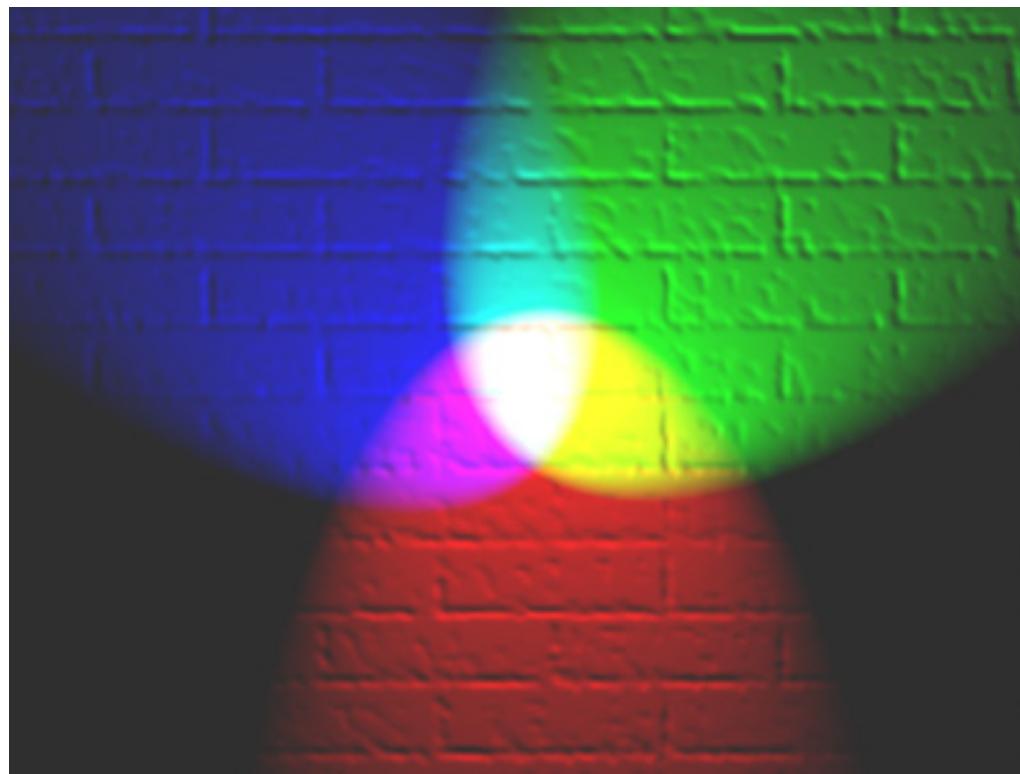
- $\text{im}(1,1,1)$  = top-left pixel value in R-channel
- $\text{im}(y, x, b)$  = y pixels down, x pixels to right in the b<sup>th</sup> channel

`imread(filename)` returns a uint8 image (values 0 to 255)

- **Convert to double format (values 0 to 1) with `im2double()`. (important ! )**

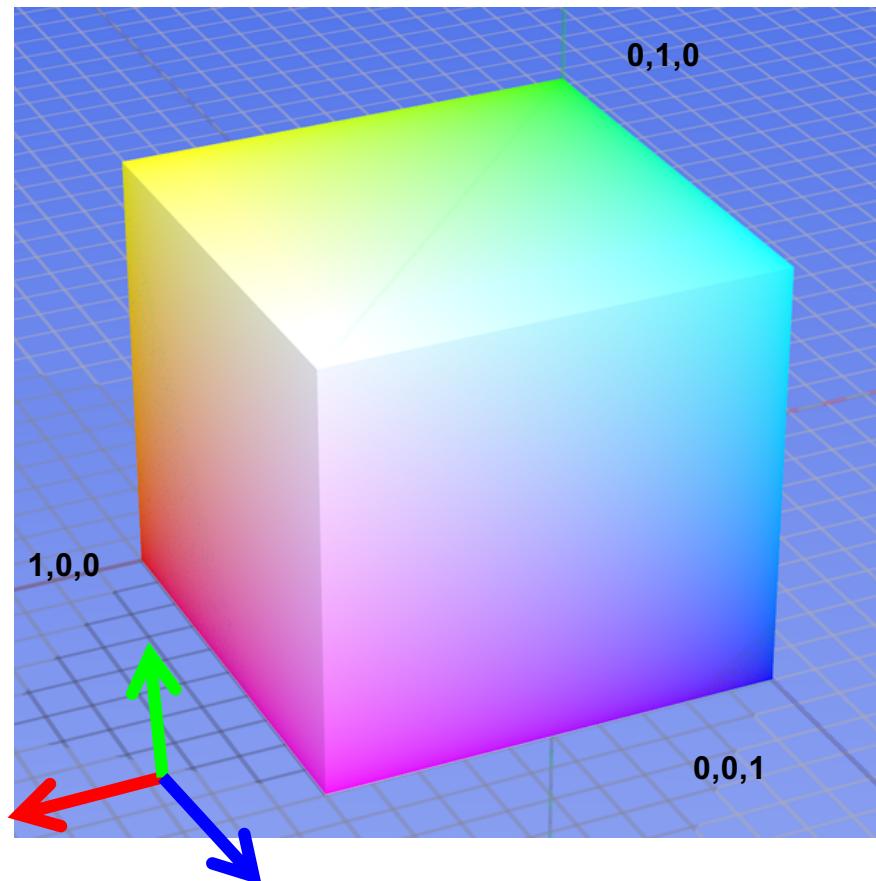
		column														
														R	G	B
row																
		0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99				
		0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91				
		0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92	0.92	0.92	0.99	
		0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95	0.95	0.91	0.91	0.95
		0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85	0.91	0.92	0.92	0.99
		0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33	0.97	0.95	0.95	0.91
		0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74	0.79	0.85	0.91	0.92
		0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93	0.45	0.33	0.97	0.95
		0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99	0.49	0.74	0.79	0.85
		0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	0.82	0.93	0.45	0.33
		0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	0.90	0.99	0.49	0.74
													0.79	0.73	0.90	0.67
													0.33	0.61	0.69	0.79
													0.73	0.93	0.97	0.82
													0.99	0.93	0.90	0.99
													0.91	0.94	0.89	0.49
													0.41	0.78	0.78	0.77
													0.78	0.77	0.89	0.99
													0.99	0.93	0.99	0.93
													0.91	0.94	0.89	0.49

# Color spaces



# RGB space

Default color space



Some drawbacks

- Strongly correlated channels
- Not perceptually meaningful.



R  
(G=0,B=0)



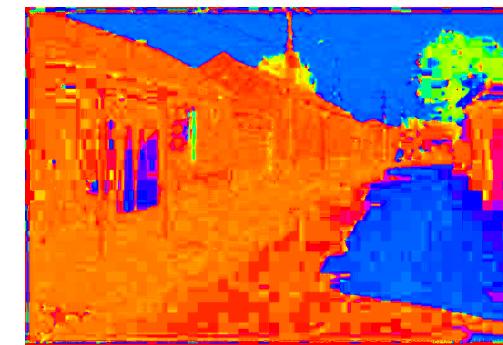
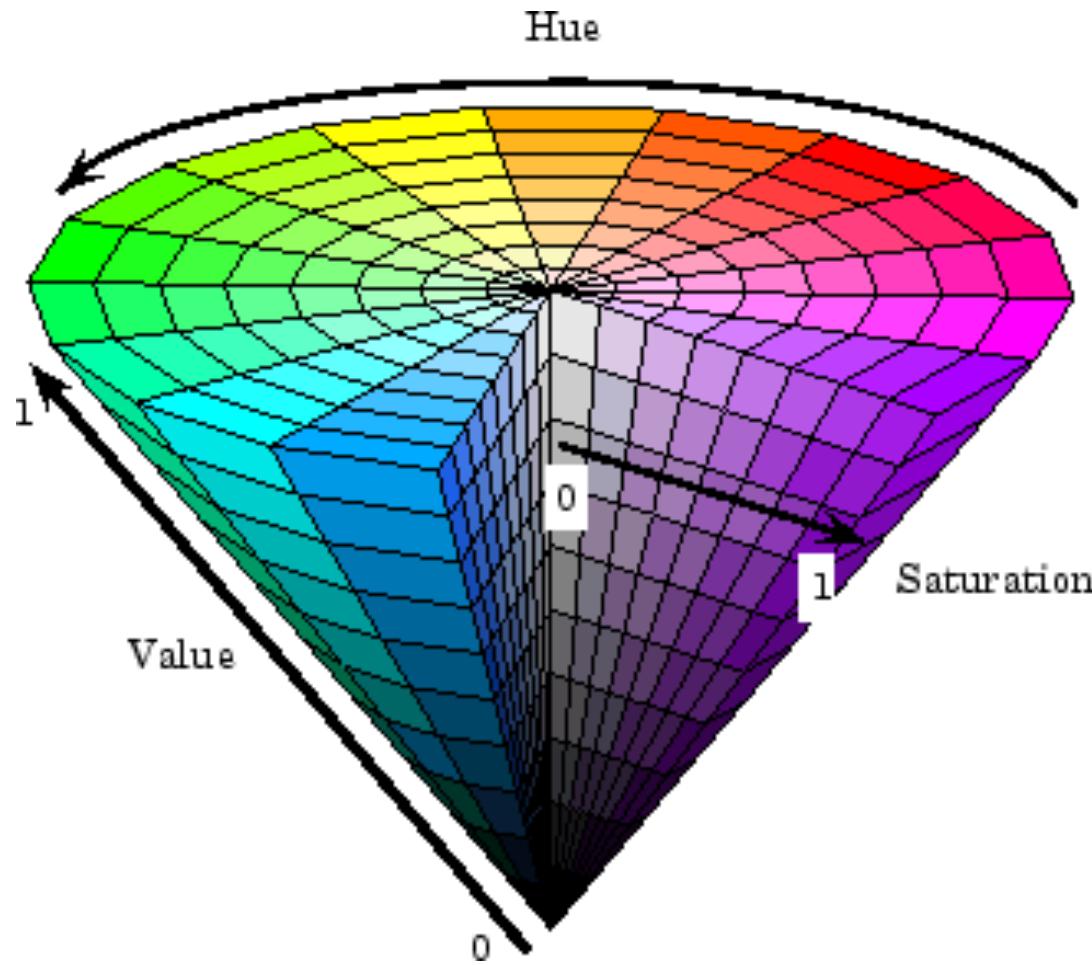
G  
(R=0,B=0)



B  
(R=0,G=0)

Image from: [http://en.wikipedia.org/wiki/File:RGB\\_color\\_solid\\_cube.](http://en.wikipedia.org/wiki/File:RGB_color_solid_cube.)

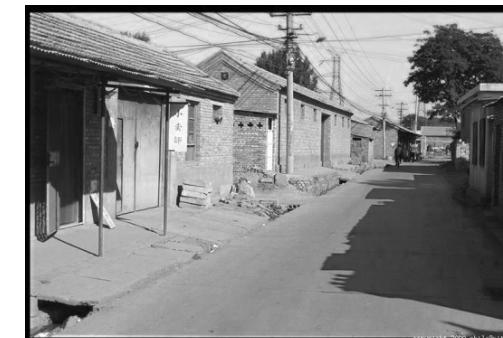
# HSV space



H  
( $S=1, V=1$ )



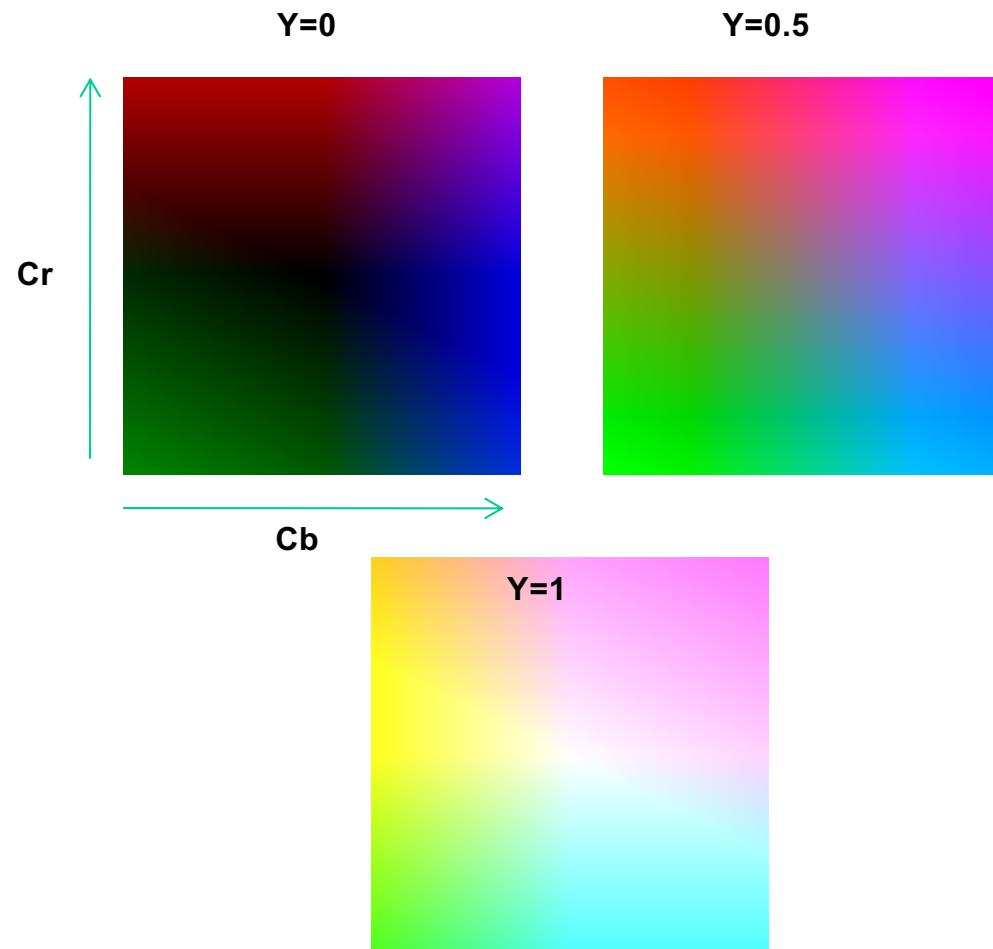
S  
( $H=1, V=1$ )



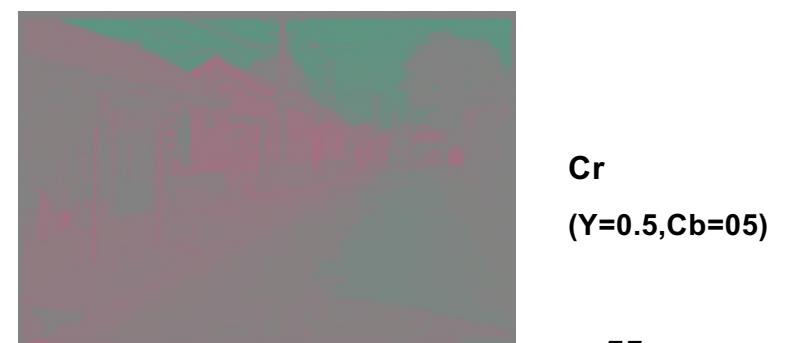
V  
( $H=1, S=0$ )

Intuitive color space

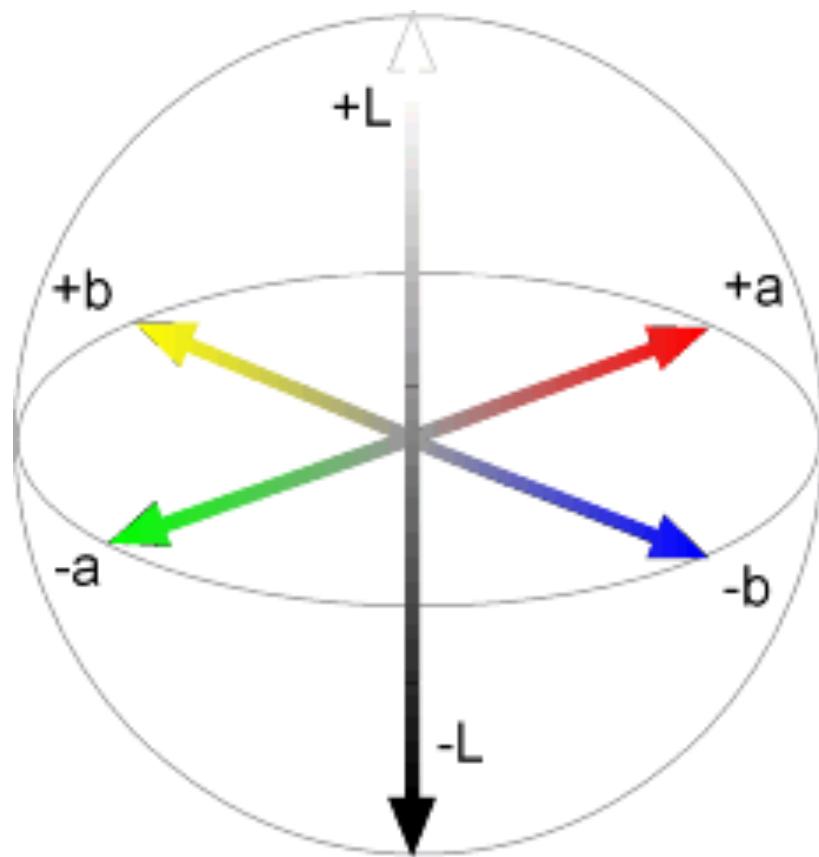
# YCbCr (YUV) space



Fast to compute, good for compression, used by TV



# CIE-L<sup>\*</sup>a<sup>\*</sup>b<sup>\*</sup> space



“Perceptually uniform” color space



L  
(a=0,b=0)



a  
(L=65,b=0)



b  
(L=65,a=0)

# Conversions between different colour spaces

---

$$H = \arccos \frac{\frac{1}{2}((R-G)+(R-B))}{\sqrt{(R-G)^2 + (R-B)(G-B)}}$$

$$S = 1 - 3 \frac{\min(R, G, B)}{R+G+B}$$

$$V = \frac{1}{3}(R+G+B)$$

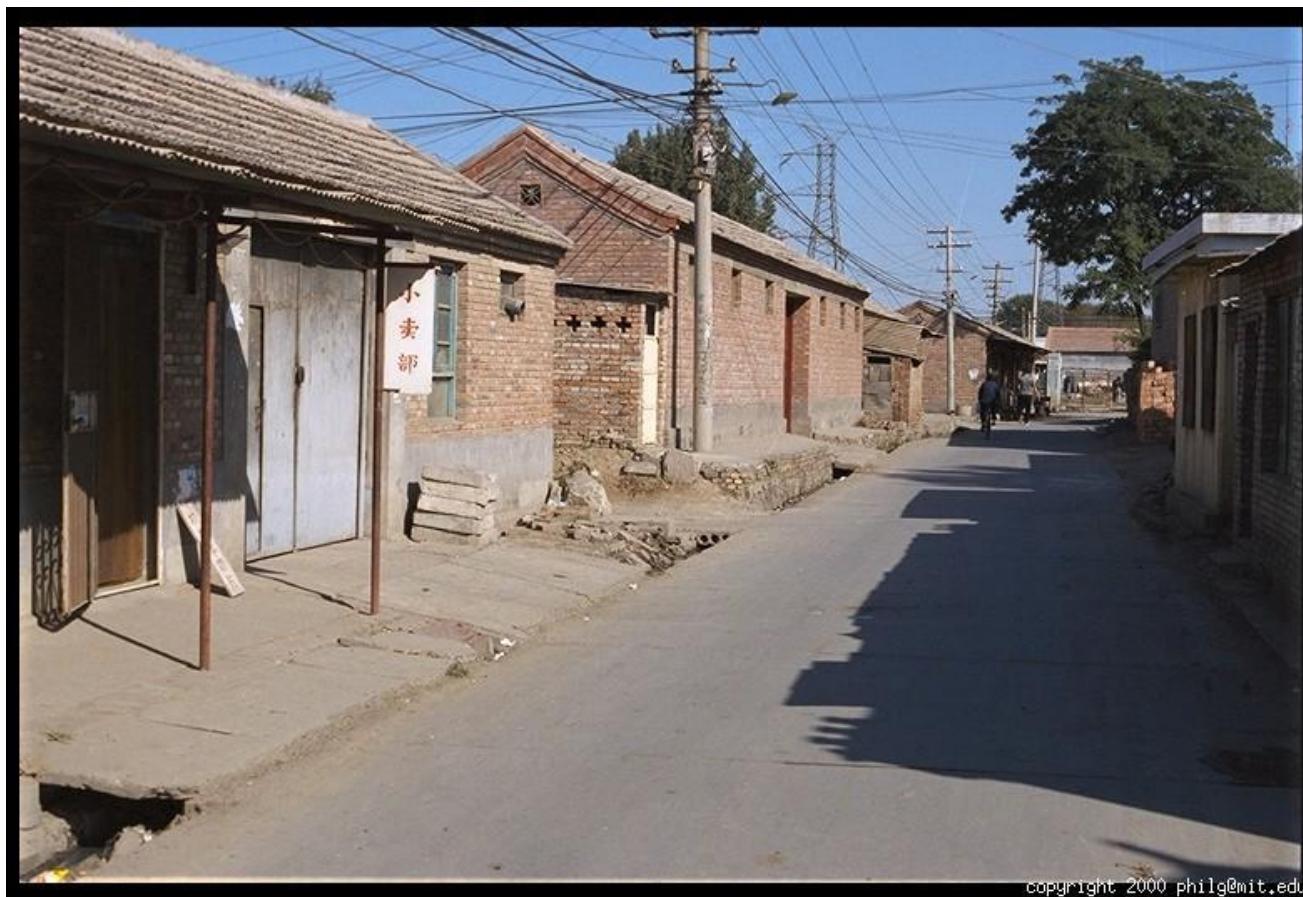
$$Y = 0.299R + 0.587G + 0.114B$$

$$C_r = R - Y$$

$$C_b = B - Y$$

# Most semantic information is however contained in the intensity band

---



Original image

# Most information is contained in intensity channel

