

Image Filtering

-Linear Filtering-

Disclaimer: Many of the slides used here are obtained from online resources (including many open lecture materials given at MIT, etc.) without due acknowledgement. They are used here for the sole purpose of classroom teaching. All the credits and all the copyrights belong to the original authors. You should not copy it, redistribute it, put it online, or use it for any, any purposes other than helping you study the course of ENGN4528/6528.

Computer Labs

- Lab enrolments are open again, please signup **before 11:59 pm Thursday. No more extensions will be given!**
- 3 Groups:
 - Clab-group-1: Thursday (10-12)
 - Clab-group-2: Friday (12-14)
 - Clab-group-3: Friday(**14-16**), new time!
- First lab assignment:
 - Due on August 23rd, 2020.
 - Late submissions will incur 10% per day (>7 days: 0 marks)
- No mid-term exam.

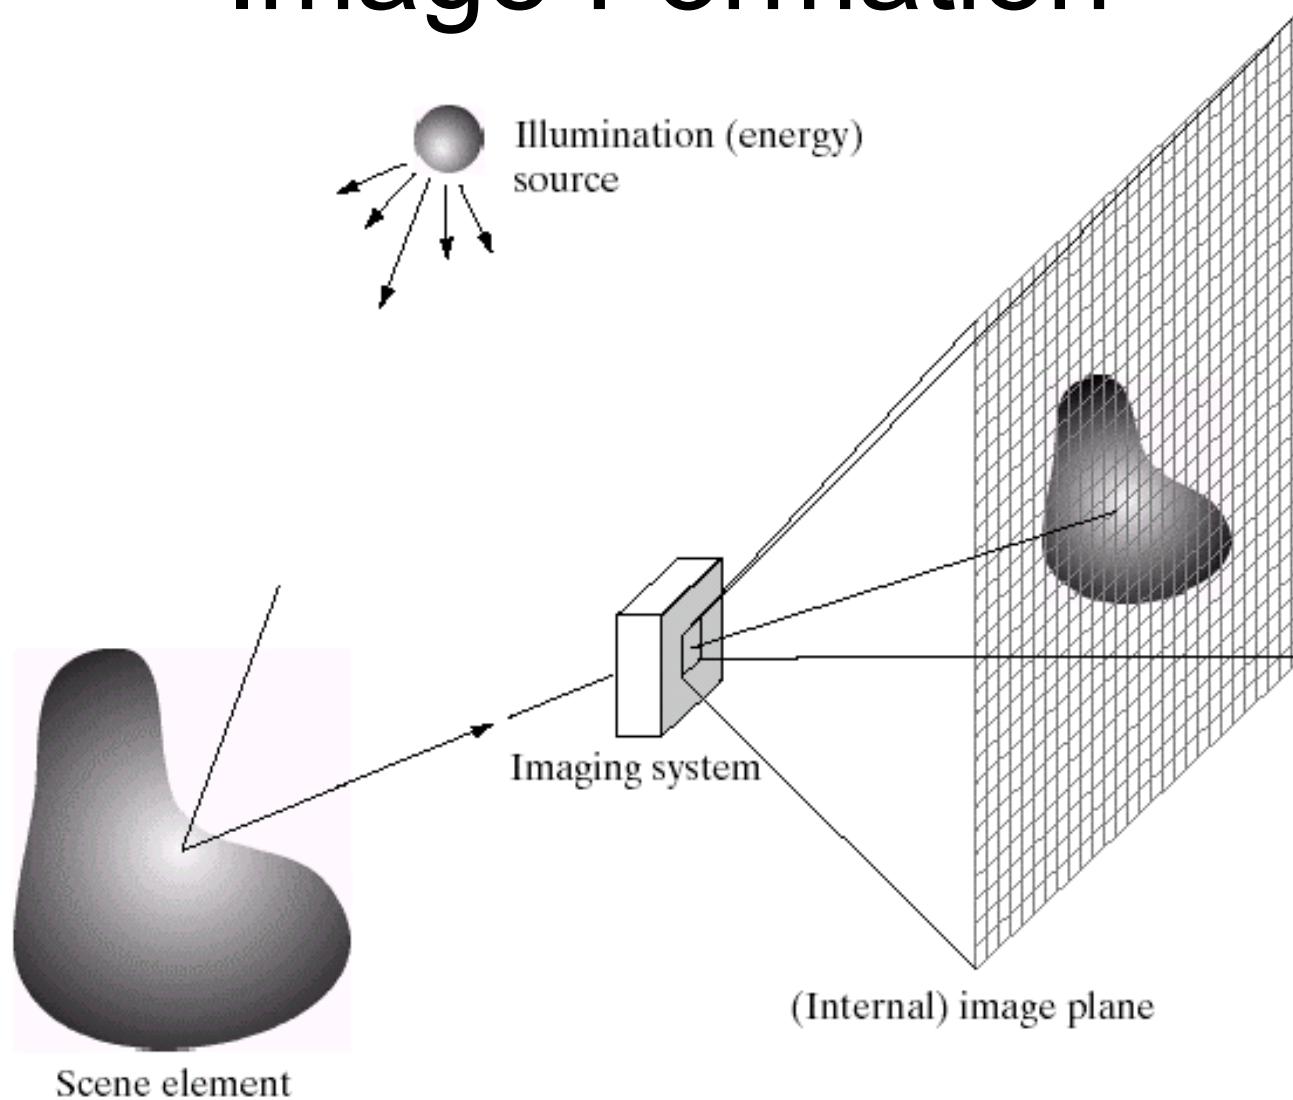
Lecture schedule (Tentative)

Week	Tuesday (2-hr)	Lecturer	Wednesday (1-hr)	Lecturer
Week-3	Image filtering	Ajanthan	Tutorial-2	Sameera
Week-4	Line fitting, corners	Richard	SIFT	Richard
Week-5	Segmentation, clustering	Richard	Tutorial-3, clab-2 intro	Sameera
Week-6	PCA, eigen faces	Richard/Ajanthan	TBD	TBD
Week-7	3D-vision-A	Richard	3D-vision-B	Richard
Week-8	3D-vision-C	Richard	Tutorial-4, clab-3 intro	Ryan
Week-9	Stereo, SLAM	Richard/Ajanthan	TBD	TBD
Week-10	Classification basics, SVM	Richard/Ajanthan	TBD	TBD
Week-11	Deep learning-A	Ajanthan	Deep learning-B	Ajanthan
Week-12	Guest lecture	TBD	Review	Ajanthan

Plan for today

- Review Photometric Image Formation
- Image Noise
- Linear Filtering
 - Correlation & Convolution
 - Gaussian Filter
 - Edge Detection

Image Formation



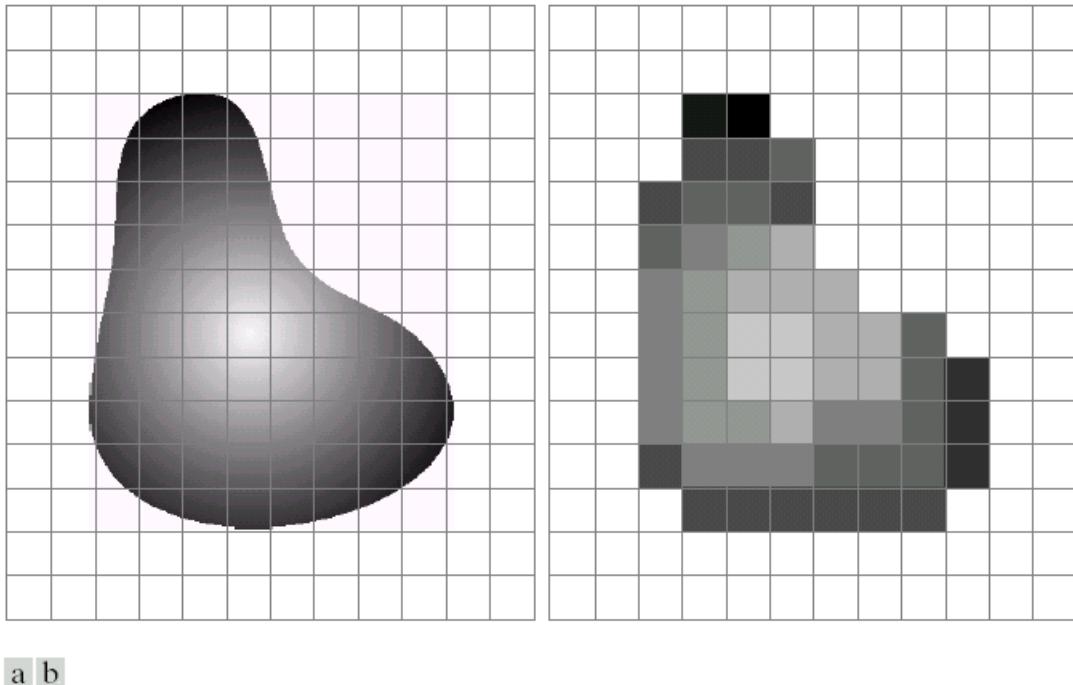
Digital camera



A digital camera replaces film with a sensor array

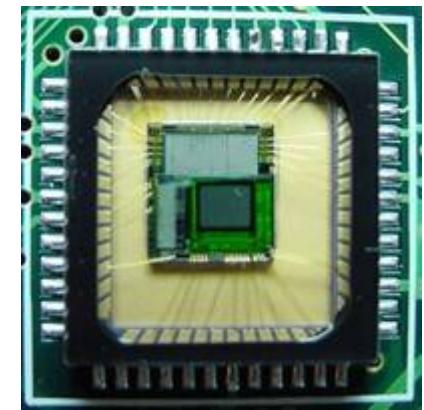
- Each cell in the array is light-sensitive diode that converts photons to electrons
- <http://electronics.howstuffworks.com/digital-camera.htm>

Digital images



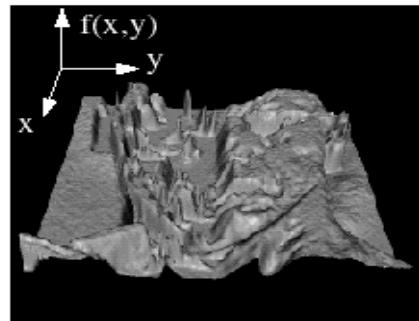
a b

FIGURE 2.17 (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.



Digital images

- **Sample** the 2D space on a regular grid
- **Quantize** each sample (round to nearest integer)
- Image thus represented as a matrix of integer values.

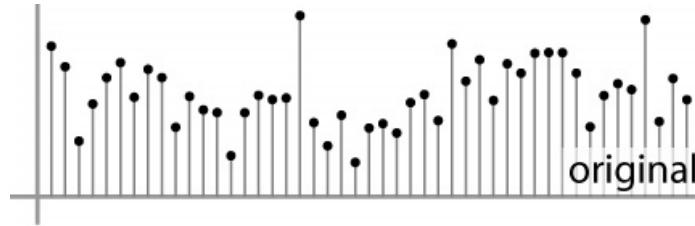


j →

i ↓

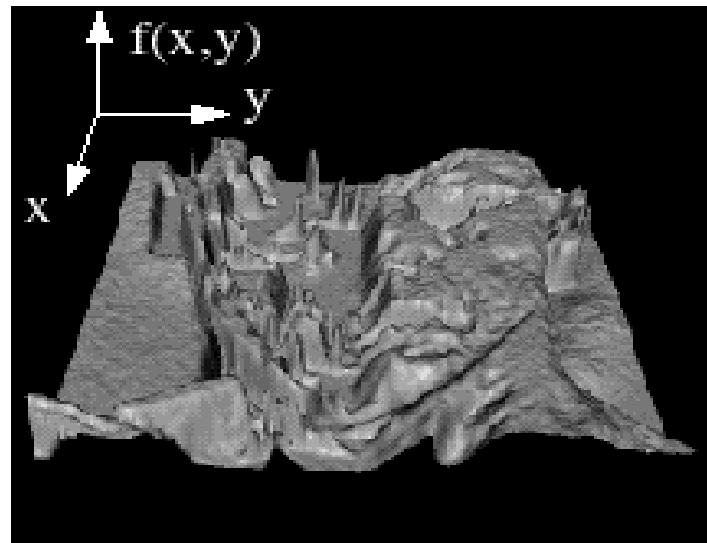
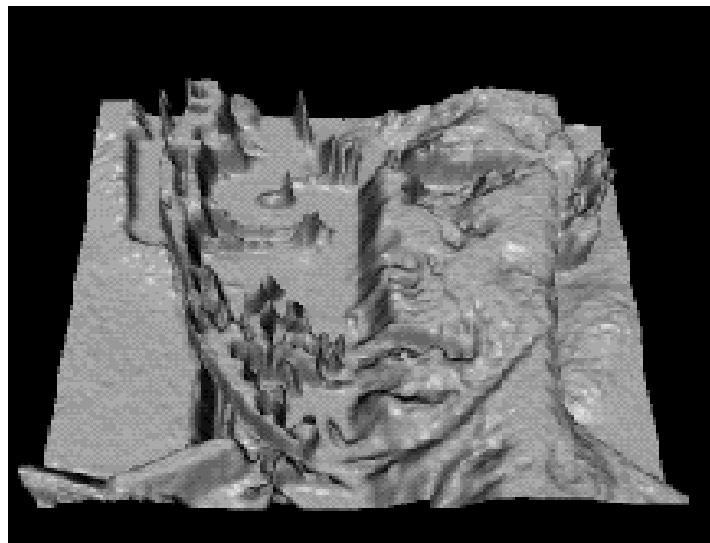
62	79	23	119	120	105	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30

2D

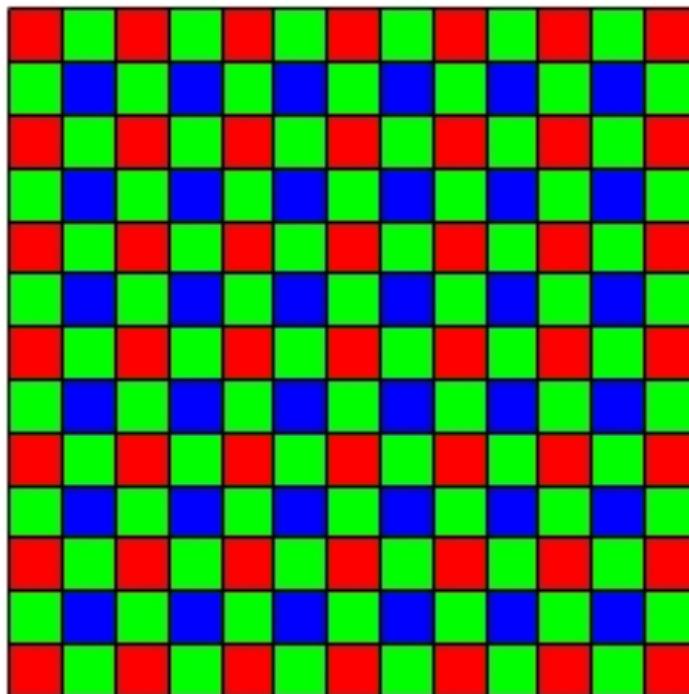


1D

Images as functions



Digital color images



Bayer filter

© 2000 How Stuff Works

Digital color images

Color images,
RGB color
space



R



G



B

Images as functions

- We can think of an **image** as a function, f , from \mathbb{R}^2 to \mathbb{R} :
 - $f(x, y)$ gives the **intensity** at position (x, y)
 - Realistically, we expect the image only to be defined over a rectangle, with a finite range:
 - $f: [a, b] \times [c, d] \rightarrow [0, 1]$

Images in Matlab

- Images represented as a matrix
- Suppose we have a NxM RGB image called “im”
 - $\text{im}(1,1,1)$ = top-left pixel value in R-channel
 - $\text{im}(y, x, b)$ = y pixels down, x pixels to right in the bth channel
 - $\text{im}(N, M, 3)$ = bottom-right pixel in B-channel
- `imread(filename)` returns a uint8 image (values 0 to 255)
 - Convert to double format (values 0 to 1) with `im2double`

row	column	R	G	B
0.92	0.93	0.94	0.97	0.62
0.95	0.89	0.82	0.89	0.56
0.89	0.72	0.51	0.55	0.51
0.96	0.95	0.88	0.94	0.56
0.71	0.81	0.81	0.87	0.57
0.49	0.62	0.60	0.58	0.50
0.86	0.84	0.74	0.58	0.51
0.96	0.67	0.54	0.85	0.48
0.69	0.49	0.56	0.66	0.43
0.79	0.73	0.90	0.67	0.33
0.91	0.94	0.89	0.49	0.41
		0.65	0.45	0.56
		0.79	0.73	0.90
		0.91	0.94	0.89
		0.65	0.45	0.56
		0.79	0.73	0.90
		0.91	0.94	0.89
		0.65	0.45	0.56
		0.79	0.73	0.90
		0.91	0.94	0.89

Images as matrices

Result of averaging 100 similar snapshots



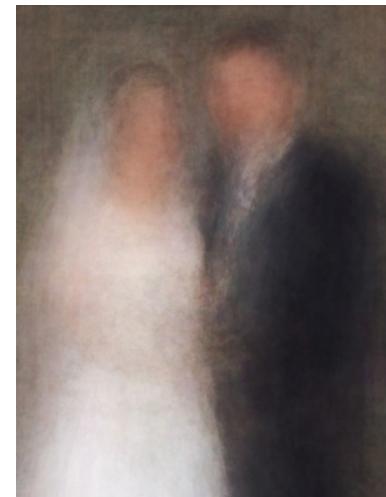
Little Leaguer



Kids with Santa



The Graduate



Newlyweds

From: *100 Special Moments*, by Jason Salavon (2004)
<http://salavon.com/SpecialMoments/SpecialMoments.shtml>

Image Noise

Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise

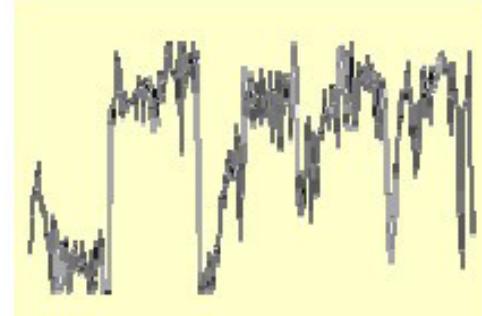
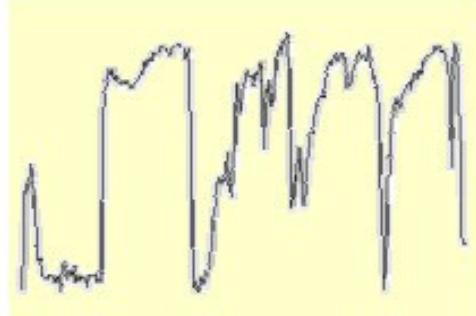
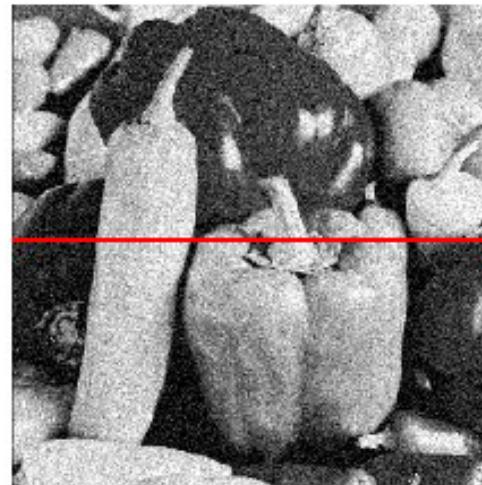
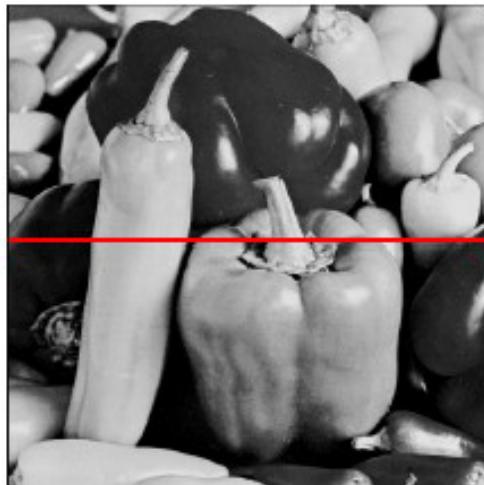


Impulse noise



Gaussian noise

Gaussian noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

```
>> noise = randn(size(im)).*sigma;  
>> output = im + noise;
```

What is impact of the sigma?

Fig: M. Hebert

Effect of
sigma on
Gaussian
noise:

Image
shows the
noise values
themselves.

$\sigma = 1$

Effect of
sigma on
Gaussian
noise:

Image
shows the
noise values
themselves.

sigma=4

Effect of
sigma on
Gaussian
noise:

sigma=1



This shows
the noise
values
added to the
raw
intensities of
an image.

Effect of
sigma on
Gaussian
noise:

Image
shows the
noise values
themselves.

sigma=
16

Effect of
sigma on
Gaussian
noise

sigma=16



This shows
the noise
values
added to the
raw
intensities of
an image.

Noise Reduction

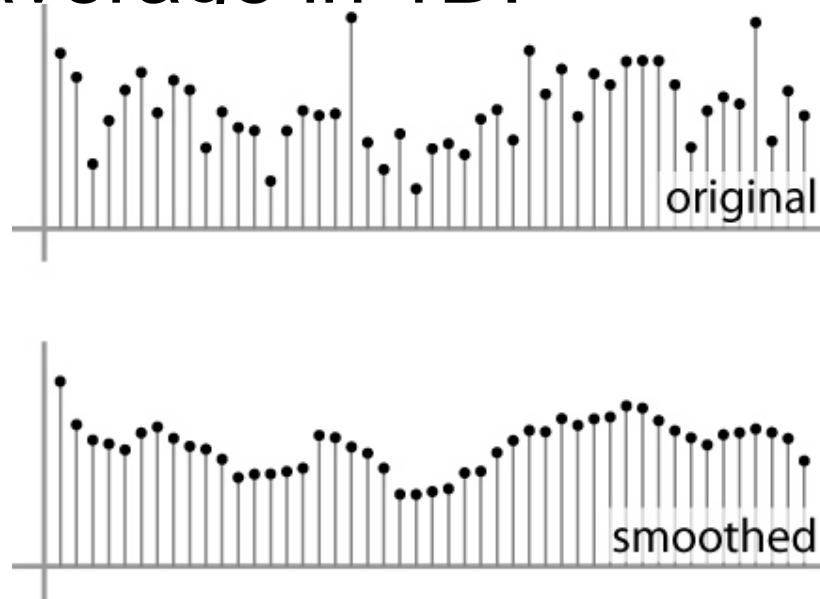
- How should we reduce the noise?

First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
 - Expect pixels to be like their neighbors
 - Expect noise processes to be independent from pixel to pixel

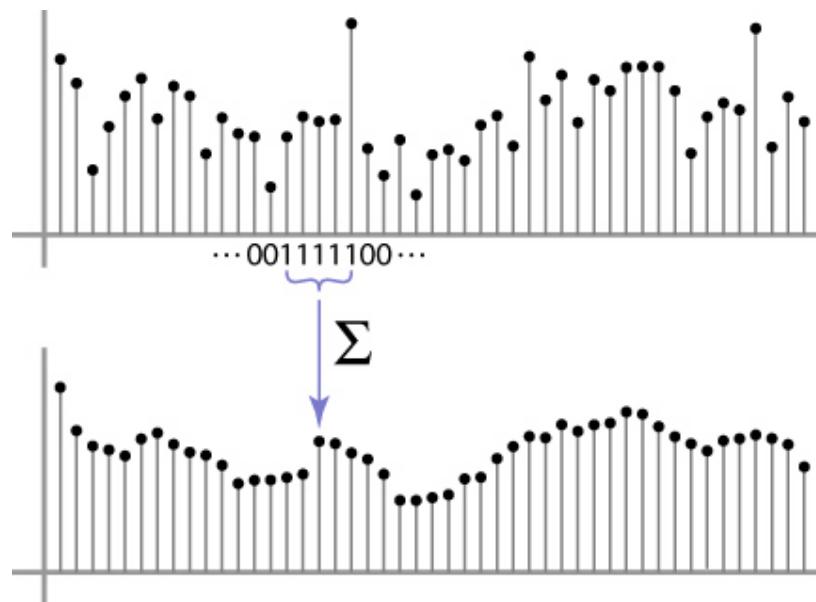
First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



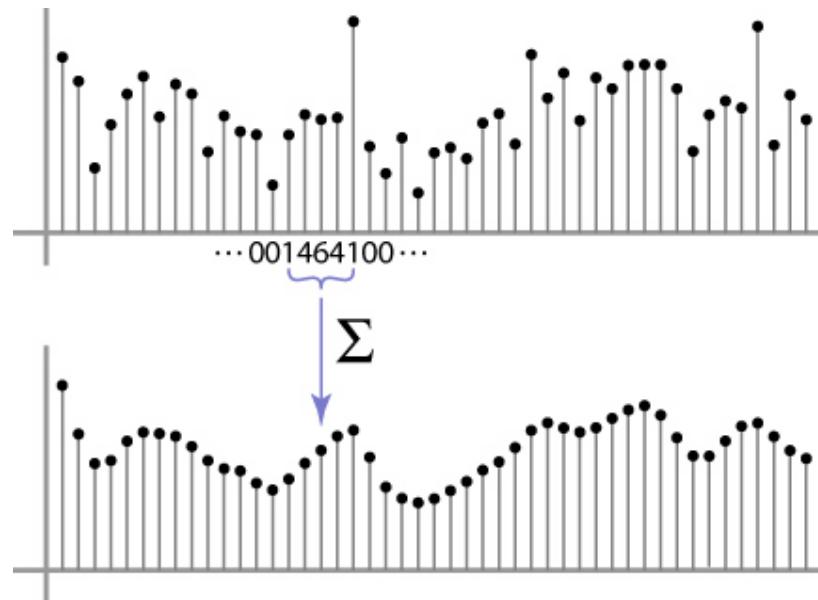
Weighted Moving Average

- Can add weights to our moving average
- *Weights* $[1, 1, 1, 1, 1] / 5$



Weighted Moving Average

- Non-uniform weights $[1, 4, 6, 4, 1] / 16$



Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

			0							

Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

0	10									

Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

			0	10	20					

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0 10 20 30

Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Correlation filtering

Say the averaging window size is $(2k+1) \times (2k+1)$:

$$G[i, j] = \frac{1}{(2k+1)^2} \underbrace{\sum_{u=-k}^k \sum_{v=-k}^k}_{\text{Attribute uniform weight to each pixel}} F[i + u, j + v] \underbrace{\qquad\qquad\qquad}_{\text{Loop over all pixels in neighborhood around image pixel } F[i,j]}$$

Now generalize to allow **different weights** depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i + u, j + v]$$

Correlation filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called cross-correlation, denoted $G = H \otimes F$

Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter “kernel” or “mask” $H[u, v]$ is the prescription for the weights in the linear combination.

Averaging filter

- What values belong in the kernel H for the moving average example?

$F[x, y]$								
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	0	0
0	0	0	90	90	90	90	0	0
0	0	0	90	90	90	90	0	0
0	0	0	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0
0	0	0	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0



$H[u, v]$								
1	1	1						
1	?	1						
1	1	1						

$$\frac{1}{9}$$

“box filter”

$G[x, y]$								
	0	20	40	60	60			

$$G = H \otimes F$$

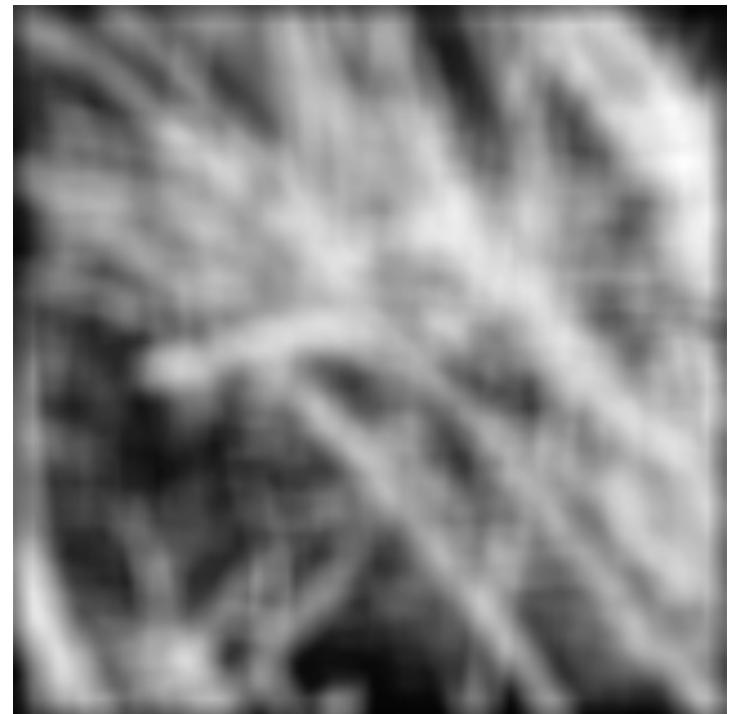
Smoothing by averaging



depicts box filter:
white = high value, black = low value



original



filtered

What if the filter size was 5×5 instead of 3×3 ?

Image Filtering

- Image filtering: compute the function value of local neighborhood at each pixel.
- Very useful
 - **Enhance images**
 - De-noise, resize, increase contrast, etc.
 - **Extract information from images**
 - Texture, edges, distinctive points, etc.
 - **Detect patterns**
 - Template matching

Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

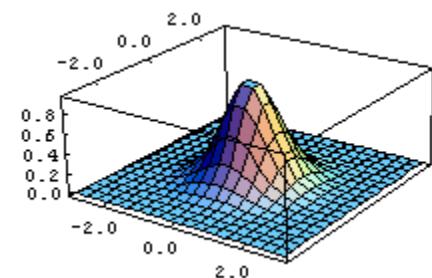
$F[x, y]$

$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$

$H[u, v]$

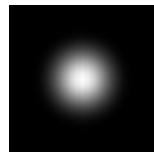
- This kernel is an approximation of a 2d Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



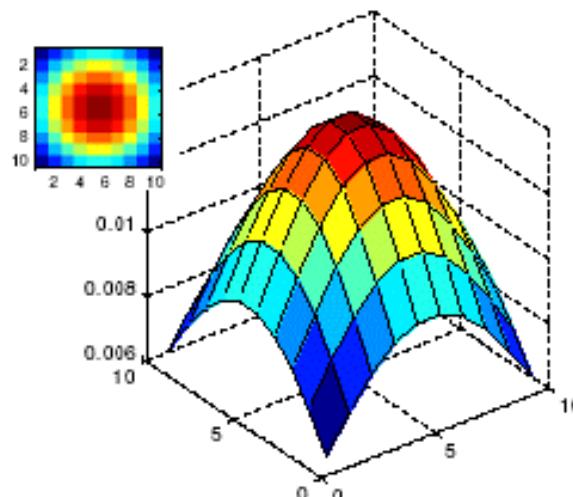
- Removes high-frequency components from the image (“low-pass filter”).

Smoothing with a Gaussian

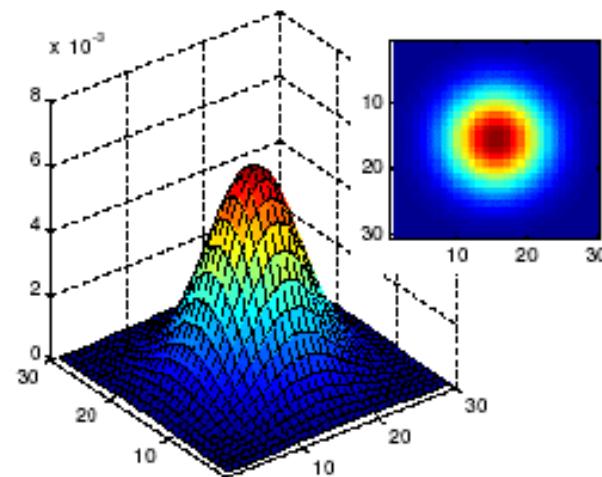


Gaussian filters

- What parameters matter here?
- **Size** of kernel or mask
 - Note, Gaussian function has infinite support, but discrete filters use finite kernels



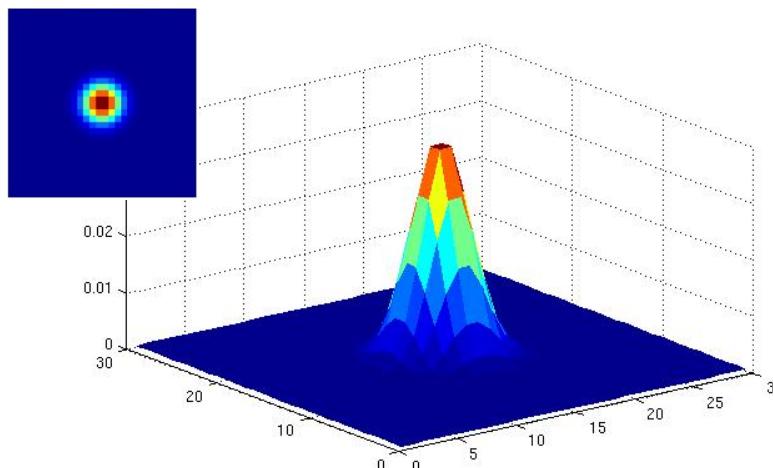
$\sigma = 5$ with
10 x 10
kernel



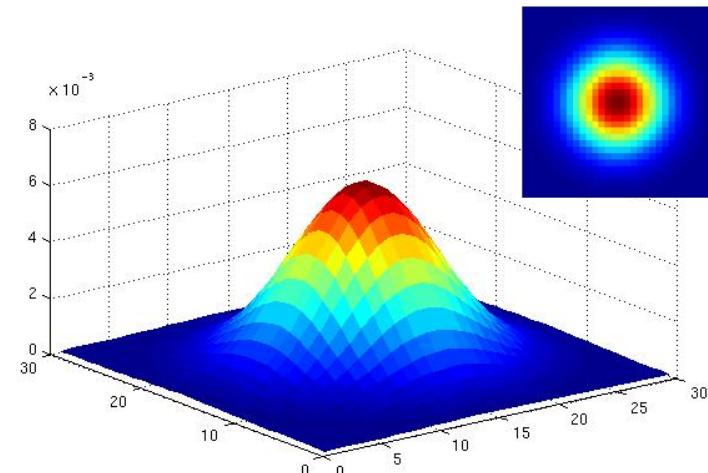
$\sigma = 5$ with
30 x 30
kernel

Gaussian filters

- What parameters matter here?
- **Variance** of Gaussian: determines extent of smoothing



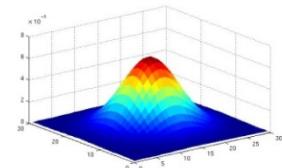
$\sigma = 2$ with
30 x 30
kernel



$\sigma = 5$ with
30 x 30
kernel

Matlab

```
>> hsize = 10;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```

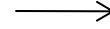


```
>> mesh(h);
```



```
>> imagesc(h);
```

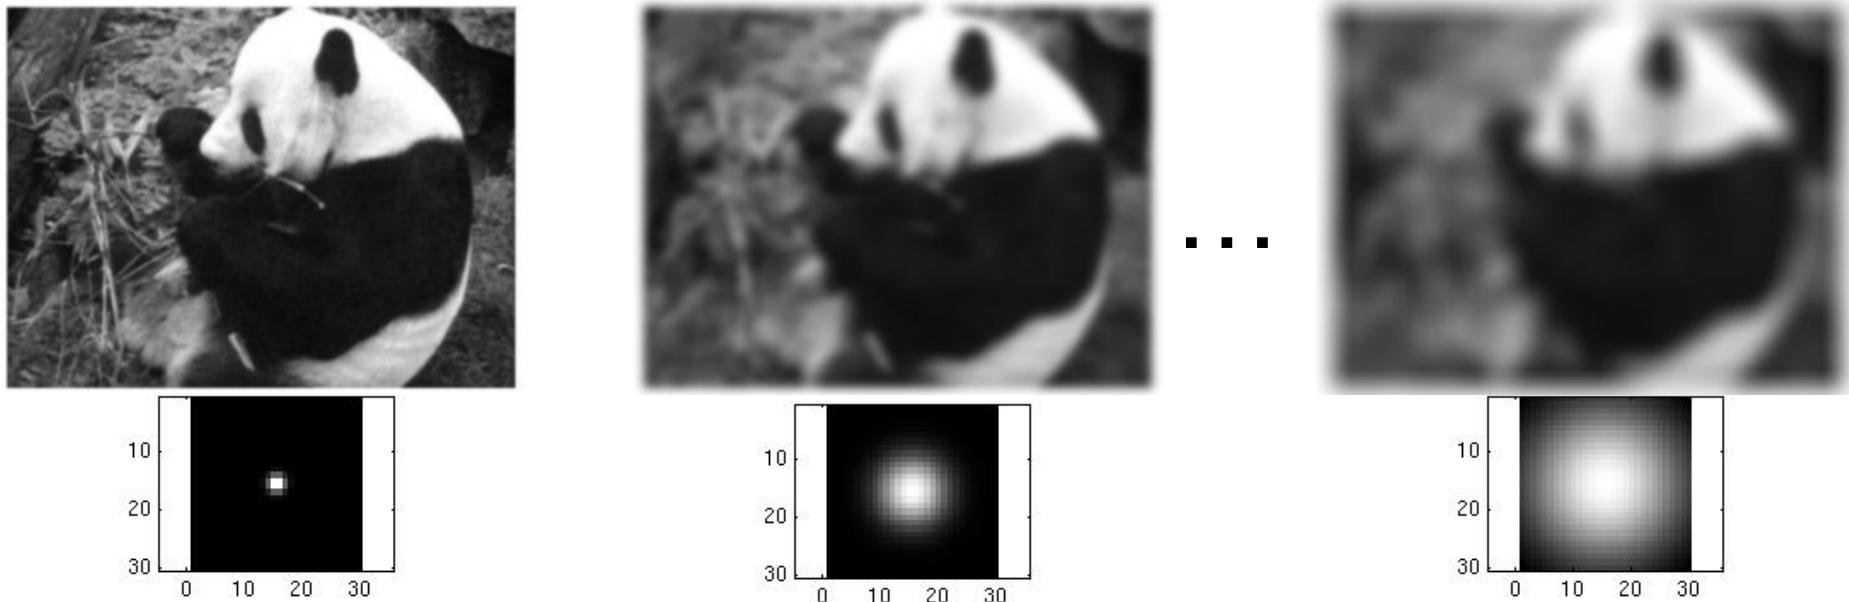
```
>> outim = imfilter(im, h); % correlation  
>> imshow(outim);
```



outim

Smoothing with a Gaussian

Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

•Slide credit:
Kristen Grauman

Convolution

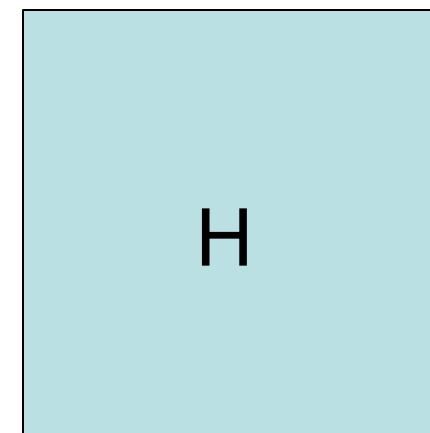
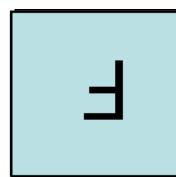
- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$



*Notation for
convolution
operator*



Convolution

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation

$$G(i, j) = \sum_{u=0}^M \sum_{v=0}^N F(u, v) H(i - u, j - v)$$

$$G = F \star H$$



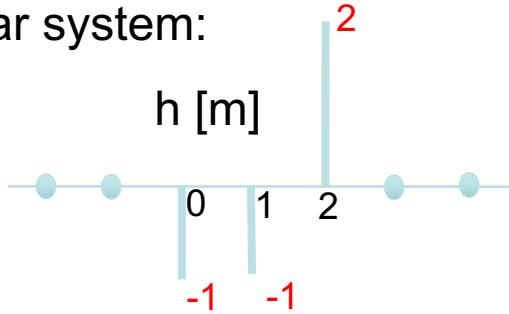
*Notation for
convolution
operator*

Convolution vs Cross correlation

- **Cross correlation:** The measurement of **similarity** between the two signals H and F.
- **Convolution:** The measurement of the **effect** of a signal (H) on the other signal (F).
- When the filter (H) is symmetric, the results are identical.

1-D linear filtering

Linear system:



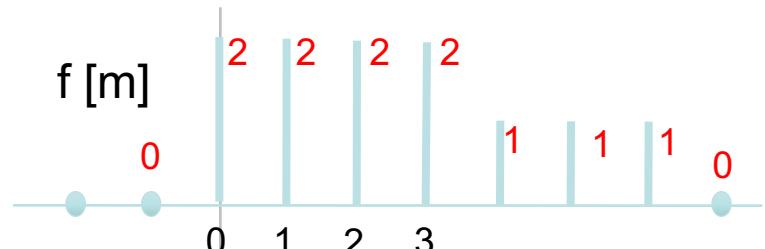
Output?

$$g(m = 0) = \sum_k h(-k)f(k)$$

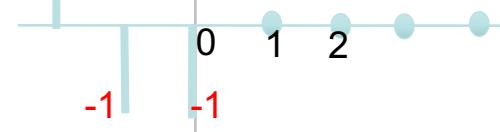
$$g(m = 1) = \sum_k h(1 - k)f(k)$$

$$g(m = 2) = \sum_k h(2 - k)f(k)$$

Input:



$h[-k]$



$$g[m=0] = -2$$

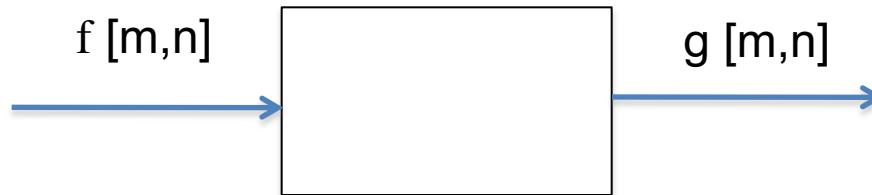
$h[1-k]$



$$g[m=1] = -4$$

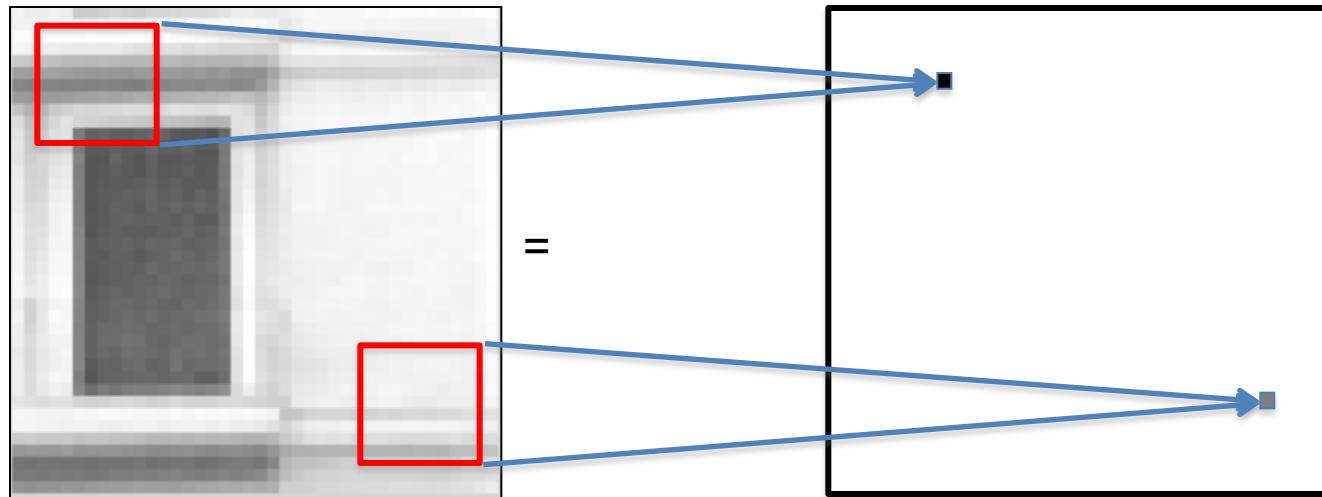
$$g[m=2] = 0$$

Discrete linear system



In vision, many times, we are interested in operations that are spatially invariant.
For a linear spatially invariant system:

$$g[m, n] = h \star f = \sum_{k,l} h[m - k, n - l]f[k, l]$$



Matlab's Filter vs. Convolution

- 2D filtering

- `f=filter2(h,g);` or
`f=imfilter(h,g);`

`h=filter` `g=image`



$$f[m, n] = \sum_{k,l} h[k, l] g[m + k, n + l]$$

- 2D convolution

- `f=conv2(h,g);`

$$f[m, n] = \sum_{k,l} h[k, l] g[m - k, n - l]$$

Properties of convolution

- **Shift invariant:**
 - Operator behaves the same everywhere, i.e. the value of the output depends on the pattern in the image neighborhood, not the position of the neighborhood.
- **Superposition:**
 - $h * (f_1 + f_2) = (h * f_1) + (h * f_2)$

Properties of convolution

- Commutative:

$$f * g = g * f$$

- Associative

$$(f * g) * h = f * (g * h)$$

- Distributes over addition

$$f * (g + h) = (f * g) + (f * h)$$

- Scalars factor out

$$kf * g = f * kg = k(f * g)$$

- Identity:

$$\text{unit impulse } e = [\dots, 0, 0, 1, 0, 0, \dots]. \quad f * e = f$$

Separability

- In some cases, filter is separable, and we can factor into two steps:
 - Convolve all rows with a 1D filter
 - Convolve all columns with a 1D filter

$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix} = \begin{bmatrix} 0 & \frac{1}{3} & 0 \\ 0 & \frac{1}{3} & 0 \\ 0 & \frac{1}{3} & 0 \end{bmatrix} \circ \begin{bmatrix} 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ 0 & 0 & 0 \end{bmatrix}$$

Separability example

- 2D convolution (center location only)(65)

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = 2 + 6 + 3 = 11$$
$$= 6 + 20 + 10 = 36$$
$$= 4 + 8 + 6 = 18$$

$$65$$

Separability example

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} 2 + 6 + 3 = 11 \\ 6 + 20 + 10 = 36 \\ 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

- The filter is factored into a product of 1D filters.

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Separability example

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} 2 + 6 + 3 = 11 \\ 6 + 20 + 10 = 36 \\ 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

- Perform convolutions along rows.

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 11 & & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array}$$

Source: K. Grauman

Separability example

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} 2 + 6 + 3 = 11 \\ 6 + 20 + 10 = 36 \\ 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & 11 & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array}$$

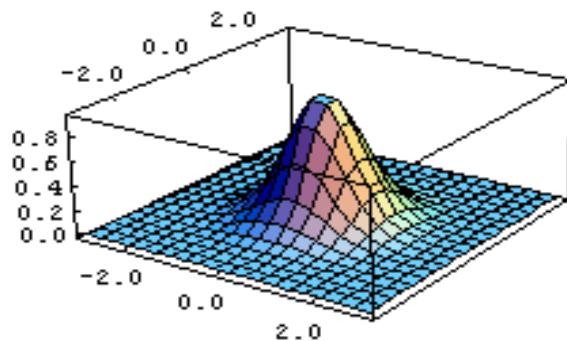
- Perform convolutions along columns.

$$\begin{array}{|c|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline & 11 & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & 65 & \\ \hline & & \\ \hline \end{array}$$

Source: K. Grauman

This kernel is an approximation of
a 2d Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



Properties of Gaussian filter

- Convolution with self gives another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into the product of two 1D Gaussians

Separability of the Gaussian filter

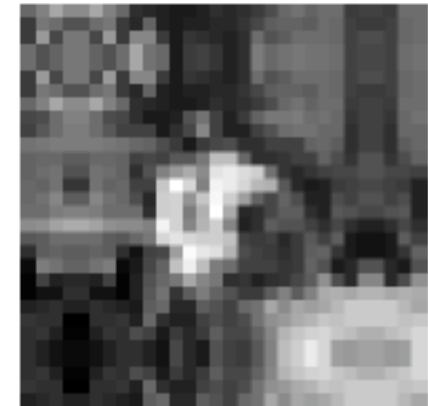
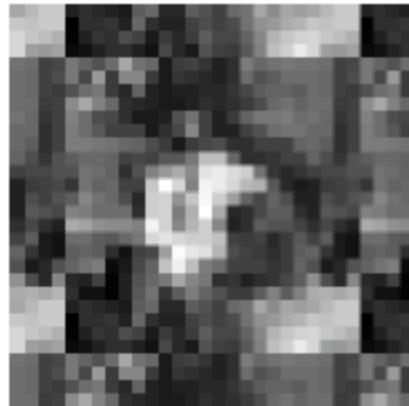
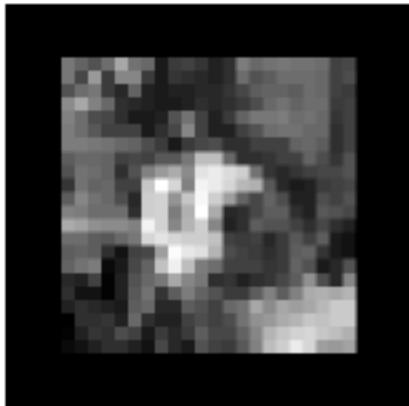
$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

- Separability means that a 2D convolution can be reduced to two 1D convolutions (one among rows and one among columns)
 - What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 - $O(n^2 m^2)$
 - What if the kernel is separable?
 - $O(n^2 m)$

How to handle borders

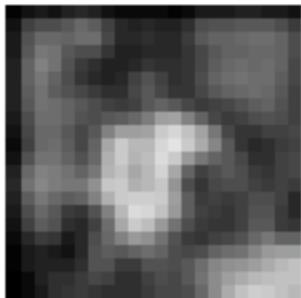


zero

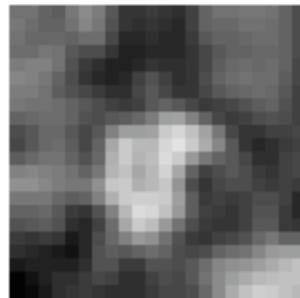
wrap

clamp

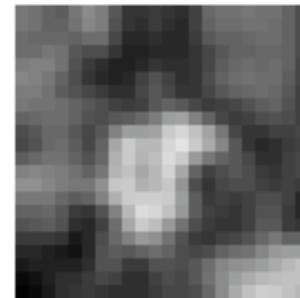
mirror



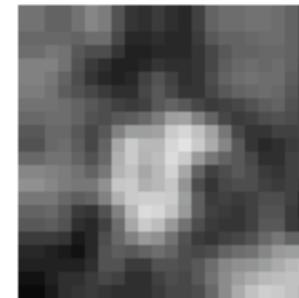
blurred: zero



normalized zero



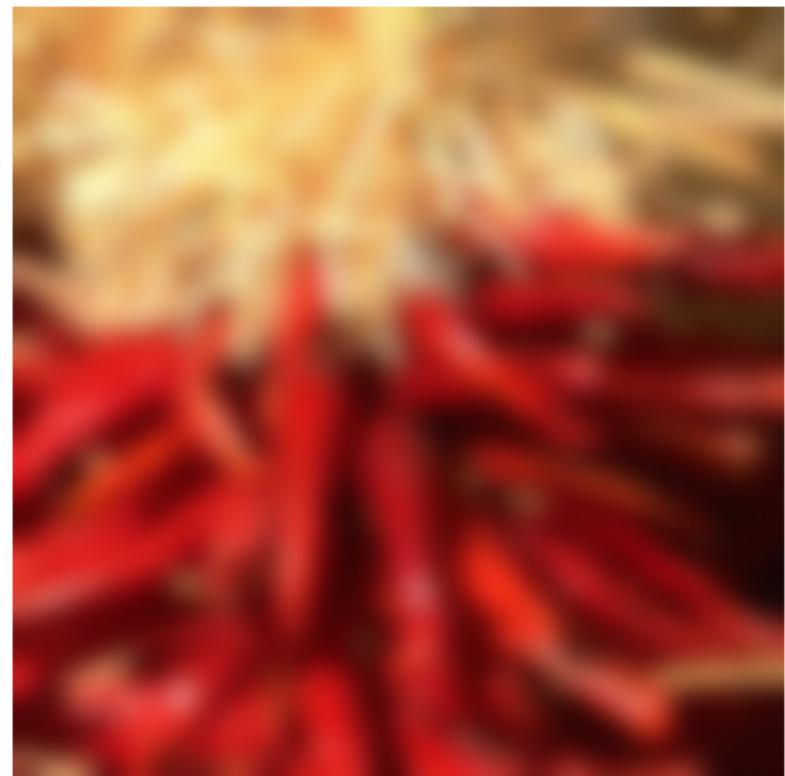
clamp



mirror

From Szeliski, Computer Vision, 2010

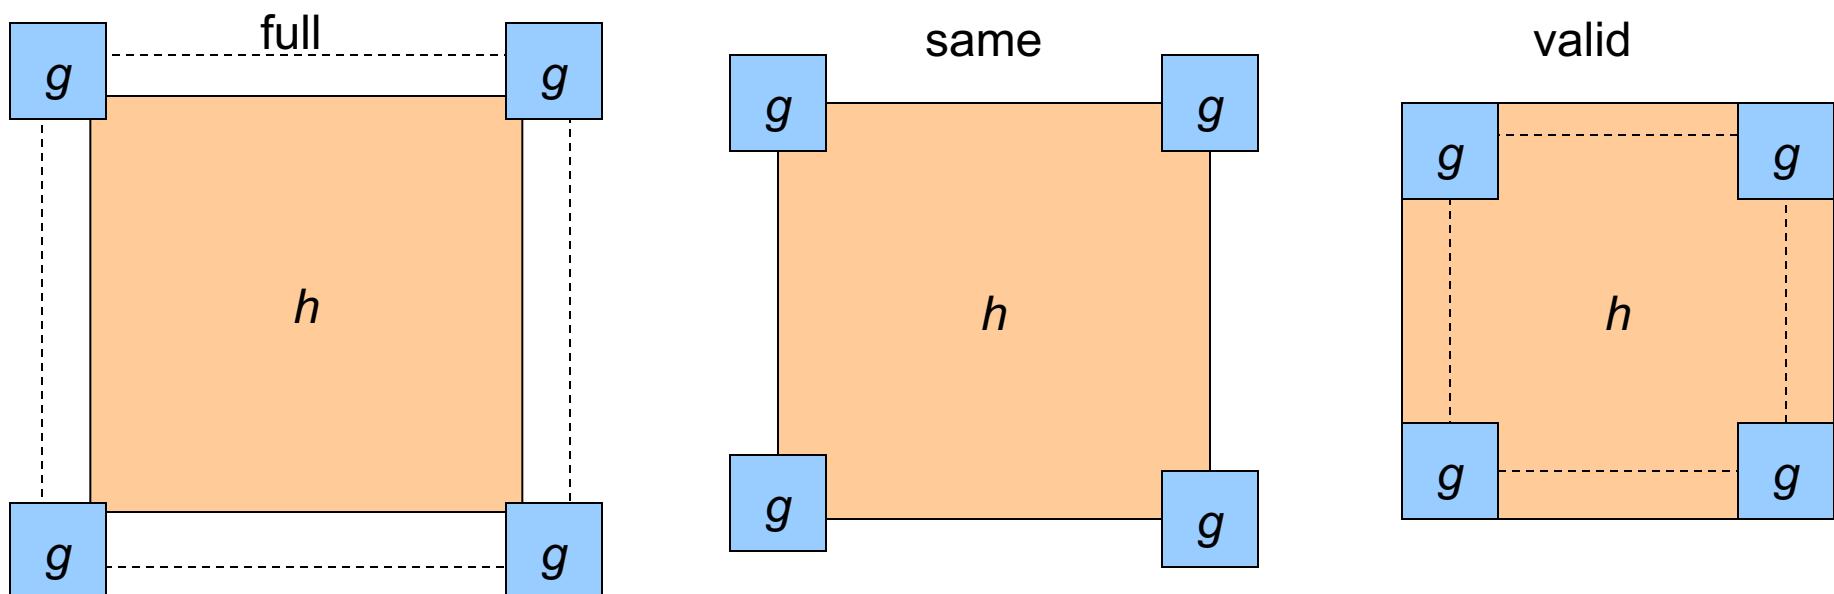
Example: Border handling



Source: S. Marschner

Matlab's border handling

- What is the size of the output?
- MATLAB: `filter2(g, h, shape)`
 - `shape = 'full'`: output size is sum of sizes of `h` and `g`
 - `shape = 'same'`: output size is same as `h`
 - `shape = 'valid'`: output size is difference of sizes of `h` and `g`



Discrete Filtering

- Linear filter: Weighted sum of pixels over rectangular neighborhood—*kernel* defines weights
- Think of kernel as template being matched by **correlation** (Matlab: `imfilter`, `filter2`)
- **Convolution**: Correlation with kernel rotated 180° (Matlab: `conv2`)

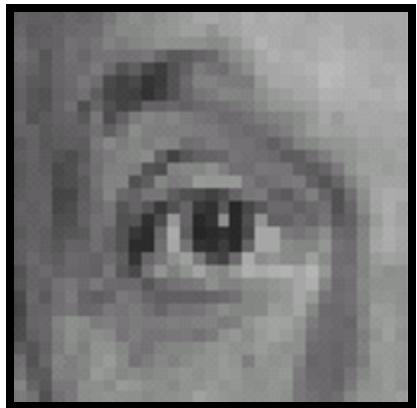
1	-1	-1
1	2	-1
1	1	1

Practice with linear filters

A taxonomy of useful filters

- Impulse, Shifts,
- Blur
 - Box filter
 - Gaussian filter
 - Bilateral filter
 - Asymmetrical filter: motion blur
- Edges
 - [-1 1]
 - Derivative filter
 - Derivative of a gaussian
 - Oriented filters
 - Gabor filter
 - Quadrature filters: phase and magnitude.
 - Elongated edges: filling gaps...

Practice with linear filters

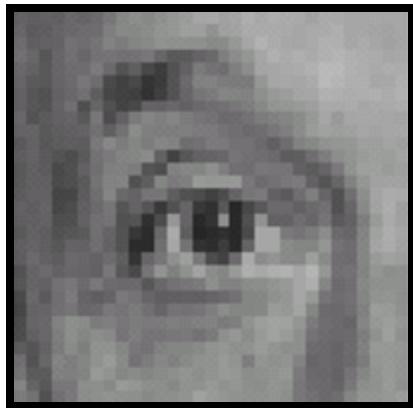


Original

0	0	0
0	1	0
0	0	0

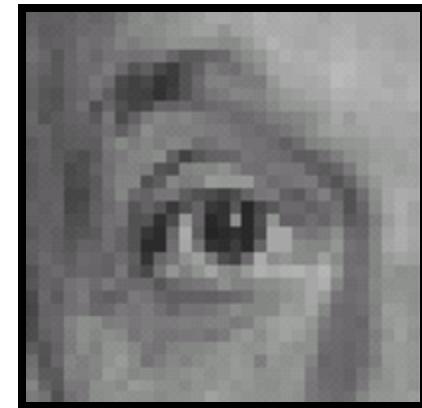
?

Practice with linear filters



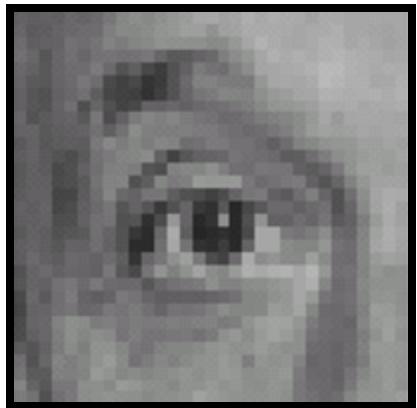
Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters

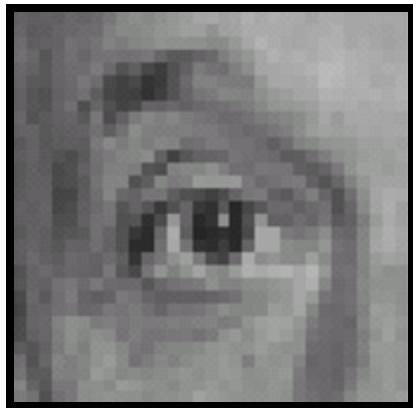


Original

0	0	0
0	0	1
0	0	0

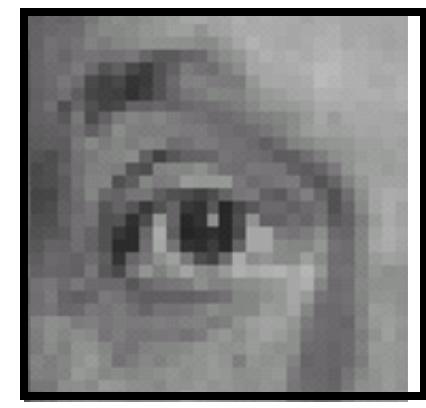
?

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

Practice with linear filters



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

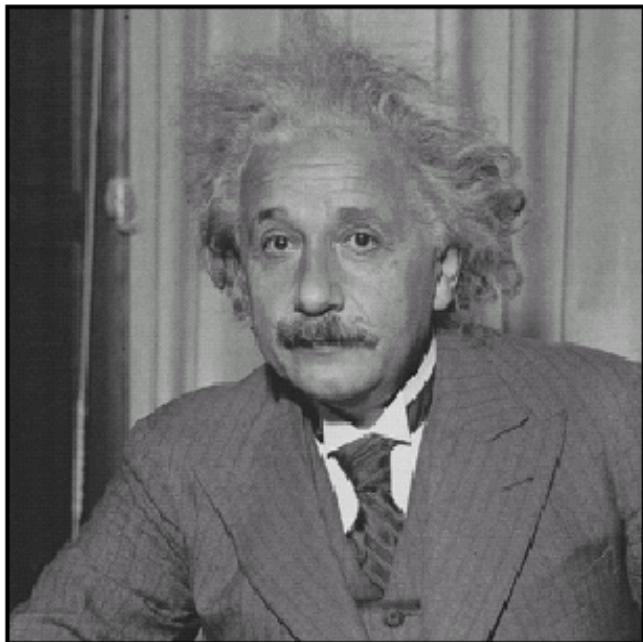


Original

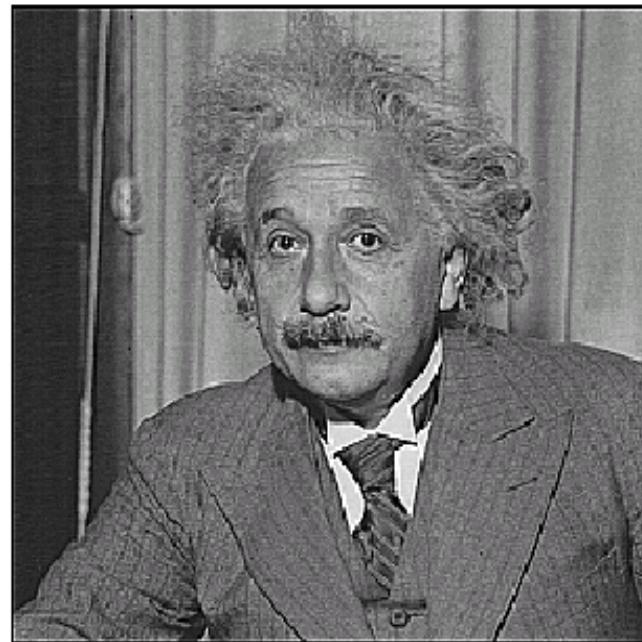
Sharpening filter

- Accentuates differences with local average

Sharpening

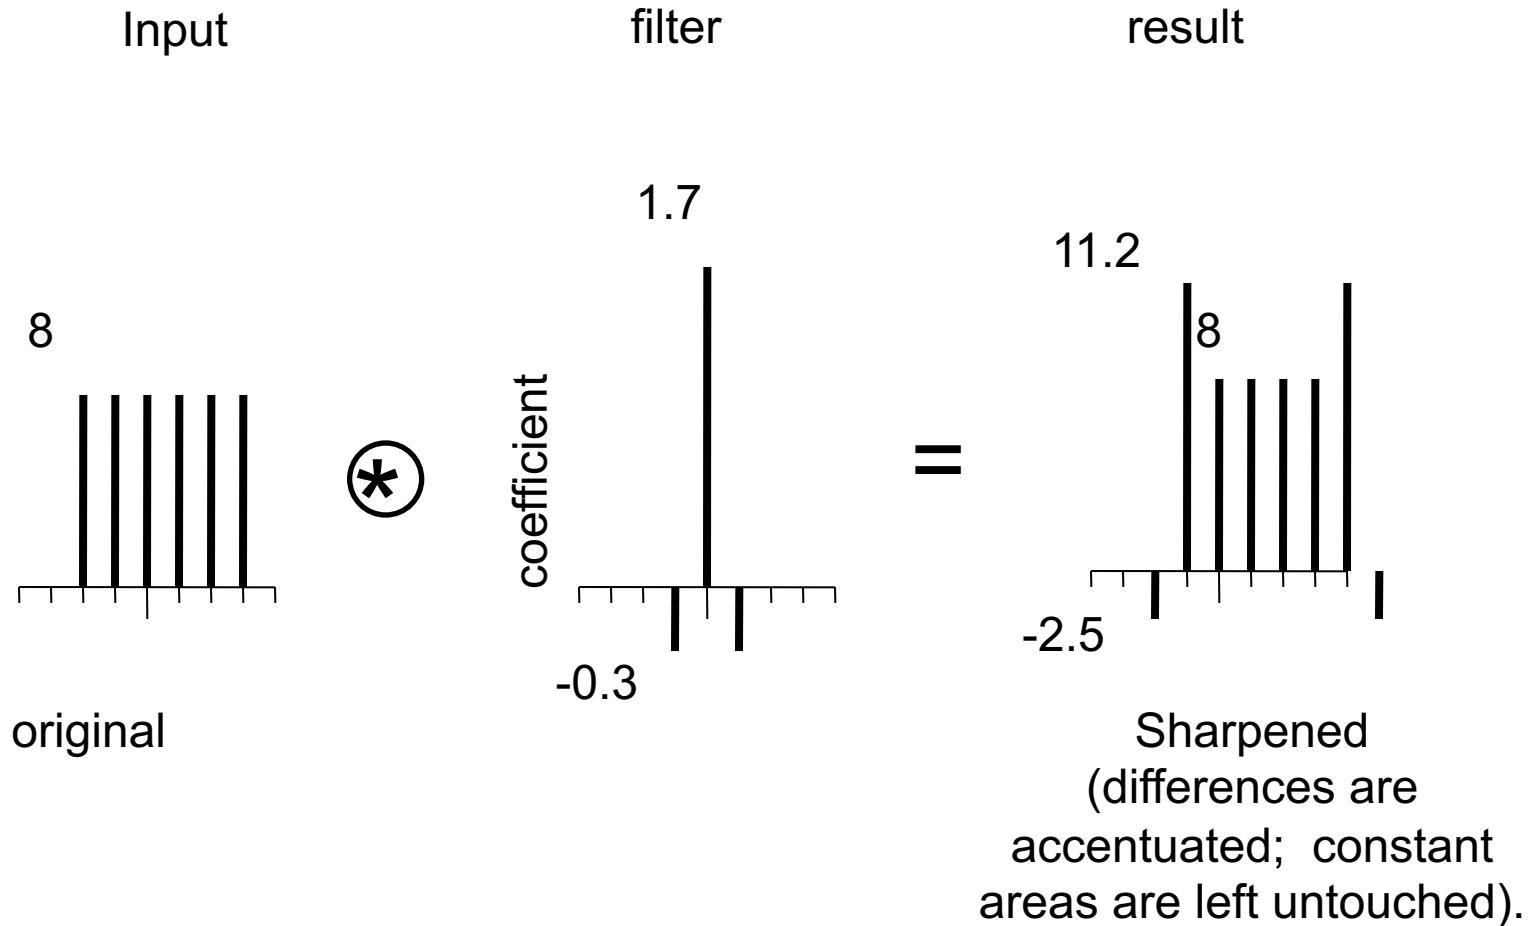


before



after

1-D sharpening example



$$[-1 \ 1]$$

$$\frac{\partial \mathbf{I}}{\partial x} \simeq \mathbf{I}(x, y) - \mathbf{I}(x - 1, y)$$



$g[m,n]$

\otimes

$$[-1, 1]$$

$=$

$h[m,n]$



$f[m,n]$

$$[-1 \ 1]^T$$

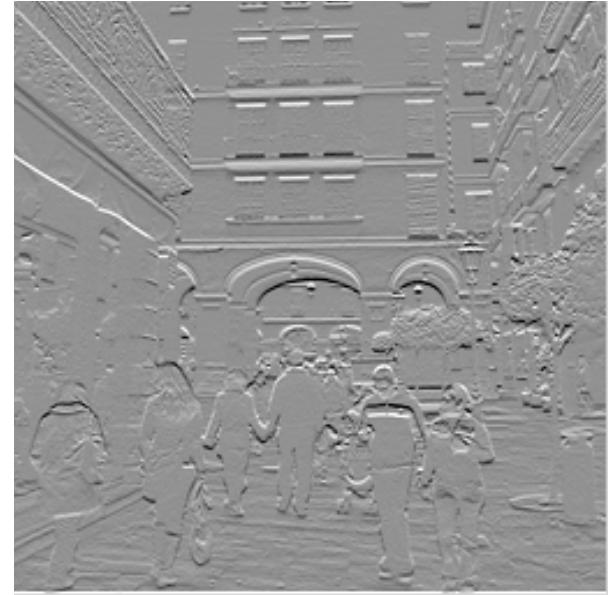


$g[m,n]$



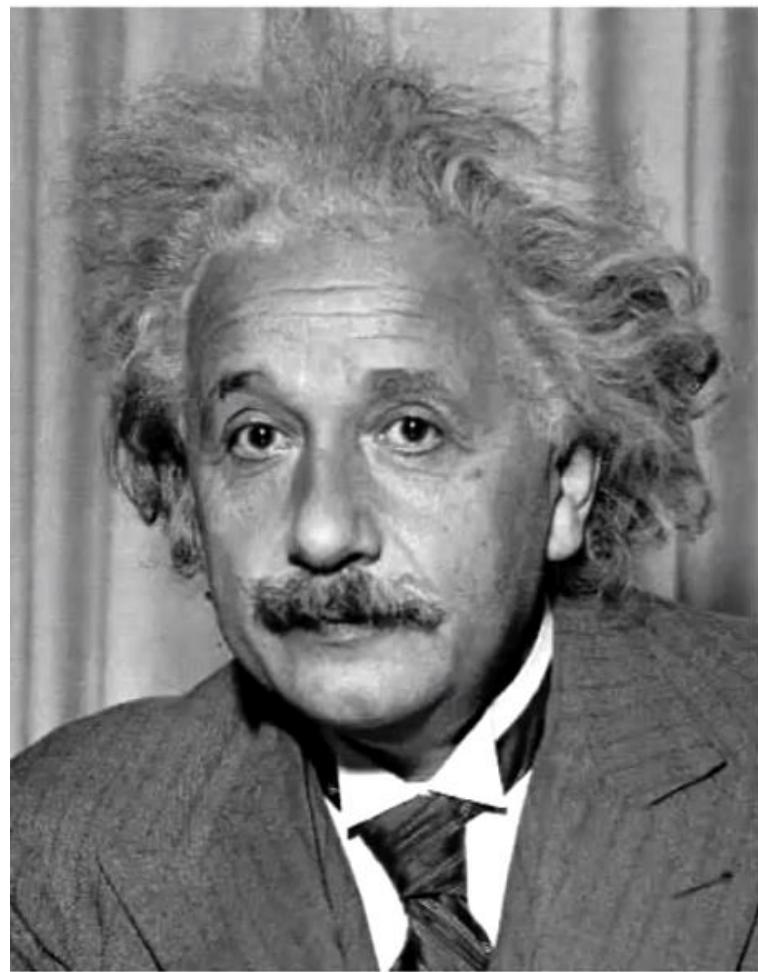
$$[-1, 1]^T =$$

$h[m,n]$



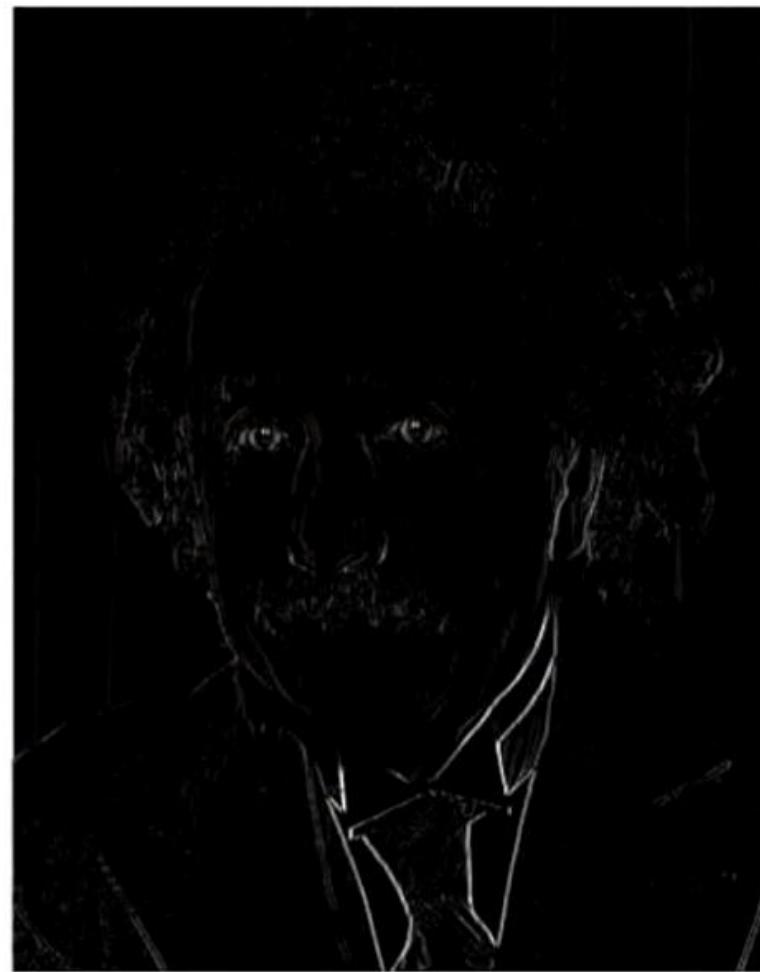
$f[m,n]$

Derivative filters



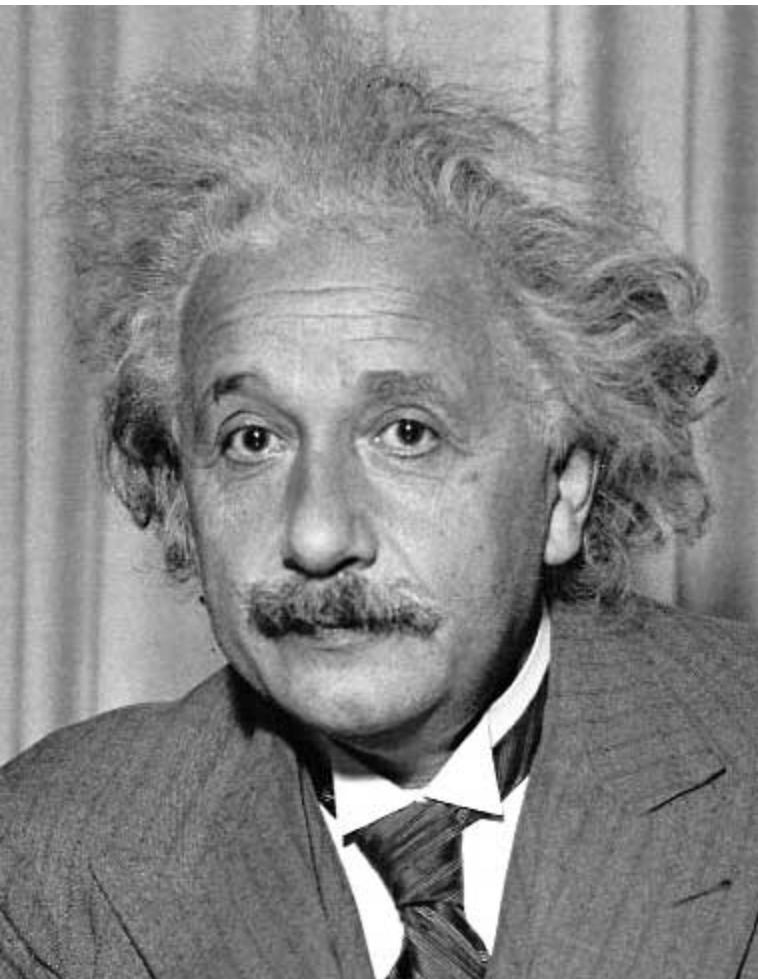
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

Derivative filters



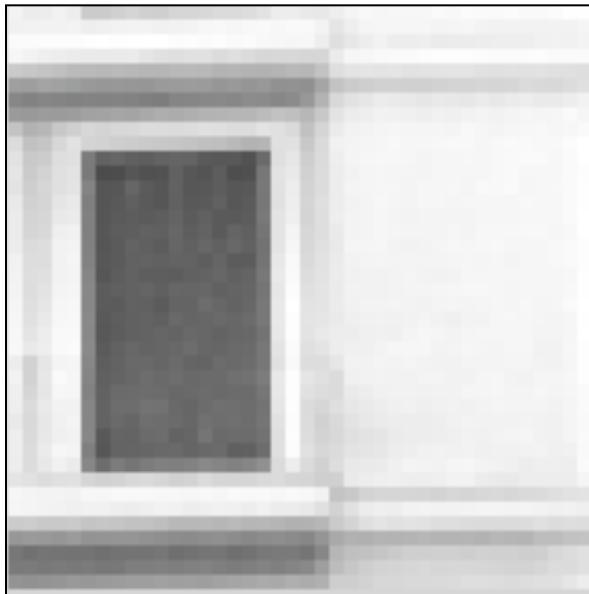
1	2	1
0	0	0
-1	-2	-1

Sobel



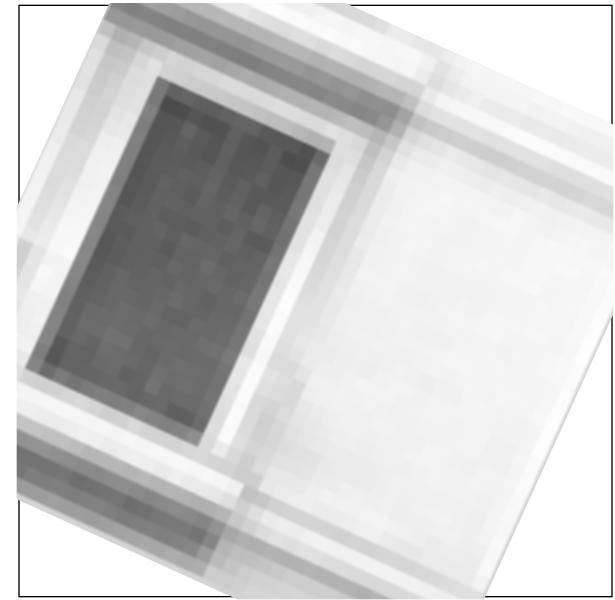
Horizontal Edge
(absolute value)

Image rotation

 \otimes

?

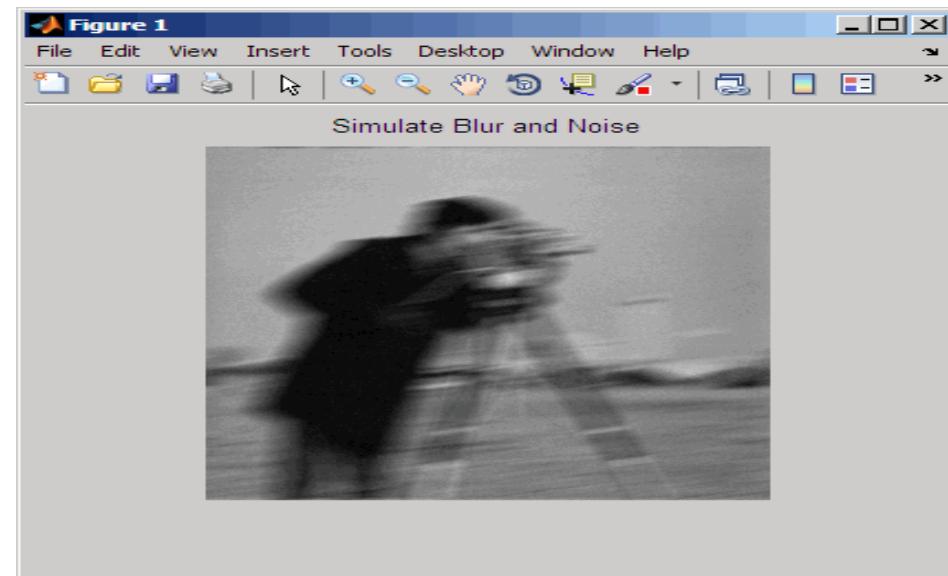
=

 $h[m,n]$  $g[m,n]$ $f[m,n]$

It is linear, but not a spatially invariant operation. There is not convolution.

How could we synthesize motion blur?

```
theta = 30; len = 20;  
fil = imrotate(ones(1, len), theta, 'bilinear');  
fil = fil / sum(fil(:));  
figure(2), imshow(imfilter(im, fil));
```



Motion blur filter



$g[m,n]$



$h[m,n]$



$f[m,n]$

Motion blur filter



$g[m,n]$

\otimes

$h[m,n]$



=



$f[m,n]$

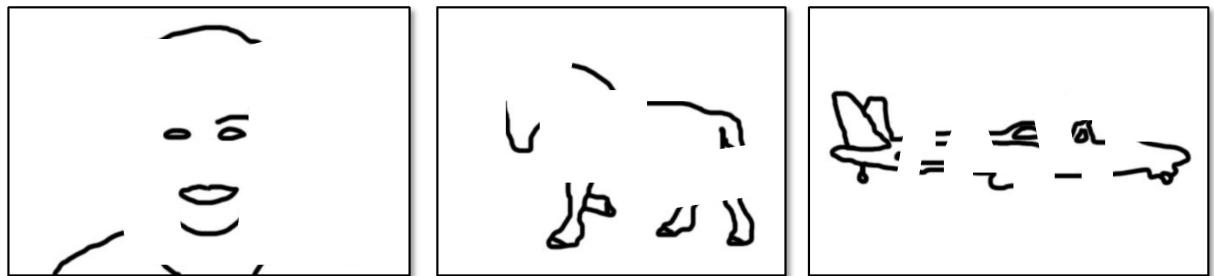
Three views of linear filtering

- Linear filters in the spatial domain
 - Filter is a mathematical operation on matrix
 - Smoothing, sharpening, measuring texture
- Linear filters in the **frequency domain [Self-study]**
 - Filtering is a way to modify the frequencies of images
 - Denoising, sampling, image compression
- Linear filters as **templates matching**
 - Filtering is a way to match a template to the image
 - Detection, registration

Edge Detection

Edge detection

- **Goal:** map image from 2d array of pixels to a set of curves or line segments or contours.
- **Why?**



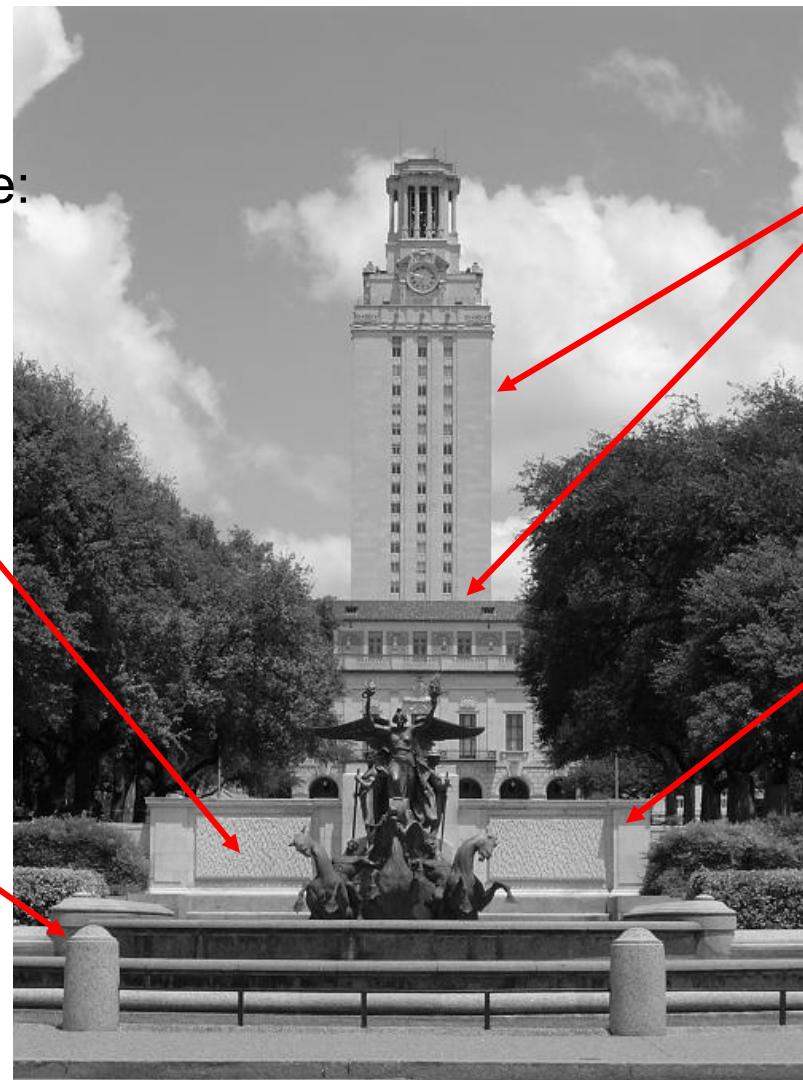
•Figure from J. Shotton et al., PAMI 2007

- **Main idea:** look for strong gradients, post-process

What causes an edge?

- Reflectance change: appearance information, texture

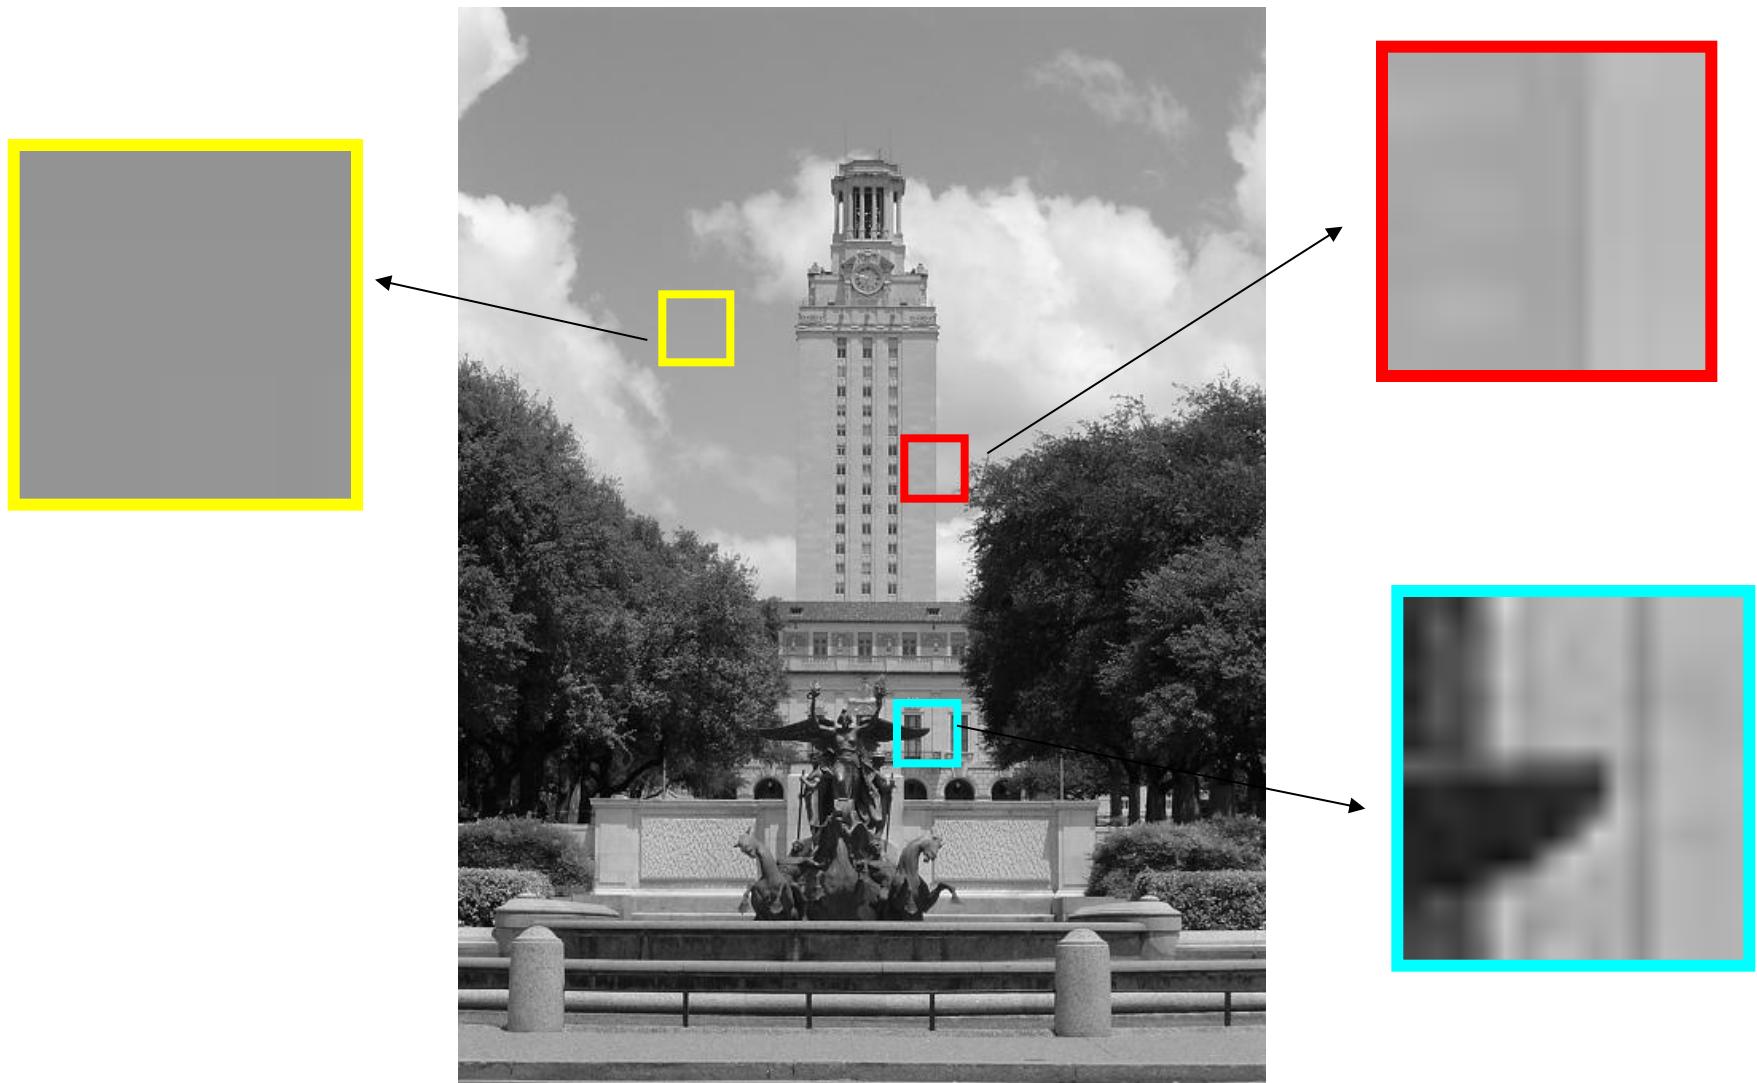
- Change in surface orientation: shape



- Depth discontinuity: object boundary

- Cast shadows

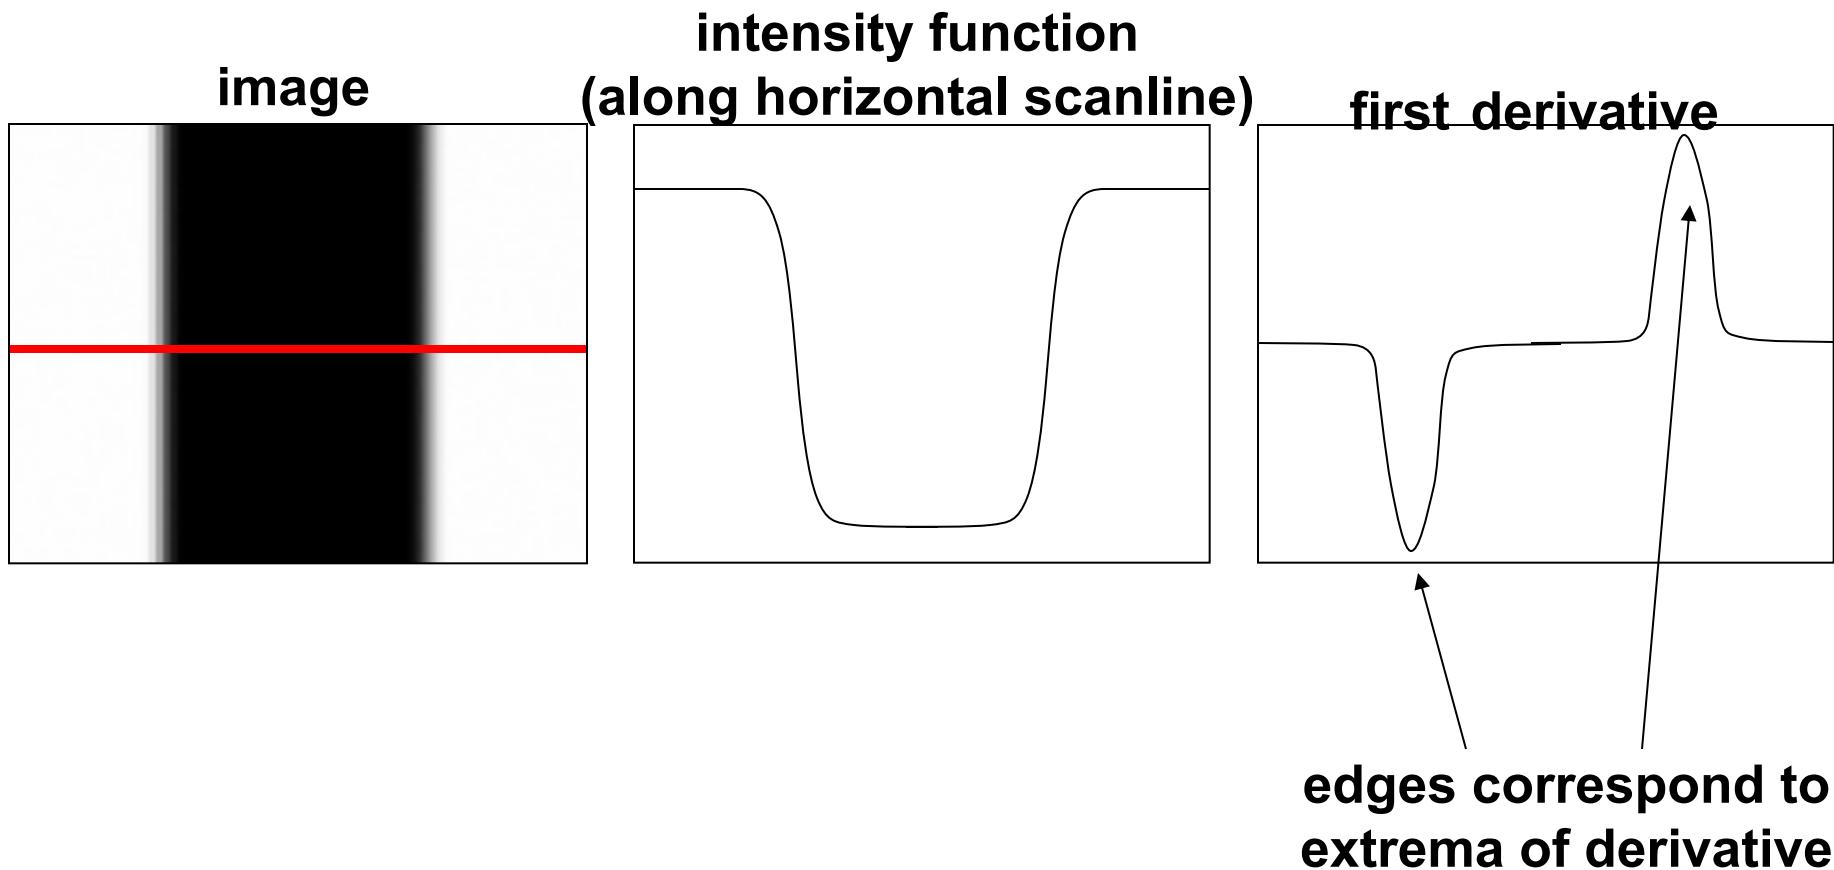
Edges/gradients and invariance



•Slide credit:
Kristen Grauman

Derivatives and edges

- An edge is a place of rapid change in the image intensity function.



Methods

- Using first order derivative: *Look for extrema*
 - Sobel operator : `edge(I, 'sobel');`
 - Prewitt, Roberts,...
 - Derivative of Gaussian
- Using second order derivative: *Look for zero-crossings*
 - Laplacian : isotropic
 - Second directional derivative $\nabla^2 \mathbf{I}$
 - LOG/DOG (Laplace of Gaussian/Difference of Gaussians)

Common edge detectors

- Basic gradient edge detectors
 - **Sobel**, Prewitt, Gaussian derivative ...
- **LoG/Marr** edge detector
- **Canny** edge detector

Derivatives with convolution

For 2D function, $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

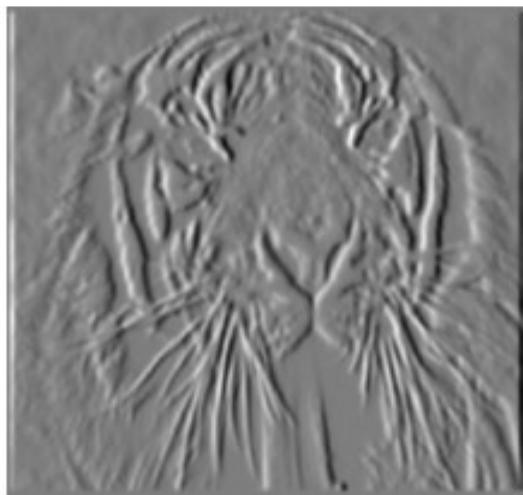
To implement above as convolution, what would be the associated filter?

Partial derivatives of an image



$$\frac{\partial f(x, y)}{\partial x}$$

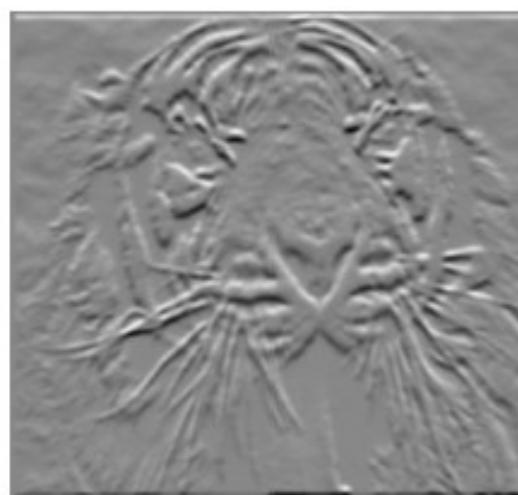
-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

or

-1	1
1	-1



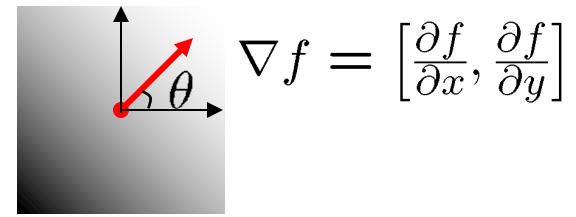
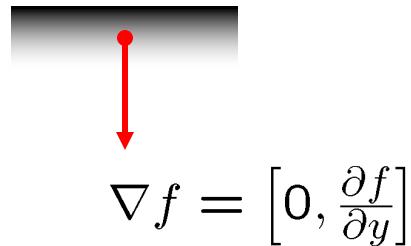
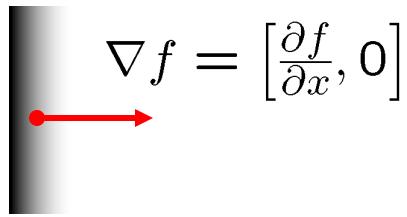
- Which shows changes with respect to x?
 - (showing filters for correlation)

Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity



- The **gradient direction** (orientation of edge normal) is given by:

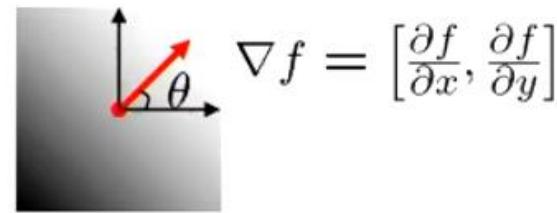
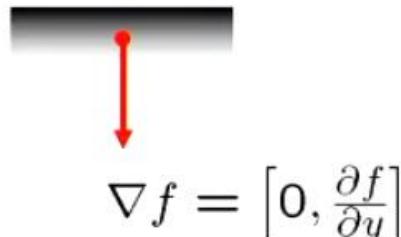
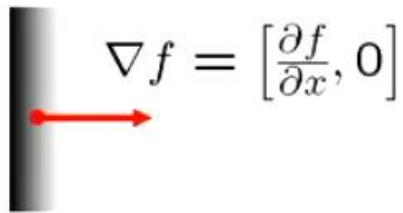
$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity



- The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- The **edge strength** is given by the gradient magnitude

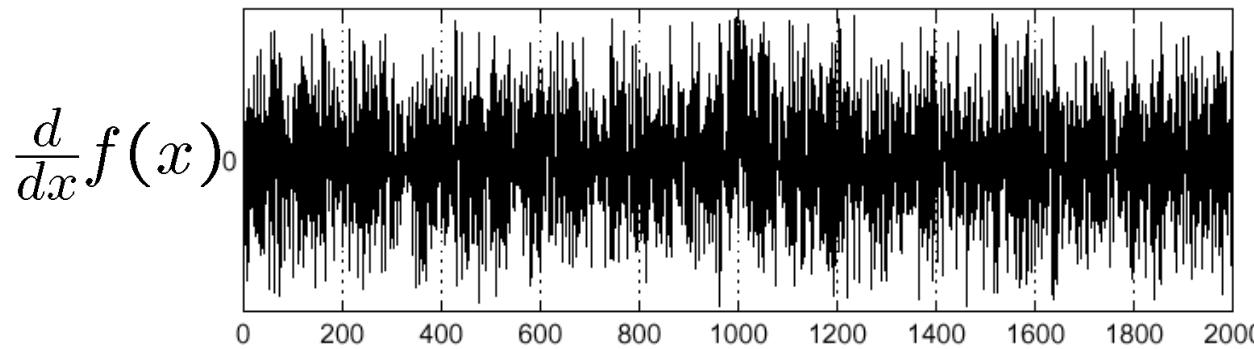
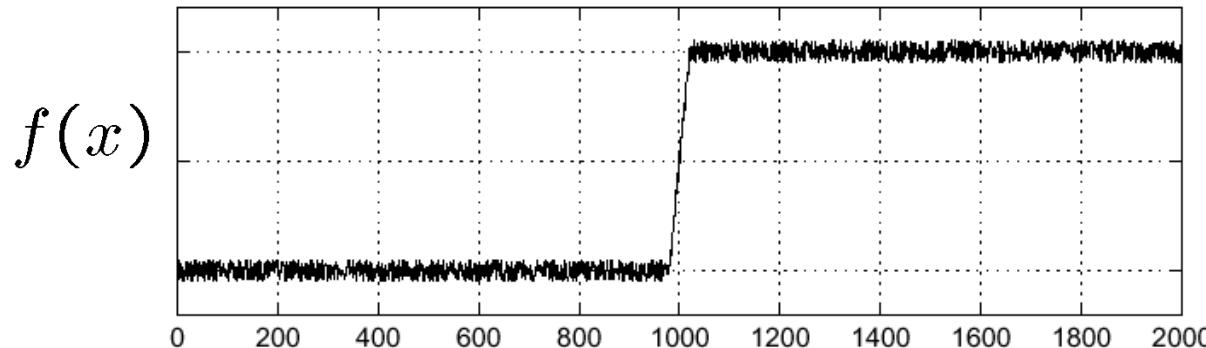
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$



Effects of noise

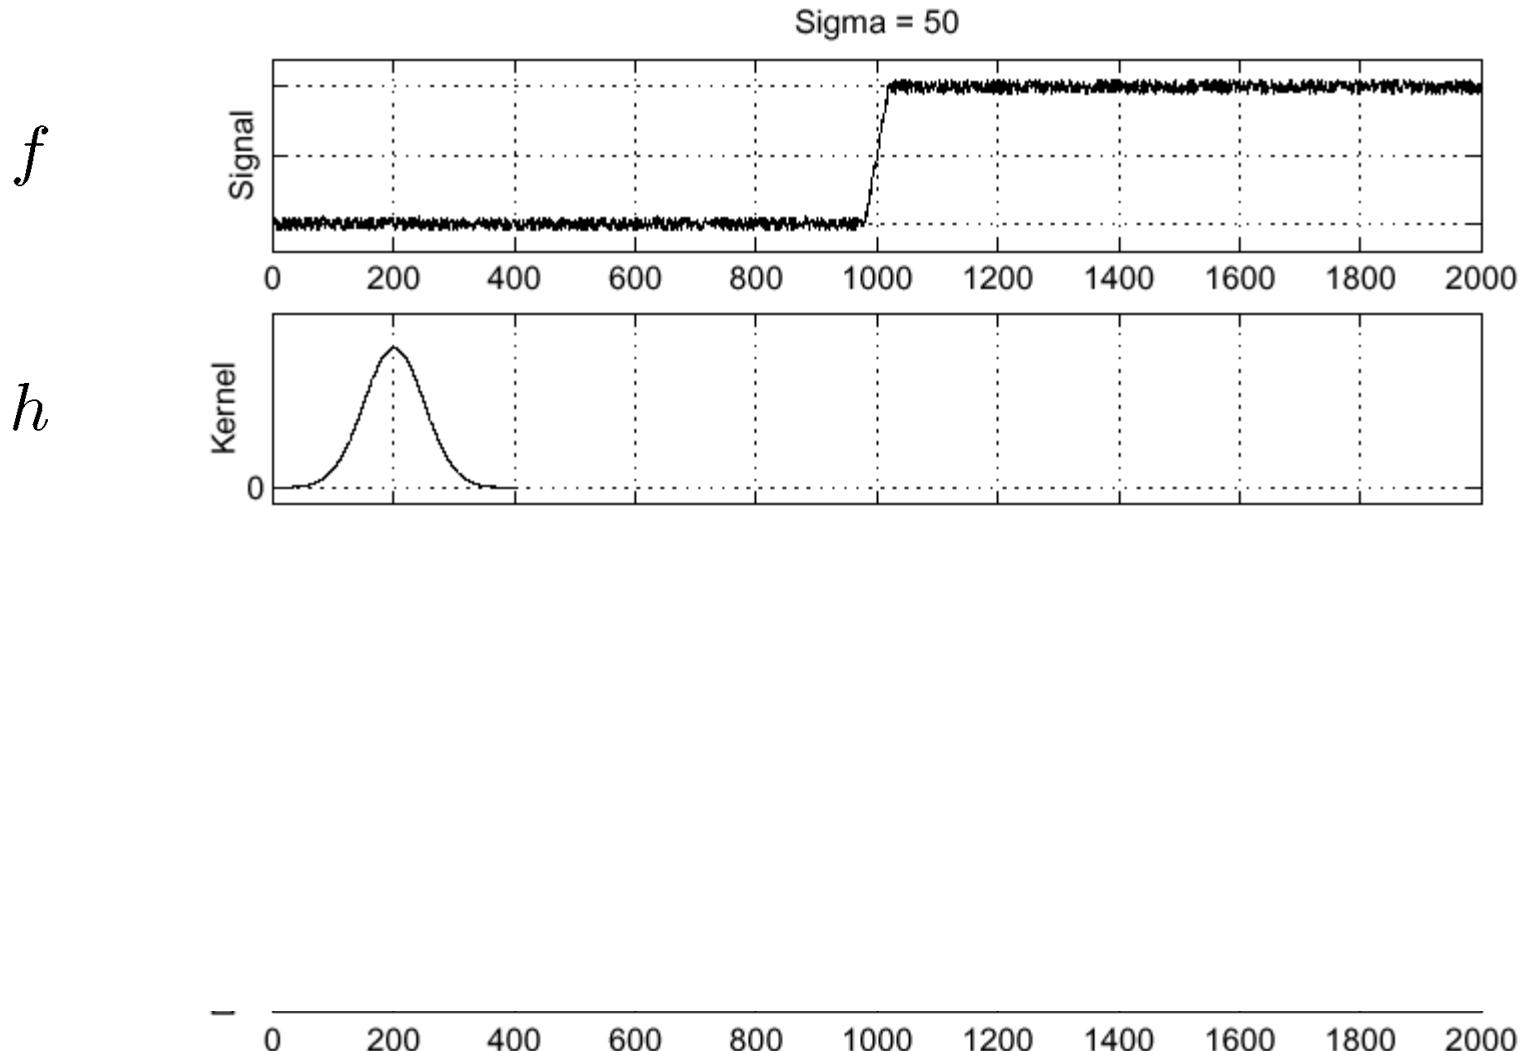
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



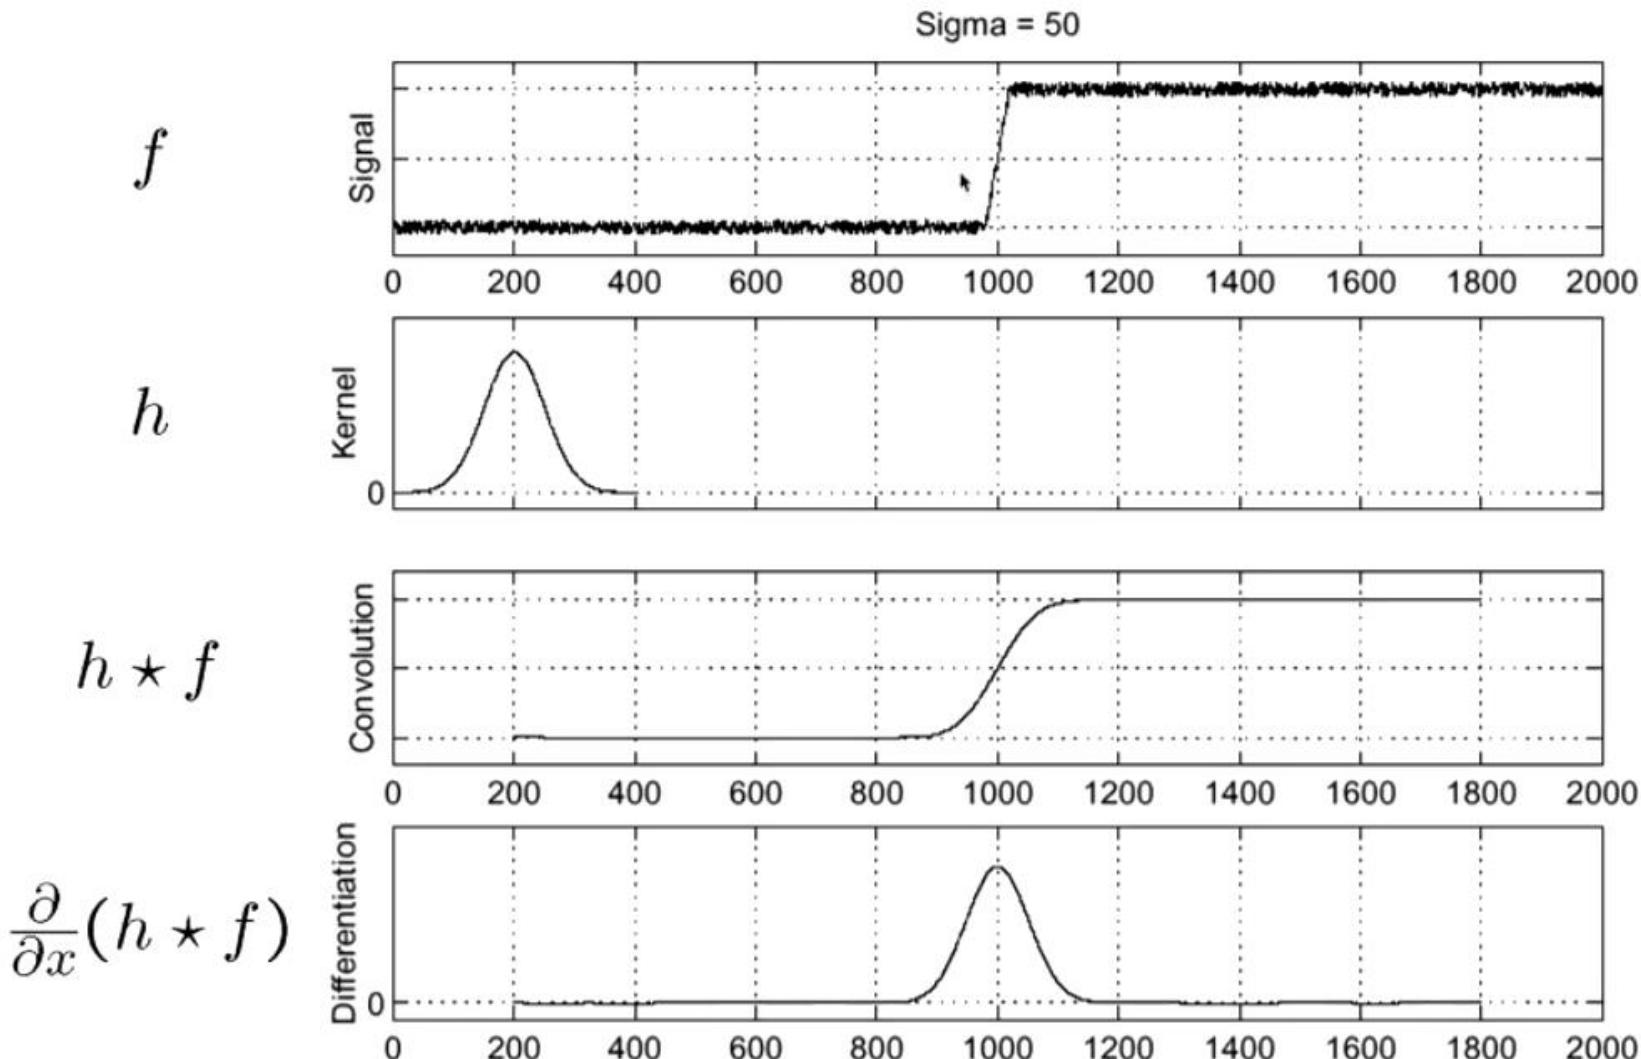
- Where is the edge?

Solution: smooth first



- Where is the edge?
- Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Solution: smooth first



- Where is the edge?

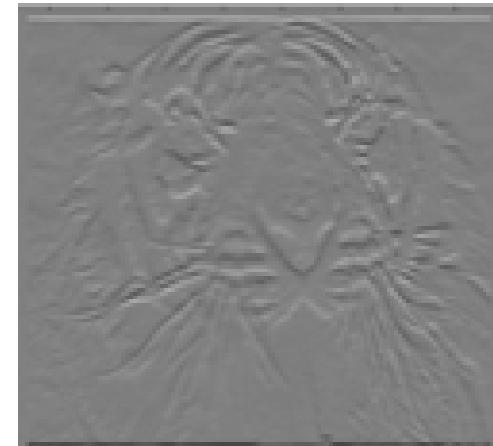
Gradient Edge Detectors

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

```
>> My = fspecial('sobel');  
>> outim = imfilter(double(im), My);  
>> imagesc(outim);  
>> colormap gray;
```



Eg: Convolute an image with a Sobel kernel

K

1	0	-1
2	0	-2
1	0	-1

• Rotate ↓

-1	0	1
-2	0	2
-1	0	1

I

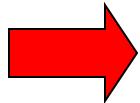
0	0	2	2
0	0	2	2
0	0	2	2
0	0	2	2

-1	0	1
-2	0	2
-1	0	1

Step 1

0	0	2	2
0	0	2	2
0	0	2	2
0	0	2	2

-1	0	1		
-2	0	0	2	2
-1	0	0	2	2
0	0	1	2	
0	0	2	2	



0			

I

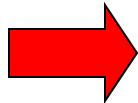
I'

-1	0	1
-2	0	2
-1	0	1

Step 2

0	0	2	2
0	0	2	2
0	0	2	2
0	0	2	2

-1	0	1	
0	0	4	2
0	0	2	2
0	0	2	2
0	0	2	2



0	6		

I

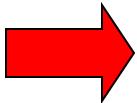
I'

-1	0	1
-2	0	2
-1	0	1

Step 3

0	0	2	2
0	0	2	2
0	0	2	2
0	0	2	2

	-1	0	1
0	0	0	4
0	0	0	2
0	0	2	2
0	0	2	2



0	6	6	

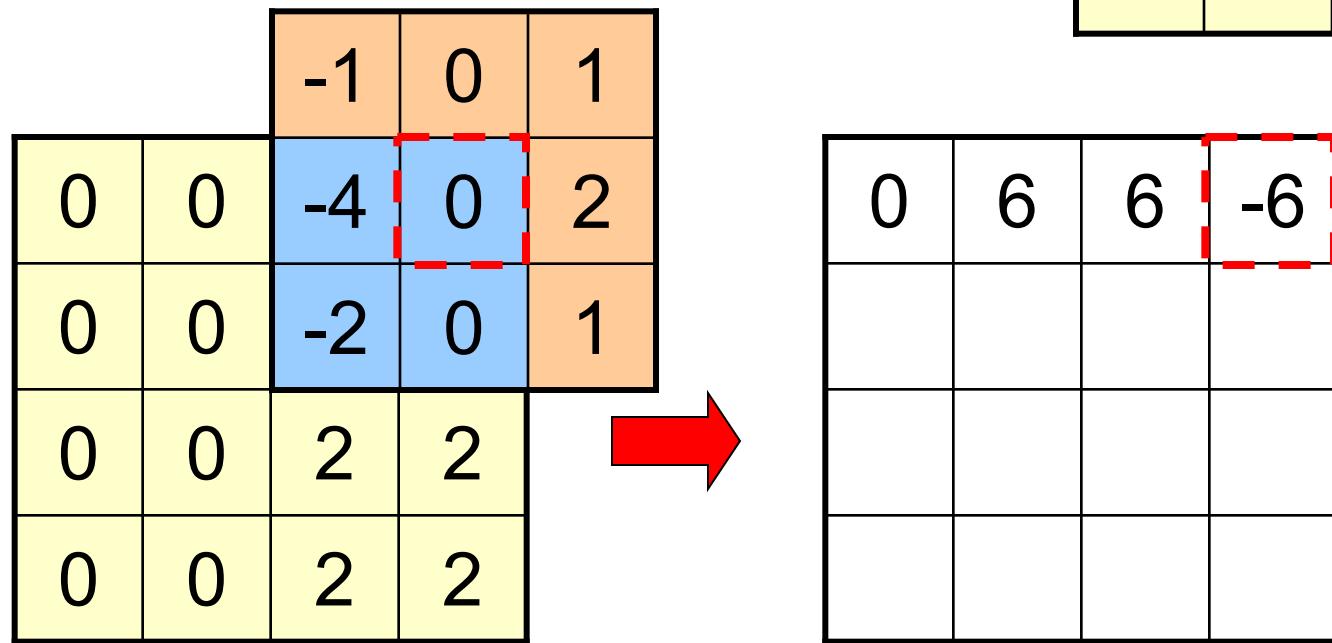
I

I'

-1	0	1
-2	0	2
-1	0	1

Step 4

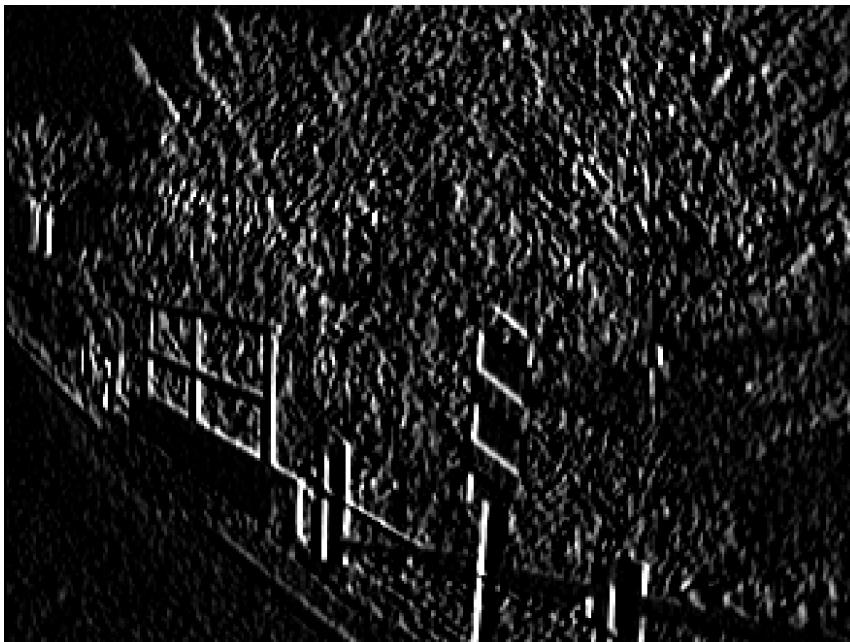
0	0	2	2
0	0	2	2
0	0	2	2
0	0	2	2



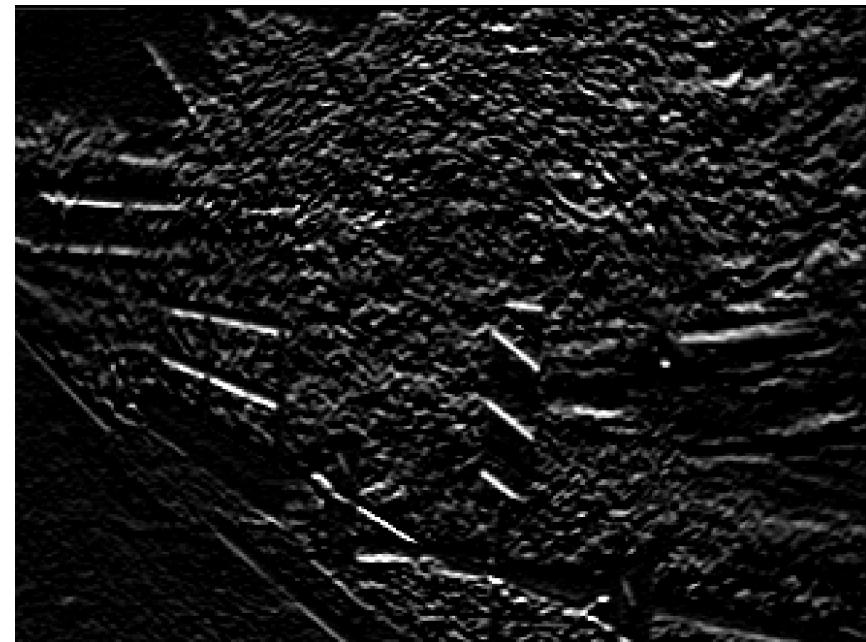
I

I'

Sobel edge detection: example



- Vertical edges

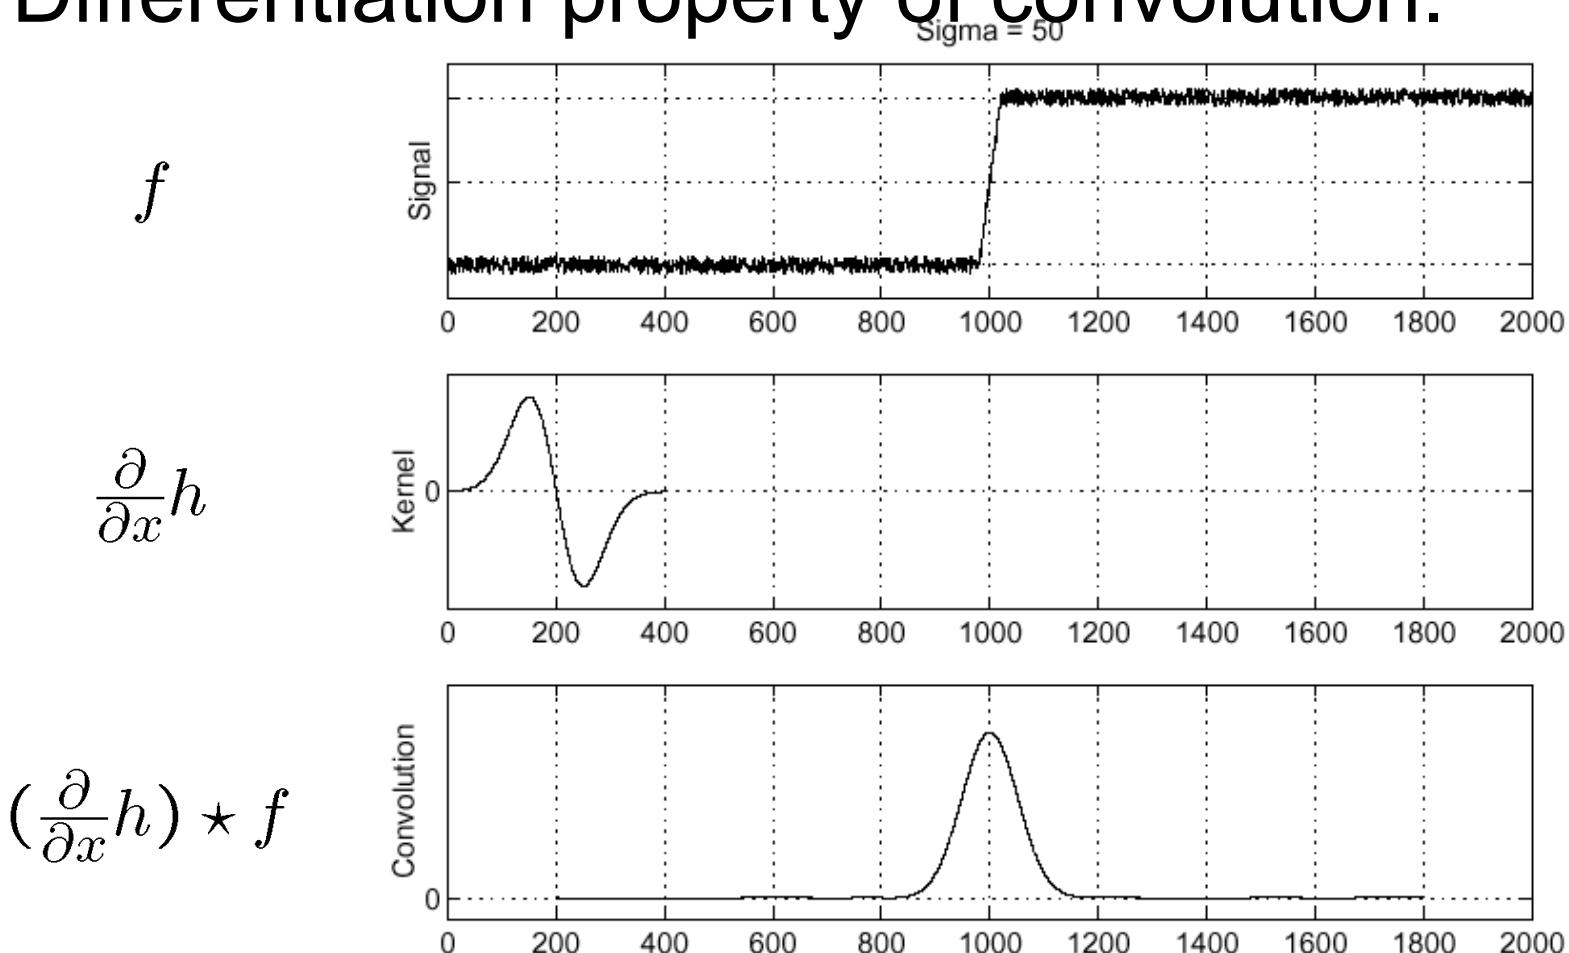


- Horizontal edges

Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

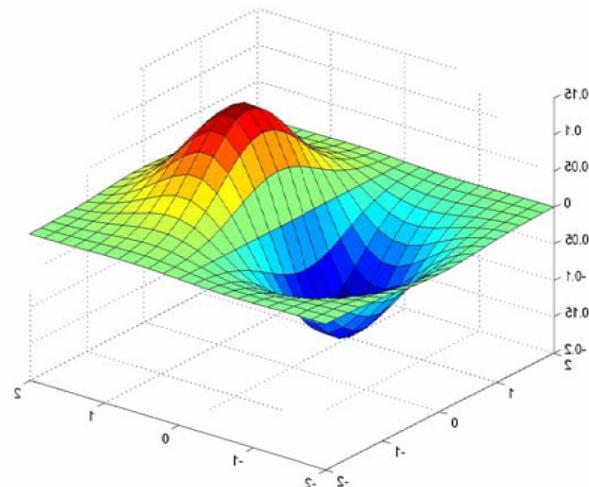
Differentiation property of convolution.



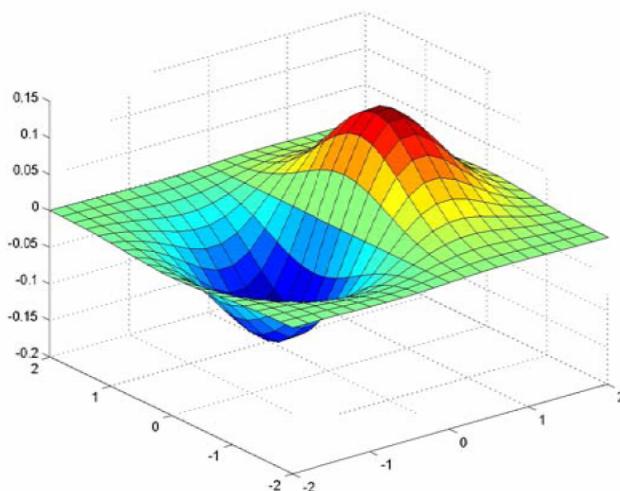
Derivative of Gaussian filters

$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

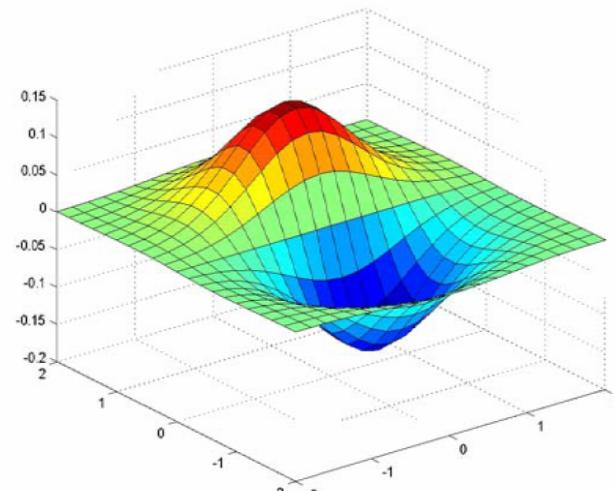
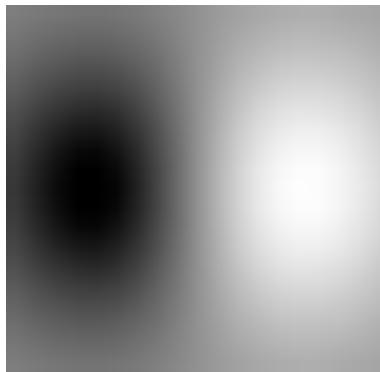
$$\begin{bmatrix} \cdot & 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ \cdot & 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ \cdot & 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ \cdot & 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ \cdot & 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} \otimes \begin{bmatrix} 1 & -1 \end{bmatrix}$$



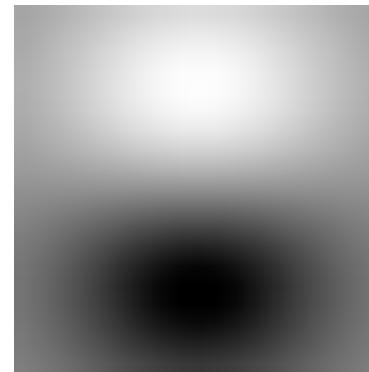
Derivative of Gaussian filters



•x-direction



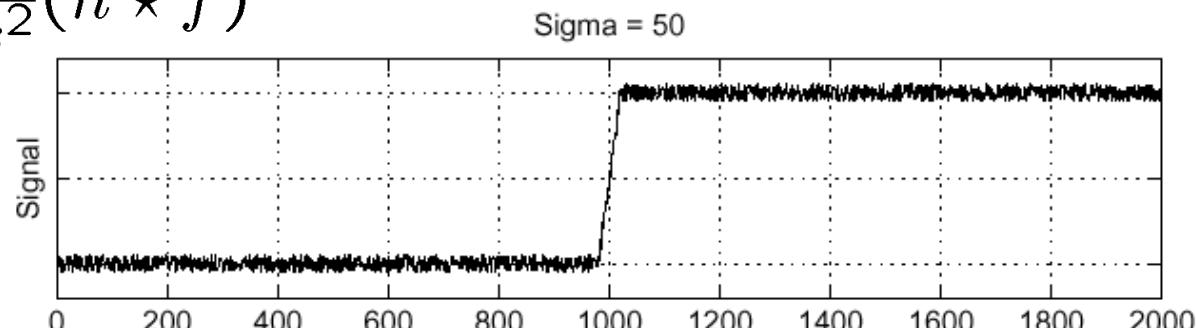
•y-direction



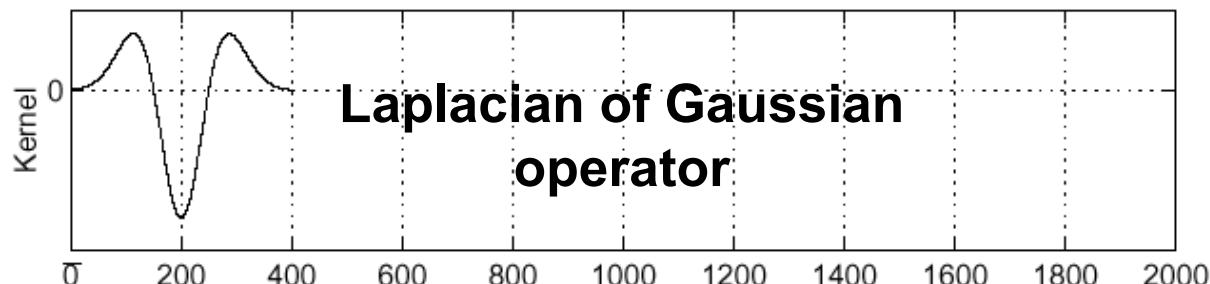
Laplacian of Gaussian

Consider $\frac{\partial^2}{\partial x^2}(h \star f)$

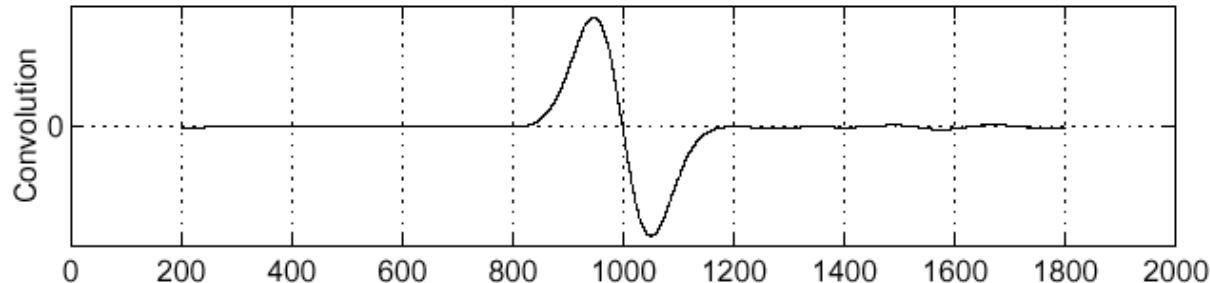
f



$\frac{\partial^2}{\partial x^2} h$



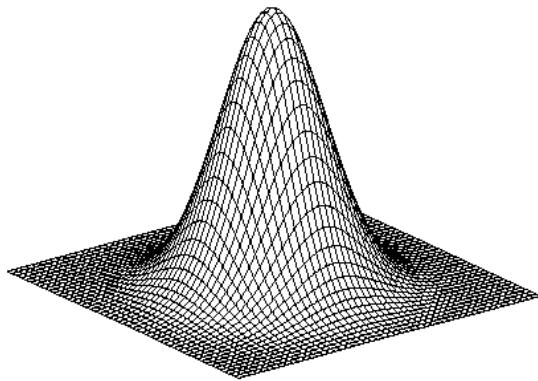
$(\frac{\partial^2}{\partial x^2} h) \star f$



- Where is the edge?
- Zero-crossings of bottom graph

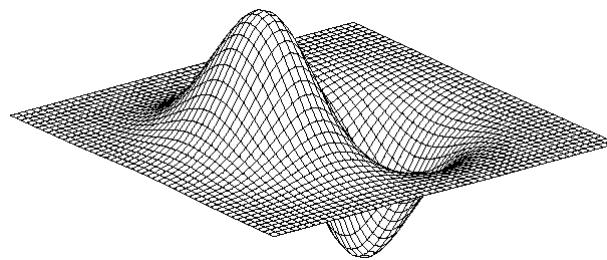
• Slide credit: Steve Seitz

2D edge detection filters



Gaussian

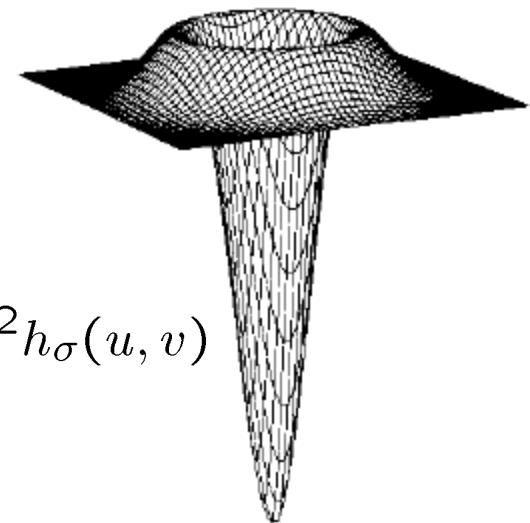
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



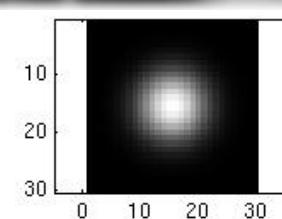
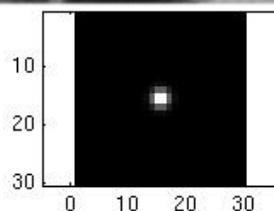
$$\nabla^2 h_\sigma(u, v)$$

- ∇^2 is the Laplacian operator:

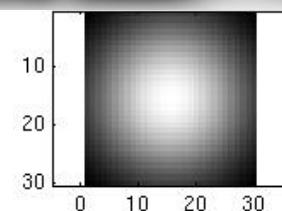
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Smoothing with a Gaussian

- Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



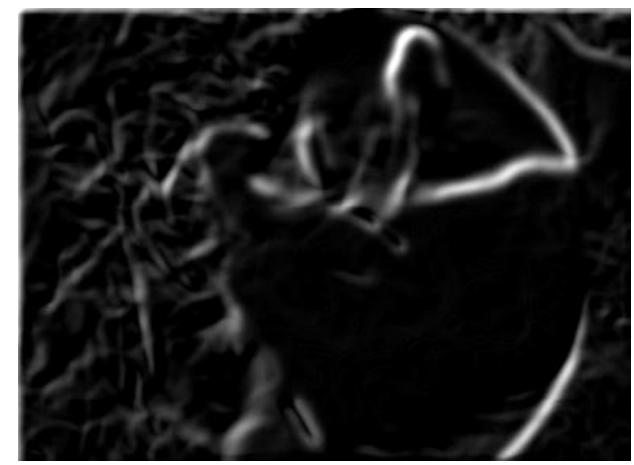
• ...



Effect of σ on derivatives



$\sigma = 1$ pixel

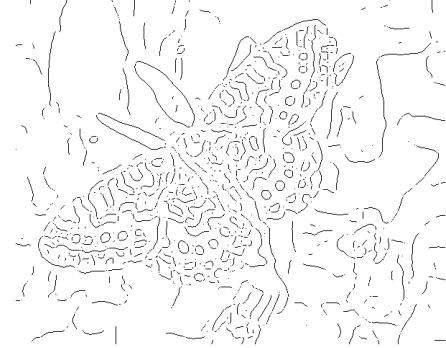


$\sigma = 3$ pixels

- The apparent structures differ depending on Gaussian's scale parameter.
- Larger values: larger scale edges detected
- Smaller values: finer features detected



Gradients -> edges



Primary edge detection steps:

1. Smoothing: suppress noise
2. Edge enhancement: filter for contrast
3. Edge localization

Determine which local maxima from filter output
are actually edges vs. noise

- Threshold, Thin

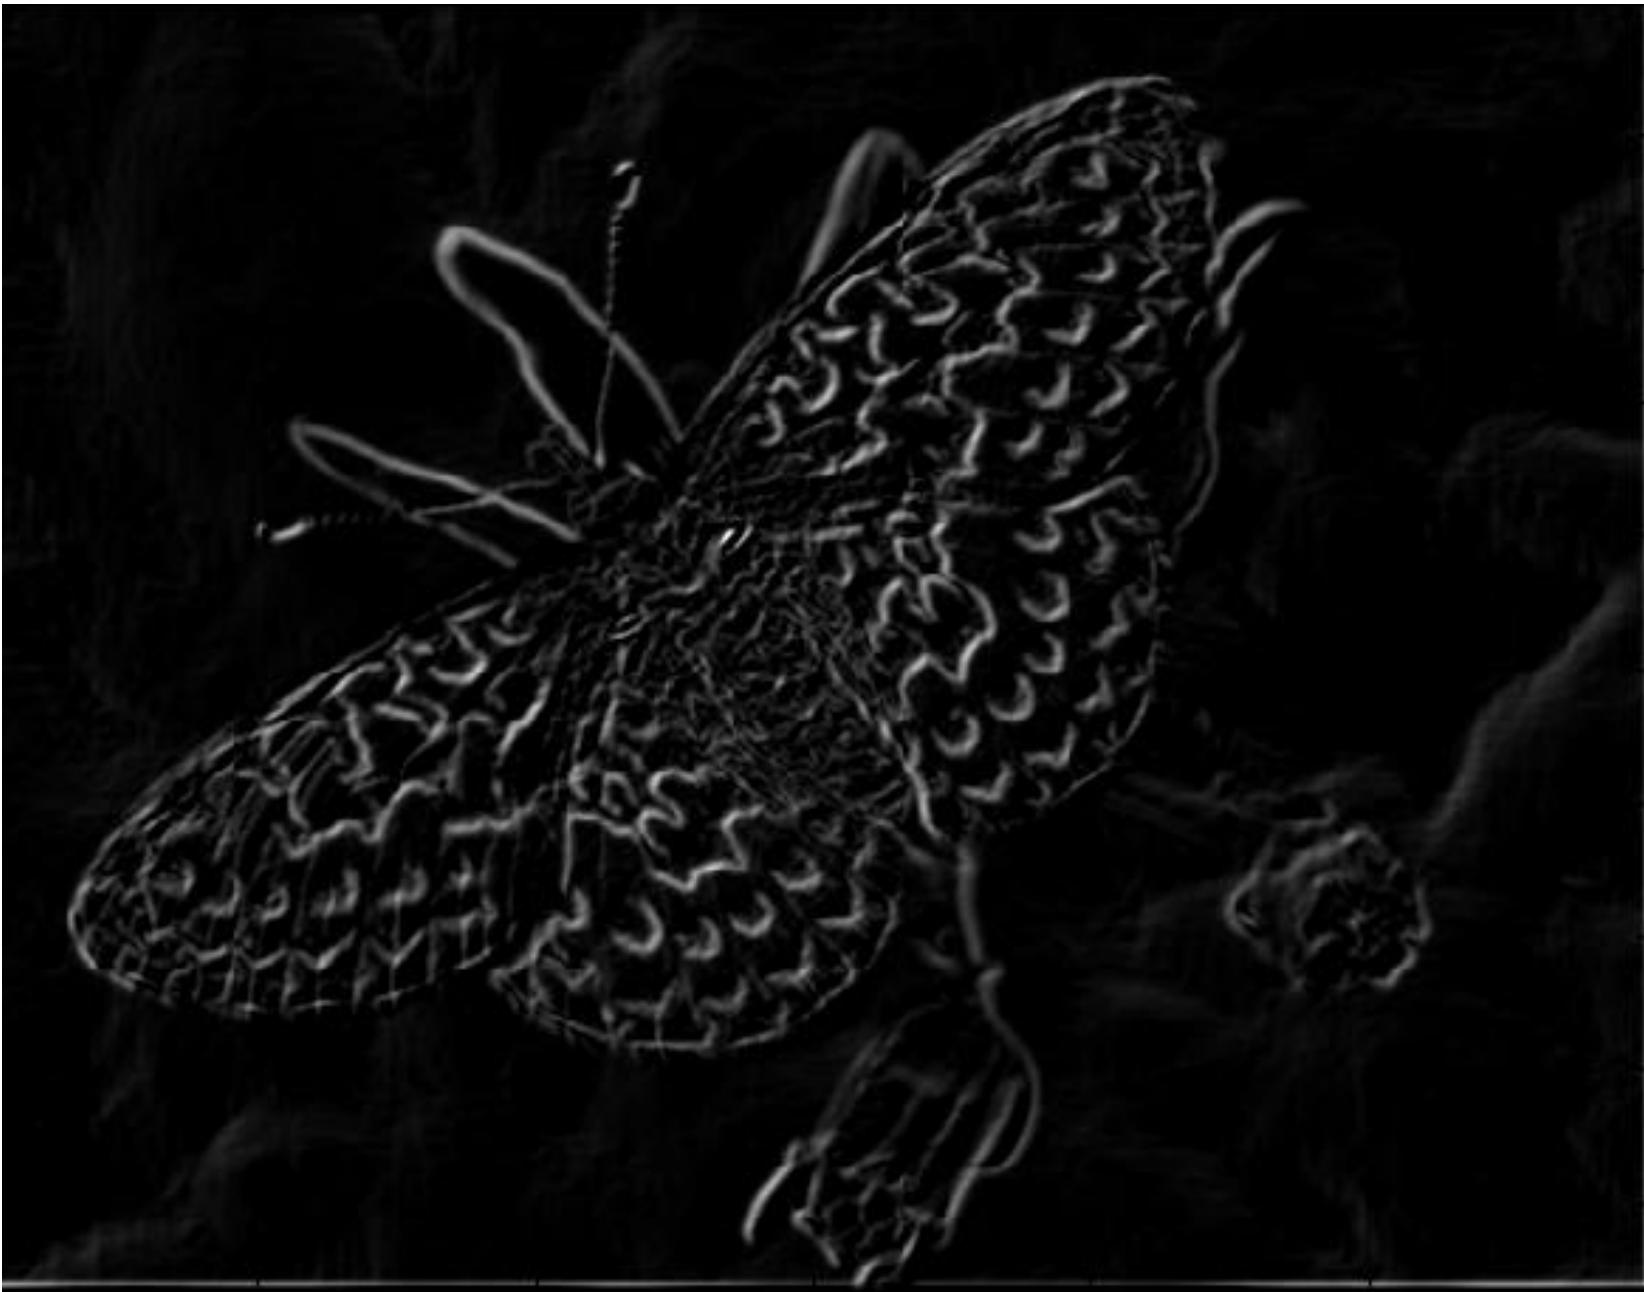
Thresholding

- Choose a threshold value t
- Set any pixels less than t to zero (off)
- Set any pixels greater than or equal to t to one (on)

Original image



Gradient magnitude image



Thresholding gradient with a lower threshold



Thresholding gradient with a higher threshold



Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny')` ;
- `>>help edge`

The Canny edge detector



original image (Lena)

The Canny edge detector



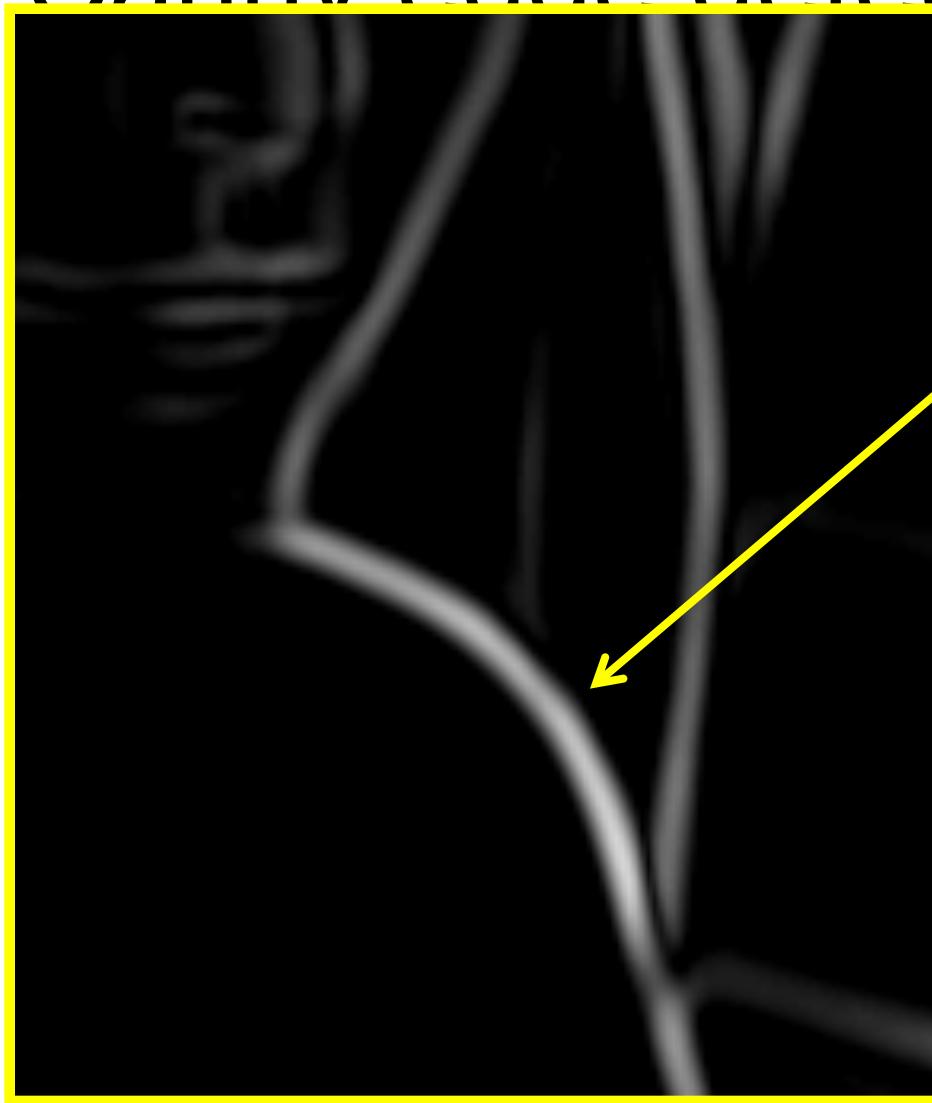
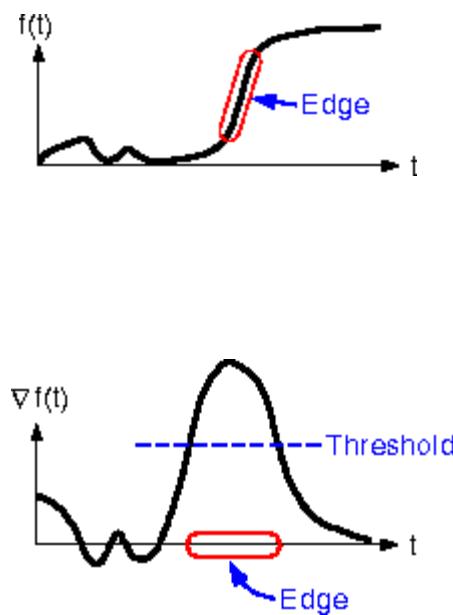
- norm of the gradient

The Canny edge detector



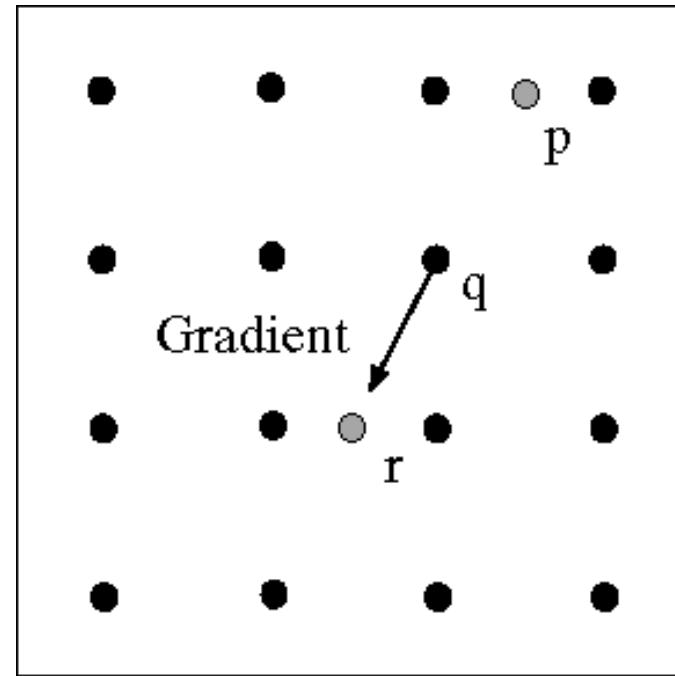
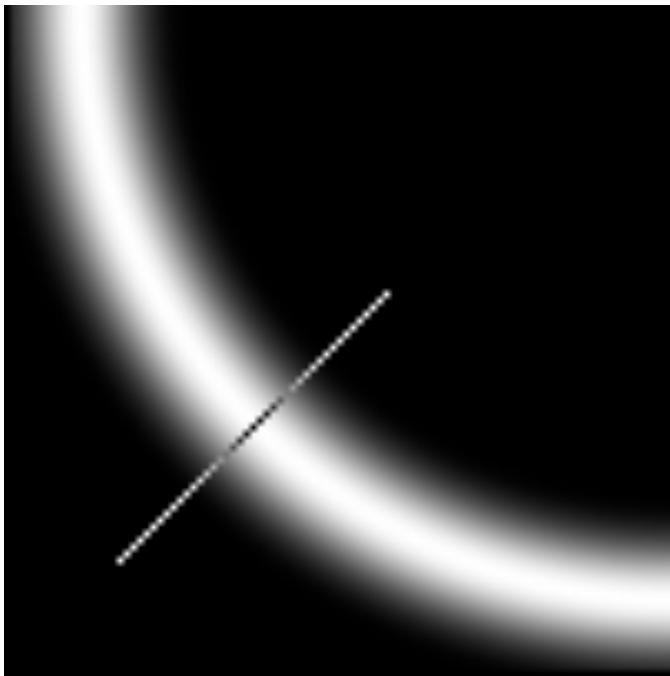
- thresholding

The Canny edge detector



- How to turn these thick regions of the gradient into curves?

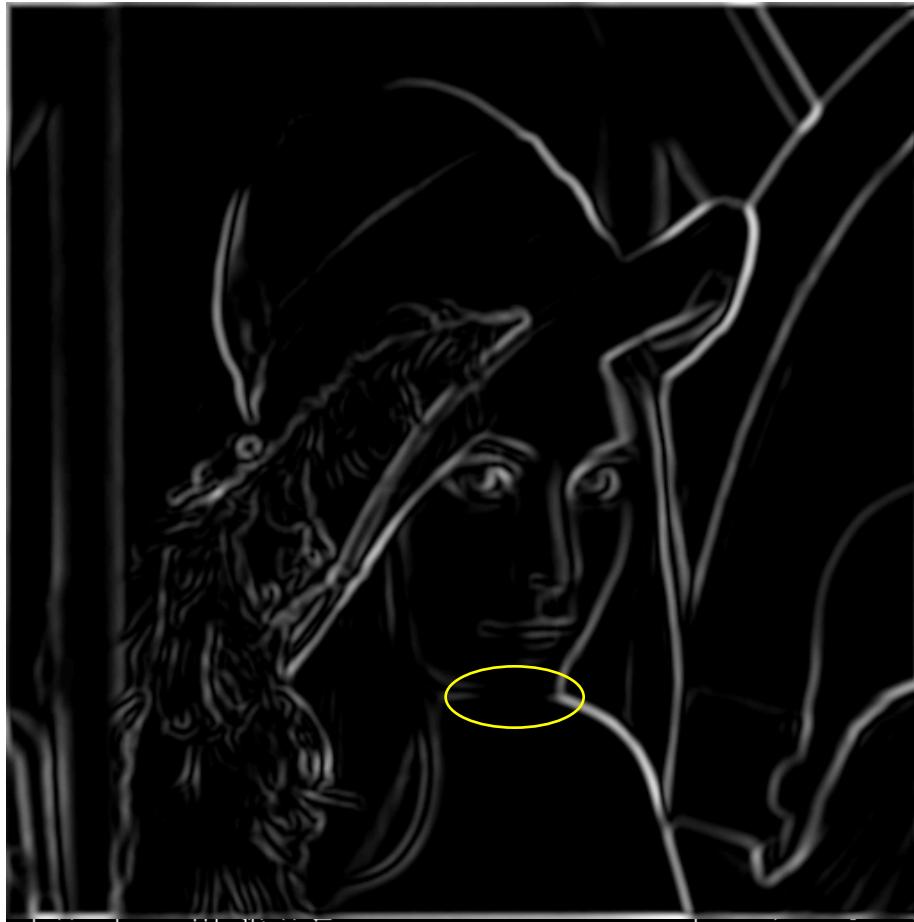
Non-maximum suppression



Check if pixel is local maximum along gradient direction, select single max across width of the edge

- requires checking interpolated pixels p and r

The Canny edge detector



- thinning
- (non-maximum suppression)

- Problem:
pixels along
this edge
didn't
survive the
thresholding

Hysteresis thresholding

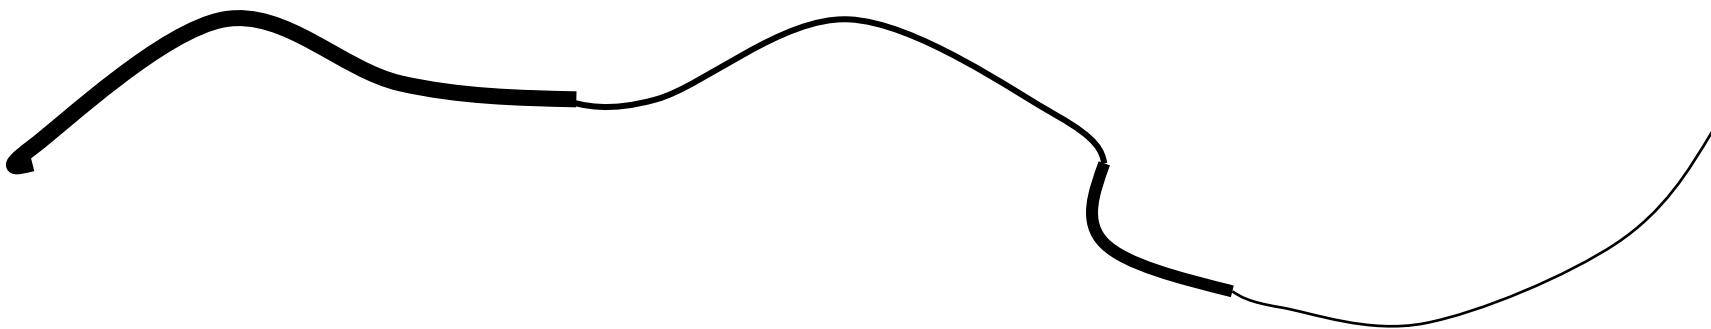
- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



•Credit: James Hay

Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them.



Final Canny Edges



•Credit: Jame

Recap: Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny')` ;
- `>>help edge`

Mask properties

- Smoothing
 - Values positive
 - Sum to 1 → constant regions same as input
 - Amount of smoothing proportional to mask size
 - Remove “high-frequency” components; “low-pass” filter
- Derivatives
 - Opposite signs used to get high response in regions of high contrast
 - Sum to 0 → no response in constant regions
 - High absolute value at points of high contrast